

[Draw your reader in with an engaging abstract. It is typically a short summary of the document. When you're ready to add your content, just click here and start typing.]

# Thor2023

[Document subtitle]

Robert Finch

---

## Table of Contents

Thor2023 .....	9
Nomenclature .....	9
Little Endian vs big Endian .....	11
Endian .....	11
Block Diagram .....	12
Programming Model .....	13
Register File .....	13
Number of Registers .....	<b>Error! Bookmark not defined.</b>
Rn – General Purpose Registers .....	13
Vn – Vector Registers .....	<b>Error! Bookmark not defined.</b>
Predicate Registers .....	15
Mask Registers (vm0 to vm7) .....	15
Vector Length (VL register) .....	15
Code Address Registers .....	16
LC - Loop Counter (reg 55) .....	16
SR - Status Register (CSR 0x?004) .....	16
Special Purpose Registers .....	18
Operating Modes .....	24
Exceptions .....	25
External Interrupts .....	25
Effect on Machine Status .....	25
Exception Stack .....	25
Reset .....	26
Precision .....	27
Memory Management .....	28
Bank Swapping .....	28
The Page Map .....	28
Regions .....	28
PMA - Physical Memory Attributes Checker .....	28
Overview .....	28
Region Table Description .....	29
Page Management Table - PMT .....	31
Overview .....	31

Location .....	31
PMTE Description .....	31
Access Control List.....	31
Share Count.....	31
Access Count .....	31
Key .....	32
Privilege Level.....	32
N.....	32
M.....	32
E.....	32
AL .....	32
C.....	32
Page Tables .....	33
Intro.....	33
Hierarchical Page Tables .....	33
Inverted Page Tables.....	33
The Simple Inverted Page Table.....	34
Hashed Page Tables .....	34
Shared Memory.....	35
Specifics: Thor2023 Page Tables.....	35
Thor2023 Hash Page Table Setup.....	35
Thor2023 Hierarchical Page Table Setup .....	38
TLB – Translation Lookaside Buffer.....	41
Overview.....	41
Size / Organization.....	41
TLB Entries - TLBE .....	43
Small TLB Entries - TLBE .....	43
What is Translated?.....	43
Page Size.....	43
Ways .....	44
Management.....	44
Accessing the TLB.....	44
?RWX <sub>3</sub> .....	45
CACHE <sub>4</sub> .....	45
TLB Entry Replacement Policies.....	45

Flushing the TLB .....	45
Reset.....	46
Card Table.....	46
Overview .....	46
Organization.....	46
Location .....	47
Operation.....	47
Sample Write Barrier .....	47
System Memory Map.....	47
Instruction Set .....	48
Overview.....	48
Predicated Instruction Execution .....	48
Instruction Descriptions .....	49
Scalar Instructions Layout .....	49
Vector Instruction Layout .....	49
Register-Register Vector Decode.....	49
Register-Immediate Vector Decode.....	49
Opcode Maps .....	49
Major Opcode .....	50
{R2} Operations.....	50
{BIT – Func3} .....	50
{SHIFT – Func5}.....	51
{FLT2} Operations .....	51
{FLT1} Operations .....	51
{AMO} Operations.....	52
Operand Swapping.....	53
Operand Sizes .....	53
Arithmetic Operations.....	54
Representations .....	54
Arithmetic Operations.....	55
Logical Operations.....	55
Immediate Operate Functions .....	56
<b>Major Opcode</b> .....	56
ABS – Absolute Value.....	57
ADD - Addition .....	58

AND – Bitwise And.....	59
BMAP – Byte Map .....	60
CHK – Check Register Against Bounds .....	61
CLMUL – Carry-less Multiply .....	63
CMP - Comparison .....	64
CMPS.B – Signed Byte Comparison .....	65
CMPU.B – Unsigned Byte Comparison .....	66
CNTLZ – Count Leading Zeros .....	67
CNTLO – Count Leading Ones .....	67
CNTPOP – Count Population .....	68
CSR – Control and Special Registers Operations .....	69
DIVS – Signed Division .....	70
DIVU – Unsigned Division.....	71
EOR – Bitwise Exclusive Or .....	72
ENOR – Bitwise Exclusive Nor.....	73
PFX – Constant Postfix.....	82
MODS – Signed Modulus.....	74
MODU – Unsigned Modulus.....	75
MULS – Multiply Signed .....	76
MULU – Unsigned Multiplication.....	78
NAND – Bitwise And and Invert.....	79
NOR – Bitwise Or and Invert .....	80
OR – Bitwise Or.....	81
REVBIT – Reverse Bit Order .....	82
SEQ – Set if Equal .....	87
SGE – Set if Greater Than or Equal.....	88
SGT – Set if Greater Than .....	89
SLE – Set if Less Than or Equal.....	90
SLT – Set if Less Than .....	91
SNE – Set if Not Equal .....	92
SQRT – Square Root .....	93
SUB - Subtraction .....	94
Vector Arithmetic Operations .....	96
VADD - Addition .....	97
VAND – Bitwise And.....	98

VSHLV – Shift Vector Left .....	99
VSHRV – Shift Vector Right.....	100
Floating-Point Operations .....	101
Precision.....	101
Representations .....	101
Rounding Modes .....	102
Operand Sizes .....	103
FABS – Absolute Value.....	104
FADD –Float Addition .....	105
FCMP - Comparison .....	106
FDIV –Float Division .....	108
FEQ – Float Set if Equal .....	109
FNE – Float Set if Not Equal .....	109
FLE – Float Set if Less Than or Equal.....	113
FLT – Float Set if Less Than .....	114
FMUL –Float Multiplication .....	115
FNEG – Negate Value .....	116
FSCALEB –Scale Exponent .....	117
FSUB –Float Subtraction.....	118
FTRUNC – Truncate Fraction .....	119
ORF – Bitwise Or to Float .....	120
String Operations .....	121
Representations .....	121
CHRNDX – Character Index .....	122
Bit Manipulation Operations.....	129
CLR – Clear Bit Field .....	130
COM – Complement Bit Field.....	131
DEP – Deposit Bit Field.....	132
EXTS – Extract Signed Bit Field.....	133
EXTU – Extract Bit Field .....	134
SBX – Sign Bit Extend .....	135
SET – Set Bit Field .....	136
Shift and Rotate Operations .....	137
ASL – Arithmetic Shift Left .....	138
ASR – Arithmetic Shift Right.....	139

LSL – Logical Shift Left.....	140
LSLAND – Logical Shift Left and And.....	141
LSLOR – Logical Shift Left and Or .....	142
LSLXOR – Logical Shift Left and Exclusive Or.....	143
LSR – Logical Shift Right .....	144
ROL – Rotate Left .....	145
ROR – Rotate Right .....	146
Flow Control Instructions .....	147
The Branch Set.....	147
Bcc – Conditional Branch.....	148
BRA – Unconditional Branch.....	170
BRK – Breakpoint.....	172
BSR – Branch to Subroutine.....	173
DBcc – Decrement and Branch.....	174
FBcc – Conditional Branch.....	175
JMP – Jump to Address .....	182
JSR – Jump to Subroutine.....	184
NOP – No Operation.....	189
RTD – Return from Subroutine, Deallocate .....	190
RTS – Return from Subroutine .....	191
RTE – Return From Exception .....	192
TRAP – Trap.....	193
Memory Operations .....	194
AMADD - Addition.....	195
AMAND – Bitwise And .....	196
AMASL – Arithmetic Shift Left.....	197
AMEOR – Bitwise Exclusive Or .....	198
AMLSR – Logical Shift Right.....	199
AMMIN - Minimum .....	200
AMMINU - Minimum .....	200
AMOR – Bitwise Or .....	201
CACHE <cmd>,<ea>.....	202
CMPXCHG – Compare and Exchange.....	203
FLOAD Rn,<ea>.....	204
FSTORE Ra,<ea> .....	205

LA Ra,<ea>.....	206
LOAD Rn,<ea>.....	207
LOADG Gn,<ea>.....	208
LOADZ Rn,<ea> .....	209
STORE Ra,<ea> .....	210
STOREPTR Ra,<ea> .....	212
STOREG Gt,<ea> .....	213
STORE_PAIR Rb, Rc, d[Ra].....	215
Vector Specific Instructions.....	216
V2BITS .....	216
Cryptographic Accelerator Instructions .....	217
AES64DS – Final Round Decryption .....	217
AES64DSM – Middle Round Decryption .....	217
AES64ES – Final Round Encryption.....	218
AES64ESM – Middle Round Encryption.....	218
SHA256SIG0 .....	219
SHA256SIG1 .....	219
SHA256SUM0.....	220
SHA256SUM1 .....	220
SHA512SIG0 .....	221
SHA512SIG1 .....	221
SHA512SUM0.....	222
SHA512SUM1 .....	222
SM3P0.....	222
SM3P1.....	223
SM4ED .....	224
SM4KS.....	224
Modifiers.....	225
ATOM.....	225
CARRY .....	227
VMASK .....	229
PRED .....	230
ROUND .....	232
MPU Hardware .....	233
PIC – Programmable Interrupt Controller .....	233



---

Overview .....	233
System Usage .....	233
Priority Resolution .....	233
Config Space .....	233
Registers .....	234
Control Register .....	235
PIT – Programmable Interval Timer .....	236
Overview .....	236
System Usage .....	236
Config Space .....	236
Parameters .....	237
Registers .....	238
Programming .....	240
Interrupts .....	241

# Thor2023

## Preface

## Who This Book is For

This book describes the Thor2023 ISA. It is for anyone interested in instruction set architectures.

## Motivation

The author desired a CPU core supporting 128-bit floating-point operations for the precision. He also wanted a core he could develop himself. The simplest approach to supporting 128-bit floats is to use 128-bit wide registers, which leads to 128-bit wide busses in the CPU and just generally a 128-bit design. It is not the author's goal to develop a 128-bit machine. There are good ways of obtaining 128-bit floating-point precision on 64-bit or even 32-bit machines, but it adds some complexity. Complexity is something the author must manage to get the project done and a flat 128-bit design would be simpler. Efficiency is being traded off for design simplicity. Some of the most efficient designs are 32-bit.

The processor presented here isn't the smallest, most efficient, and fastest RISC processor. It's also not a simple beginner's example. Those weren't my goals. Instead, it offers reasonable performance with an easy-to-understand state machine and hopefully design simplicity. It's also designed around the idea of using a simple compiler. Some operations like multiply and divide could have been left out and supported with software generated by a compiler rather than having hardware support. But I was after a simple compiler design. There's lots of room for expansion in the future. I chose a 64 bit design supporting 128-bit ops in part anticipating more than 4GB of memory available sometime down the road. A 64-bit architecture is doable in FPGA's today, although it uses double or more the resources that a 32-bit design would.

## About the Author

First a warning: I'm an enthusiastic hobbyist like yourself, with a ton of experience. I've spent a lot of time at home doing research and implementing several soft-core processors, almost maniacally. One of the first cores I worked on was a 6502 emulation. I then went on to develop the Butterfly32 core. Later the Raptor64. I have about 25 years professional experience working on banking applications at a variety of language levels including assembler. So, I have some real-world experience developing complex applications. I also have a diploma in electronics engineering technology. Some of the cores I work on these days are too complex and too large to do at home on an inexpensive FPGA. I await bigger, better, faster boards yet to come. To some extent larger boards have arrived. The author is a bit wary of larger boards. Larger FPGAs increase build times by their nature.

# Nomenclature

There has been some mix-up in the naming of load and store instructions as computer systems have evolved. A while ago, a “word” referred to a 16-bit quantity. This is reflected in the mnemonics of instructions where move instructions are qualified with a “.w” for a 16-bit move. Some machines referred to 32-bits as a word. Times have changed and 64-bit workstations are now more common. In the author’s parlance a word refers to the word size of a machine, which may be 16, 32, 64 bits or some other size. What does “.w” or “.d”, and “.l” refer to? To some extent it depends on the architecture.

The ISA refers to primitive object sizes following the convention suggested by Knuth of using Greek.

8	byte	LDB, STB	UTF8 usage
16	wyde	LDW, STW	
32	tetra	LDT, STT	
64	octa	LDO, STO	
128	hexi	LDH, STH	

The register used to address instructions is referred to as the instruction pointer or IP register. The instruction pointer is a synonym for program counter or PC register.

# Little Endian vs big Endian

One choice to make is whether the architecture is little endian or big endian. There's a never-ending argument by computer folks as to which endian is better. In reality they are about the same or there wouldn't be an argument. In a little-endian architecture, the least significant byte is stored at the lowest memory address. In a big-endian architecture the most significant byte is stored at the lowest memory address. The author is partial to little endian machines; it just seems more natural to him although he knows people who swear by the opposite. Whichever endian is chosen, often the machine has instructions(s) for converting from one endian to the other. The author does not bother with endian conversion; it's a feature that he probably wouldn't use. Some implementations even allow the endian of the machine to be set by the user. This seems like overkill to the author. The endian of data is important because some file types depend on data being in little or big-endian format. Thor is a little-endian machine.

## Endian

Thor2023 is a little-endian machine. The difference between big endian and little endian is in the ordering of bytes in memory. Bits are also numbered from lowest to highest for little endian and from highest to lowest for big endian.

Shown is an example of a 32-bit word in memory.

Little Endian:

Address	3	2	1	0
Byte	3	2	1	0

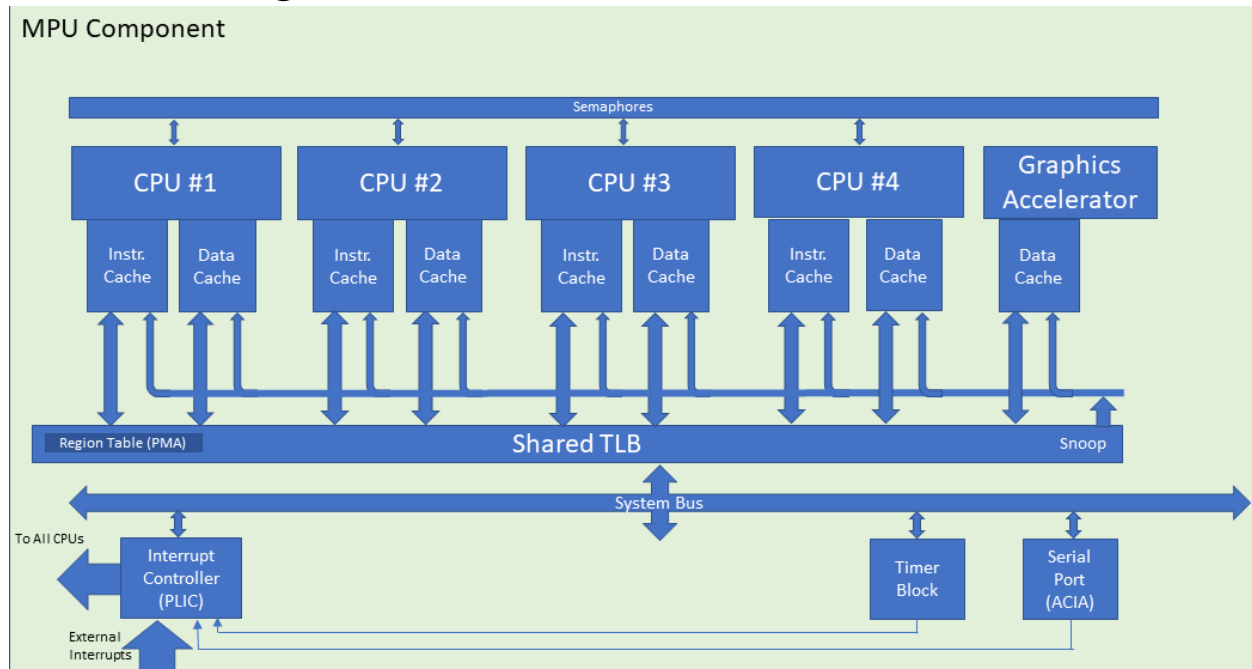
Big Endian:

Address	3	2	1	0
Byte	0	1	2	3

For Thor2023 the root opcode is in byte zero of the instruction and bytes are shown from right to left in increasing order. As the following table shows.

Address 3	Address 2	Address 1	Address 0
Byte 3	Byte 2	Byte 1	Byte 0
			▼
31 24	23 16	15 8	7 5 4 0
Constant <sub>8</sub>	Raspec <sub>8</sub>	Rtspec <sub>8</sub>	Sz <sub>3</sub> Opcode <sub>5</sub>

# Block Diagram



# Programming Model

## Register File

### Rn – General Purpose Registers

The register file contains 64 128-bit general purpose registers.

The register file is *unified* and may hold either integer or floating-point values. The stack pointer, register 63, is banked with a separate stack pointer for each operation mode. Registers may be loaded or stored individually or in groups of eight 64-bit values.

Register r53 is special in that when read it refers to the program counter's current value, used to form PC relative addresses. When written it refers to the stack canary register. Attempting to load the register from memory causes a stack canary check instead. The loaded value will be compared against the canary and an exception will occur if they differ.

Register r0 is special in that it always reads as a zero. Note that it may be inverted to read as -1 for some instructions.

## Register ABI

Regno	ABI	Group Reg	ABI Usage
0	0	AG0	Always zero
1	A0		First argument / return value register
2	A1		Second argument / return value register
3	A2		Third argument register
4 to 7	A3 to A6		Argument registers
8 to 15	T0 to T7	TG0	Temporary register, caller save
16 to 31	S0 to S15	SG0, SG1	Saved register, register variables

32 to 39	VM0 to VM7	VMG	Vector mask
40 to 43	A7 to A10	G5	Argument register
44 to 47	T8 to T11		Temporaries
48 to 51			unassigned
52	TS	G6	Thread state pointer
53	PC / SC		Program counter; LOAD does canary check
54	CTA		Card table address
55	LC		Loop counter
56	LR0		Subroutine link register #0; branch subroutine specific
57	LR1	G7	Subroutine link register #1; milli-code routines
58	LR2		Subroutine link register #2
59	LR3		Subroutine link register #3
60	GP1		Global Pointer #1 (RO data segment)
61	GP0		Global Pointer #0 (Data segment)
62	FP		Frame Pointer
63	SP		Stack Pointer
63	ASP		Application/User Stack pointer
63	SSP		Supervisor Stack pointer
63	HSP		Hypervisor Stack pointer
63	MSP		Machine stack pointer

## Predicate Registers

The original Thor machine had 16 four-bit dedicated predicate registers. Thor2023 by contrast stores predicate conditions in general purpose registers. Any GPR may be used to hold values used in predication. Original Thor predicates were a prefix byte containing the predicate register and condition present for every instruction. This has been superseded using the predicate instruction modifier, [PRED](#), which allows up to eight following instructions to be predicated in the same manner. The PRED modifier is more storage efficient than predicating every instruction with predicate bits as most instructions do not require predication.

## Mask Registers (vm0 to vm7)

Mask registers are used to mask off vector operations so that a vector instruction doesn't perform the operation on all elements of the vector. Vector instructions (loads and stores) that don't explicitly specify a mask register assume the use of mask register zero (vm0). Mask registers are a subset of the general-purpose register array, allowing instructions that operate on GPRs to operate on the mask registers. Potentially any register could be used as a mask register, the compiler will assign a register as needed. Vm0 to vm7 are just a suggestion of registers to reserve for vector masking.

Mask register specification allows the mask register to be used in an inverted form. This can be applied to r0 which will then enable all lanes of execution.

*Thor2022 had dedicated mask registers leading to additional instructions required to manipulate them.*

vm0	32	
vm1	33	
vm2	34	
vm3	35	
vm4	36	
vm5	37	
vm6	38	
vm7	39	

## Vector Length (VL register)

The vector length register controls how many elements of a vector are processed. The vector length register may not be set to a value greater than the number of elements supported by hardware. After the vector length is set a SYNC instruction should be used to ensure that following instructions will see the updated version of the length register.

Vector length has register tag #87.

15                      8      7                      0



0	Elements <sub>7..0</sub>
---	--------------------------

## Code Address Registers

Many architectures have registers dedicated to addressing code. Almost every modern architecture has a program counter or instruction pointer register to identify the location of instructions. Many architectures also have at least one link register or return address register holding the address of the next instruction after a subroutine call. There are also dedicated branch address registers in some architectures. These are all code addressing registers.

*The original Thor lumped these registers together in a code address register array. For Thor2023 some of these registers are now part of the general register file.*

It is possible to do an indirect method call using any register.

### LRn – Link Registers

There are four registers in the Thor2023 architecture reserved for subroutine linkage. These registers are used to store the address after the calling instruction. They may be used to implement fast returns for several levels of subroutines or to used to call milli-code routines. The jump to subroutine, [JSR](#), and branch to subroutine, [BSR](#), instructions update a link register. The return from subroutine, [RTS](#), instruction is used to return to the next instruction.

### PC – Program Counter

This register points to the currently executing instruction. The program counter increments as instructions are fetched, unless overridden by another flow control instruction. The program counter may be set to any byte address. There is no alignment restriction. It is possible to write position independent code, PIC, using PC relative addressing.

## LC - Loop Counter (reg 55)

The loop counter register is used in counted loops along the decrement and branch, [DBcc](#), instruction.

## SR - Status Register (CSR 0x?004)

The processor status register holds bits controlling the overall operation of the processor, state that needs to be saved and restored across interrupts. The bits have individual bit set / clear capability using the CSRRS, CSRRC instructions. Only the user interrupt enable bit is available in user mode, other bits will read as zero.

Bit		Usage
0	uie	User interrupt enable
1	sie	Supervisor interrupt enable
2	hie	Hypervisor interrupt enable
3	mie	Machine interrupt enable
4	die	Debug interrupt enable
5 to 7	ipl	Interrupt level
8	ssm	Single step mode
9	te	Trace enable
10 to 11	om	Operating mode
12 to 13	ps	Pointer size
14 to 15	~	reserved
16	mprv	memory privilege
17	~	reserved
18	dmi	<del>Decimal mode for integers</del>
19	dmf	<del>Decimal mode for float</del>
20 to 23	~	reserved
24 to 31	cpl	Current privilege level

CPL is the current privilege level the processor is operating at.

T indicates that trace mode is active.

OM processor operating mode.

PS: indicates the size of pointers in use. This may be one of 32, 64 or 128 bits.

AR: Address Range indicates the number of address bits in use. 0 = near or short (32-bit) addressing is in use. When short addressing is in use only the low order 32-bit are significant and stored or loaded to or from the stack.

IPL is the interrupt mask level

RT specifies the return type for an [RTI](#) instruction.

MPRV Memory Privilege, indicates to use previous operating mode for memory privileges

## Decimal Mode

~~Setting the 'D' flag bit 5 in the SR register sets the processor in decimal operating mode. Arithmetic operations will use BCD numbers for both source and destination operands.~~

~~Decimal mode, 'D' flag bit 4, may also be applied to floating point which will use decimal floating point operations instead of binary.~~

Decimal mode is now handled on an instruction-by-instruction basis with bits in the instruction indicating when decimal mode is in use.

# Vector Programming Model

## Register File

### Vn – SIMD Registers

The SIMD register file contains 32 512-bit registers.

Regno	ABI	ABI Usage
0		
1	VA0	First argument / return value
2	VA1	Second argument / return value
3	VA2	Third argument
4 to 15	VT0 to VT11	
16 to 27	VS0 to VS11	
28 to 31	VA3 to VA6	

### Vector Related CSRs

VGM		Global mask register
VRM		Restart mask register
VERR		Error mask register
VRGSZ		Vector register size
VED		Vector element descriptor

The number of elements is limited to 128 as that is the width of a predicate register.

### Vector Global Mask Register (VGM)

The global mask register contains predicate bits indicating which vector elements are active. Vector elements of the target are updated only when the corresponding global mask bit is set. The global mask register takes the place of the vector length register in other architectures. Normally the global mask contains a right aligned bitmask of all ones up to the number of elements to be processed.

### Vector Restart Mask Register (VRM)

The restart mask register contains a bitmask indicating the vectors elements to be processed after a restart. The restart mask register is set to all ones at the end of a vector operation.

### Vector Error Mask Register (VERR)

The vector error mask register contains a bit for each vector element indicating if an error occurred.

## Vector Register Size (VRGSZ)

The vector register size register contains the length of a vector register in bytes. Only the low order eight bits of the register are implemented, other bits read as zero, and ignore writes.

## Vector Element Description Register (VED)

This register contains bits describing an element of a vector.

127	6	5	3	2	0
	~			OT <sub>3</sub>	Size <sub>3</sub>

Size <sub>3</sub>	Bits	Bytes
0	8	1
1	16	2
2	32	4
3	64	8
4	128	16
5	256	32
6	512	64
7		reserved

OT <sub>3</sub>	Operand Type
0	Integer
1	Float
2	Decimal
3	Posit
4	Char
5 to 7	reserved

## Special Purpose Registers

### SC - Stack Canary (GPR 53)

This special purpose register is available in the general register file as register 53. The stack canary register is used to alleviate issues resulting from buffer overflows on the stack. The canary register contains a random value which remains consistent throughout the run-time of a program. In the right conditions, the canary register is written to the stack during the function's prolog code. In the function's epilog code, the value of the canary on stack is checked to ensure it is correct, if not a check exception occurs.

### [U/S/H/M]\_IE (0x?004)

See status register.

This register contains interrupt enable bits. The register is present at all operating levels. Only enable bits at the current operating level or lower are visible and may be set or cleared. Other bits will read as zero and ignore writes. Only the lower four bits of this register are implemented. The bits have individual bit set / clear capability using the CSRRS, CSRRC instructions.

63																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																					
----	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

### [U/S/H/M]\_CAUSE (CSR- 0x?006)

This register contains a code indicating the cause of an exception or interrupt. The break handler will examine this code to determine what to do. Only the low order 12 bits are implemented. The high order bits read as zero and are not updateable.

### U\_REPBUF - (CSR – 0x008)

This register contains information needed for the REP instruction that must be saved and restored during context switches and interrupts. Note that the loop counter should also be saved.

127	112	121		48	47	44	43	42	40	39		8	7	6	0
Resv	pc				Resv2	V	ICnt	Limit				resv	Ins[15:9]		

Pc: (64 bits) the address of the instruction following the REP

V: REP valid bit, 1 only if a REP instruction is active

ICnt: the current instruction count, distance from REP instruction.

Limit: a 32-bit amount to compare the loop counter against.

Ins: bits 9 to 15 of the REP instruction which contains the instruction count of instruction included in the repeat and condition under which the repeat occurs.

### [U/S/H/M]\_SCRATCH – CSR 0x?041

This is a scratchpad register. Useful when processing exceptions. There is a separate scratch register for each operating mode.

### S\_PTBR (CSR 0x1003)

This register contains the base address of the page table, which must be a multiple of 16384. Also included in this register is table parameters depth and type. Register tag #152.

95	14	13	12	11	8	7	6	5	4	3	2	1	0
Page Table Address <sub>67..14</sub>				~ <sub>2</sub>	Levels	AL <sub>2</sub>	~ <sub>2</sub>	S	~	Type			

Type: 0 = inverted page table, 1 = page table

S: 1=software managed TLB miss, 0 = hardware table walking

Levels are ignored for the inverted page table. For a normal page table gives the top entry level.

AL<sub>2</sub>: TLB entry replacement algorithm, 0=fixed,1=LRU,2=random,3=reserved

### S\_ASID (CSR 0x101F)

This register contains the address space identifier (ASID) or memory map index (MMI). The ASID is used in this design to select (index into) a memory map in the paging tables. Only the low order twelve bits of the register are implemented.

### S\_KEYS (CSR 0x1020 to 0x1027)

These eight registers contain the collection of keys associated with the process for the memory lot system. Each key is twenty-four bits in size. All eight registers are searched in parallel for keys matching the one associated with the memory page. Keyed memory enhances the security and reliability of the system.

		23	0
1020			key0
1021			key1
...			...
1027			key7

**M\_CORENO (CSR 0x3001)**

This register contains a number that is externally supplied on the coreno\_i input bus to represent the hardware thread id or the core number. It should be non-zero.

**M\_TICK (CSR 0x3002)**

This register contains a tick count of the number of clock cycles that have passed since the last reset. Note that this register should not be used for precise timing as the processor's clock frequency may vary for performance and power reasons. The TIME CSR may be used for wall-clock timing as it has its own timing source.

**M\_SEED (CSR 0x3003)**

This register contains a random seed value based on an external entropy collector. The most significant bit of the state is a busy bit.

63	60	59		16	15	0
State <sub>4</sub>			~44	seed <sub>16</sub>		

State<sub>4</sub>  
Bit

0	dead
1	test
2	valid, the seed value is valid
3	Busy, the collector is busy collecting a new seed value

**M\_BADADDR (CSR 0x3007)**

This register contains the address for a load / store operation that caused a memory management exception or a bus error. Note that the address of the instruction causing the exception is available in the EPC register.

**M\_BAD\_INSTR (CSR 0x300B)**

This register contains a copy of the exceptioned instruction.

**M\_SEMA (CSR 0x300C)**

This register contains semaphores. The semaphores are shared between all cores in the MPU.

**M\_TVEC – CSR 0x3030 to 0x3034**

These registers contain the address of the exception handling routine for a given operating level. TVEC[4] (0x3034) is used directly by hardware to form an address of the debug routine. The lower eight bits of TVEC[3] are not used. The lower bits of the exception address are determined from the operating level. TVEC[0] to TVEC[2] are used by the REX instruction.

A sync instruction should be used after modifying one of these registers to ensure the update is valid before continuing program execution.

Reg #

0x3030	TVEC[0] – user mode
0x3031	TVEC[1] - supervisor mode
0x3032	TVEC[2] – hypervisor mode
0x3033	TVEC[3] – machine mode
0x3034	TVEC[4] - debug

### M\_SR\_STACK (CSR 0x303C to CSR 0x303D)

This pair of registers contains a stack of the status register which is pushed during exception processing and popped on return from interrupt. There are only eight slots as that is the maximum nesting depth for interrupts.

	127	96	95	64	63	32	31	0
0x303C	SR3		SR2		SR1		SR0	
0x303D	SR7		SR6		SR5		SR4	

### M\_IOS – IO Select Register (CSR 0x3100)

The location of IO is determined by the contents of the IOS control register. The select is for a 1MB region. This address is a virtual address. The low order 16 bits of this register should be zero and are ignored.

63	16	15	0
Virtual Address <sub>67..20</sub>		0 <sub>16</sub>	

### M\_EPC (CSR 0x3108 to 0x310F)

This set of registers contains the address stack for the program counter used in exception handling.

Reg #	Name
0x3108	EIP0
...	
0x310F	EIP7

### AV – Application Vector Table Address

This register holds the address of the applications vector table. The vector table must be 16-byte aligned.

63	0
App Vector Table Address <sub>67..4</sub>	

### VB – Vector Base Register

The vector base register provides the location of the vector table. The vector table must be octa aligned. On reset the VBR is loaded with zero. There is a separate vector base register for each operating mode.



63	3	2	1 0
Vector Table Address <sub>63..3</sub>		~	~

## Operating Modes

The core operates in one of four basic modes: application/user mode, supervisor mode, hypervisor mode or machine mode. Machine mode is switched to when an interrupt or exception occurs, or when debugging is triggered. On power-up the core is running in machine mode. An RTI instruction must be executed to leave machine mode after power-up.

A subset of instructions is limited to machine mode.

Mode Bits	Mode
0	User / App
1	Supervisor
2	Hypervisor
3	Machine

# Exceptions

## External Interrupts

There is little difference between an externally generated exception and an internally generated one. An externally caused exception will set the exception cause code for the currently fetched instruction.

There are eight priority interrupt levels for external interrupts. When an external interrupt occurs the mask level is set to the level of the current interrupt. A subsequent interrupt must exceed the mask level to be recognized.

## Effect on Machine Status

The operating mode is always switched to machine mode on exception. It is up to the machine mode code to redirect the exception to a lower operating mode when desired. Further exceptions at the same or lower interrupt level are disabled automatically. Machine mode code must enable interrupts at some point.

## Exception Stack

The status register and program counter are pushed onto an internal stack when an exception occurs. This stack is at least 16 entries deep to allow for nested interrupts and multiply nested traps and exceptions.

## Exception Table

Vector	Usage
0	Reset value for system stack pointer
1	Reset value for program counter
2	Bus Error
3	Address Error
4	Unimplemented Instruction
5	
6	
7	
8	Privilege Violation
9	Instruction trace
10	
11	Stack Canary
12 to 23	reserved
24	Spurious interrupt
25	Auto vector #1
26	Auto vector #2
27	Auto vector #3
28	Auto vector #4
29	Auto vector #5
30	Auto vector #6

31	Auto vector #7
32	Breakpoint (BRK)
33	Debug
34	Instruction Address
33 to 63	Trap #1 to 31
	Applications Usage
64	Divide by zero
65	Overflow
66	Table Limit
67 to 511	Unassigned usage

### Reset Stack Pointer Vector (0)

This vector contains the address the machine stack pointer is set to at reset.

### Reset Vector Vector (1)

This vector contains the address that the processor begins running at.

### Bus Error Fault (2)

The bus error fault is performed if the bus error signal was active during the bus transaction. This could be due to a bad or missing device.

### Unimplemented Instruction Fault (4)

An unimplemented instruction causes this fault.

### Stack Canary Fault (11)

This fault is caused if the stack canary was overwritten. A load instruction using the canary register did not match the value in the canary register.

### Breakpoint Fault (32)

The breakpoint instruction, 0, was encountered.

### Instruction Address Fault (34)

An error occurred addressing instructions. This could be due to a bad instruction sequence, for instance executing multiple postfixes in a row.

## Reset

Reset is treated as an exception. The reset routine should exit using an RTI instruction. The status register should be setup appropriately for the return.

The core begins executing instructions at address \$00...00. All registers are in an undefined state.

## Precision

Exceptions in Thor2023 are precise. They are processed according to program order of the instructions. If an exception occurs during the execution of an instruction, then an exception field is set in the pipeline buffer. The exception is processed when the instruction commits which happens in program order. If the instruction was executed in a speculative fashion, then no exception processing will be invoked unless the instruction makes it to the commit stage.

# Memory Management

## Bank Swapping

About the simplest form of memory management is a single bank register that selects the active memory bank. This is the mechanism used on many early microcomputers. The bank register may be an eight bit I/O port supplying control over some number of upper address bits used to access memory.

## The Page Map

The next simplest form of memory management is a single table map of virtual to physical addresses. The page map is often located in a high-speed dedicated memory. An example of a mapping table is the 74LS612 chip. It may map four address bits on the input side to twelve address bits on the output side. This allows a physical address range eight bits greater than the virtual address range. A more complicated page map is something like the MC6829 MMU. It may map 2kB pages in a 2MB physical address space for up to four different tasks.

## Regions

In any processing system there are typically several different types of storage assigned to different physical address ranges. These include memory mapped I/O, MMIO, DRAM, ROM, configuration space, and possibly others. Thor2023 has a region table that supports up to eight separate regions.

The region table is a list of region entries. Each entry has a start address, an end address, an access type field, and a pointer to the PMT, page management table. To determine legal access types, the physical address is searched for in the region table, and the corresponding access type returned. The search takes place in parallel for all eight regions.

Once the region is identified the access rights for a particular page within the region can be found from the PMT corresponding to the region. Global access rights for the entire region are also specified in the region table. These rights are gated with value from the PMT and TLB to determine the final access rights.

# PMA - Physical Memory Attributes Checker

## Overview

The physical memory attributes checker is a hardware module that ensures that memory is being accessed correctly according to its physical attributes.

Physical memory attributes are stored in an eight-entry region table. Three bits in the PTE select an entry from this table. The operating mode of the CPU also determines which 32-bit set of attributes to apply for the memory region.

Most of the entries in the table are hard-coded and configured when the system is built. However, they may be modified at the address range \$F...F9F0xxx.

Physical memory attributes checking is applied in all operating modes.

The region table is accessible as a memory mapped IO, MMIO, device.

## Region Table Description

Reg	Bits		
00	128	Pmt	associated PMT address
01	128	cta	Card table address
02	128	at	Four groups of 32-bit memory attributes, 1 group for each of user, supervisor, hypervisor and machine.
03	128	...	Not used
04 to 1F		...	7 more register sets

### PMT Address

The PMT address specifies the location of the associated PMT.

### CTA – Card Table Address

The card table address is used during the execution of the store pointer, STPTR instruction to locate the card table.

### Attributes

Bitno			
0	X	may contain executable code	
1	W	may be written to	
2	R	may be read	
3	~	reserved	
4-7	C	Cache-ability bits	
8-10	G	granularity	
		G	
		0	byte accessible
		1	wyde accessible
		2	tetra accessible
		3	octa accessible
		4	hexi accessible
		5 to 7	reserved
11	~	reserved	
12-14	S	number of times to shift address to right and store for telescopic STPTR stores.	
16-23	T	device type (rom, dram, eeprom, I/O, etc)	
24-31	~	reserved	



# Page Management Table - PMT

## Overview

For the first translation of a virtual to physical address, after the physical page number is retrieved from the TLB, the region is determined, and the page management table is referenced to obtain the access rights to the page. PMT information is loaded into the TLB entry for the page translation. The PMT contains an assortment of information most of which is managed by software. Pieces of information include the key needed to access the page, the privilege level, and read-write-execute permissions for the page. The table is organized as rows of access rights table entries (PMTEs). There are as many PMTEs as there are pages of memory in the region.

For subsequent virtual to physical address translations PMT information is retrieved from the TLB.

As the page is accessed in the TLB, the TLB may update the PMT.

## Location

The page management table is in main memory and may be accessed with ordinary load and store instructions. The PMT address is specified by the region table.

## PMTE Description

There is a wide assortment of information that goes in the page management table. To accommodate all the information an entry size of 128-bits was chosen.

### Page Management Table Entry

V	N	M	~9	C	E	AL <sub>2</sub>	~16
ACL <sub>16</sub>							Share Count <sub>16</sub>
							Access Count <sub>32</sub>
PL <sub>8</sub>				Key <sub>24</sub>			

## Access Control List

The ACL field is a reference to an associated access control list.

## Share Count

The share count is the number of times the page has been shared to processes. A share count of zero means the page is free.

## Access Count

This part uses the term ‘access count’ to refer to the number of times a page is accessed. This is usually called the reference count, but that phrase is confusing because reference counting may also refer to share counts. So, the phrase ‘reference count’ is avoided. Some texts use the term



reference count to refer to the share count. Reference counting is used in many places in software and refers to the number of times something is referenced.

Every time the page of memory is accessed, the access count of the page is incremented. Periodically the access count is aged by shifting it to the right one bit.

The access count may be used by software to help manage the presence of pages of memory.

## Key

The access key is a 24-bit value associated with the page and present in the key ring of processes. The keyset is maintained in the keys CSRs. The key size of 20 bits is a minimum size recommended for security purposes. To obtain access to the page it is necessary for the process to have a matching key OR if the key to match is set to zero in the PMTE then a key is not needed to access the page.

## Privilege Level

The current privilege level is compared with the privilege level of the page, and if access is not appropriate then a privilege violation occurs. For data access, the current privilege level must be at least equal to the privilege level of the page. If the page privilege level is zero anybody can access the page.

## N

indicates a conforming page of executable code. Conforming pages may execute at the current privilege level. In which case the PL field is ignored.

## M

indicates if the page was modified, written to, since the last time the M bit was cleared. Hardware sets this bit during a write cycle.

## E

indicates if the page is encrypted.

## AL

indicates the compression algorithm used.

## C

The C indicator bit indicates if the page is compressed.

# Page Tables


## Intro

Page tables are part of the memory management system used map virtual addresses to real physical addresses. There are several types of page tables. Hierarchical page tables are probably the most common. Almost all page tables map only the upper bits of a virtual address, called a page. The lower bits of the virtual address are passed through without being altered. The page size often 4kB which means the low order 12-bits of a virtual address will be mapped to the same 12-bits for the physical address.

## Hierarchical Page Tables

Hierarchical page tables organize page tables in a multi-level hierarchy. They can map the entire virtual address range but often only a subrange of the full virtual address space is mapped. This can be determined on an application basis. At the topmost level a register points to a page directory, that page directory points to a page directory at a lower level until finally a page directory points to a page containing page table entries. To map an entire 64-bit virtual address range approximately five levels of tables are required.

## Paged MMU Mapping



## Inverted Page Tables

An inverted page table is a table used to store address translations for memory management. The idea behind an inverted page table is that there are a fixed number of pages of memory no matter how it is mapped. It should not be necessary to provide for a map of every possible address, which is what the hierarchical table does, only addresses that correspond to real pages of memory need be mapped. Each page of memory can be allocated only once. It is either allocated or it is


not. Compared to a non-inverted paged memory management system where tables are used to map potentially the entire address space an inverted page table uses less memory. There is typically only a single inverted page table supporting all applications in the system. This is a different approach than a non-inverted page table which may provide separate page tables for each process.

## The Simple Inverted Page Table

The simplest inverted page table contains only a record of the virtual address mapped to the page, and the index into the table is used as the physical page number. There are only as many entries in the inverted page table as there are physical pages of memory. A translation can be made by scanning the table for a matching virtual address, then reading off the value of the table index. The attraction of an inverted page table is its small size compared to the typical hierarchical page table. Unfortunately, the simplest inverted page table is not practical when there are thousands or millions of pages of memory. It simply takes too long to scan the table. The alternative solution to scanning the table is to hash the virtual address to get a table index directly.

## Inverted Page Table

Entry number identifies physical page number



## Hashed Page Tables

### Hashed Table Access

Hashes are great for providing an index value immediately. The issue with hash functions is that they are just a hash. It is possible that two different virtual address will hash to the same value. What is then needed is a way to deal with these hash collisions. There are a couple of different methods of dealing with collisions. One is to use a chain of links. The chain has each link in the chain pointing the to next page table entry to use in the event of a collision. The hash page table is slightly more complicated then as it needs to store links for hash chains. The second method is to use open addressing. Open addressing calculates the next page table entry to use in the event of a collision. The calculation may be linear, quadratic or some other function dreamed up. A linear probe simply chooses the next page table entry in succession from the previous one if no match occurred. Quadratic probing calculates the next page table entry to use based on squaring the count of misses.

## Clustered Hash Tables

A clustered hash table works in the same manner as a hashed page table except that the hash is used to access a cluster of entries rather than a single entry. Hashed values may map to the same cluster which can store multiple translations. Once the cluster is identified, all the entries are searched in parallel for the correct one. A clustered hash table may be faster than a simple hash table as it makes use of parallel searches. Often accessing memory returns a cache line regardless of whether a single byte or the whole cached line is referenced. By using a cache line to store a cluster of entries it can turn what might be multiple memory accesses into a single access. For example, an ordinary hash table with open addressing may take up to 10 memory accesses to find the correct translation. With a clustered table that turns into 1.25 memory accesses on average.

## Shared Memory

Another memory management issue to deal with is shared memory. Sometimes applications share memory with other apps for communication purposes, and to conserve memory space where there are common elements. The same shared library may be used by many apps running in the system. With a hierarchical paged memory management system, it is easy to share memory, just modify the page table entry to point to the same physical memory as is used by another process. With an inverted page table having only a single entry for each physical page is not sufficient to support shared memory. There needs to be multiple page table entries available for some physical pages but not others because multiple virtual addresses might map to the same physical address. One solution would be to have multiple buckets to store virtual addresses in for each physical address. However, this would waste a lot of memory because much of the time only a single mapped address is needed. There must be a better solution. Rather than reading off the table index as the physical page number, the association of the virtual and physical address can be stored. Since we now need to record the physical address multiple times the simple mechanism of using the table index as the physical page number cannot be used. Instead, the physical page number needs to be stored in the table in addition to the virtual page number.

That means a table larger than the minimum is required. A minimally sized table would contain only one entry for each physical page of memory. So, to allow for shared memory the size of the table is doubled. This smells like a system configuration parameter.

## Specifics: Thor2023 Page Tables

### Thor2023 Hash Page Table Setup

#### Hash Page Table Entries - HPTE

We have determined that a page table entry needs to store both the physical page number and the virtual page number for the translations. To keep things simple, the page table stores only the information needed to perform an address translation. Other bits of information are stored in a secondary table called the page management table, PMT. The author did a significant amount of juggling around the sizes of various fields, mainly the size of the physical and virtual page numbers. Finally, the author decided on a 192-bit HPTE format.

V	LVL/BC <sub>5</sub>	RGN <sub>3</sub>	M	A	T	S	G	SW <sub>2</sub>	CACHE <sub>4</sub>	MRWX <sub>3</sub>	HRWX <sub>3</sub>	SRWX <sub>3</sub>	URWX <sub>3</sub>
PPN <sub>31..0</sub>													
PPN <sub>63..32</sub>													
VPN <sub>37..6</sub>													
VPN <sub>69..38</sub>													
~4		ASID <sub>11..0</sub>						~2		VPN <sub>83..70</sub>			

### Fields Description

V	1	translation Valid
G	1	global translation
RGN	3	region
PPN	64	Physical page number
VPN	84	Virtual page number
RWX	3	readable, writeable, executable
ASID	12	address space identifier
LVL/BC	5	bounce count
M	1	modified
A	1	accessed
T	1	PTE type (not used)
S	1	Shared page indicator
SW	3	OS usage

The page table does not include everything needed to manage pages of memory. There is additional information such as share counts and privilege levels to take care of, but this information is better managed in a separate table.

### Small Hash Page Table Entries - SHPTE

The small HPTE is used for the test system which contains only 512MB of physical RAM to conserve hardware resources. The SHPTE is 96-bits in size.

V	LVL/BC <sub>5</sub>	RGN <sub>3</sub>	M	A	T	S	G	SW <sub>2</sub>	CACHE <sub>4</sub>	MRWX <sub>3</sub>	HRWX <sub>3</sub>	SRWX <sub>3</sub>	URWX <sub>3</sub>
VPN <sub>5..0</sub>		PPN <sub>25..0</sub>											
ASID <sub>11..0</sub>							VPN <sub>25..6</sub>						

### Page Table Groups – PTG

We want the search for translations to be fast. That means being able to search in parallel. So, PTEs are stored in groups that are searched in parallel for translations. This is sometimes referred to as a clustered table approach. Access to the group should be as fast as possible. There are also hardware limits to how many entries can be searched at once while retaining a high clock rate. So, the convenient size of 1024 bits was chosen as the amount of memory to fetch.

A page table group then contains ten SHPTE page-table entries or five HPTE entries. All entries in the group are searched in parallel for a match. Note that the entries are searched as the PTG is loaded, so that the PTG group load may be aborted early if a matching PTE is found before the load is finished.

127	0
	PTE0
	PTE1
	PTE2
	PTE3
	PTE4
	PTE5
	PTE6
	PTE7

## Size of Page Table

There are several conflicting elements to deal with, with regards to the size of the page table. Ideally, the hash page table is small enough to fit into the block RAM resources available in the FPGA. It may be practical to store the hash page table in block RAM as there would be only a single table for all apps in the system. This probably would not be practical for a hierarchical table.

About 1/3 of the block RAMs available are dedicated to MMU use. At the same time a multiple of the number of physical pages of memory should be supported to support page sharing and swapping pages to secondary storage. To support swapping pages, double the number of physical entries were chosen. To support page sharing, double that number again. Therefore, a minimum size of a page table would contain at least four times the number of physical pages for entries. By setting the size of the page table instead of the size of pages, it can be worked backwards how many pages of memory can be supported.

For a system using 512k block RAM to store PTEs.  $512k / 96 = 5,461$  entries.  $5,461 / 4 = 1,365$  physical pages. Since the RAM size is 512MB, each page would be  $512MB / 1,365 = 393kB$ . Rounding up, 512kB. Since half the pages may be in secondary storage, 1GB of address range is available. A 512kB page is probably too large to be useful, so either more block RAM is required, or the table could be place in main memory.

Since there are 5,461 entries in the table and they are grouped into groups of ten, there are 546 PTGs. To get to a page table group fast a hash function is needed then that returns a 10-bit number.

## Hash Function

The hash function needs to reduce the size of a virtual address down to a 11-bit number. The asid should be considered part of the virtual address. Including the asid an address is 76 bits. The first thing to do is to throw away the lowest fourteen bits as they pass through the MMU unaltered. We now have 62-bits to deal with. We can probably throw away some high order bits too, as a process is not likely to use the full 64-bit address range.

The hash function chosen uses the asid combined with virtual address bits 18 to 28 and bits 29 to 39. This should space out the PTEs according to the asid. Address bits 16 and 17 select one of four address ranges. the PTG supports ten PTEs. The translations where address bits 16 and 17

are involved are likely consecutive pages that would show up in the same PTG. The hash is the asid exclusively or'd with address bits 18 to 28 exclusively or'd with address bits 29 to 39.

## Collision Handling

Quadratic probing of the page table is used when a collision occurs. The next PTG to search is calculated as the hash plus the square of the miss count. On the first miss the PTG at the hash plus one is searched. Next the PTG at the hash plus four is searched. After that the PTG at the hash plus nine is searched, and so on.

## Finding a Match

Once the PTG to be searched is located using the hash function, which PTE to use needs to be sorted out. The match operation must include both the virtual address bits and the asid, address space identifier, as part of the test for a match. It is possible that the same virtual address is used by two or more different address spaces, which is why it needs to be in the match.

## Locality of Reference

The page table group may be cached in the system read cache for performance. It is likely that the same PTG group will be used multiple times due to the locality of reference exhibited by running software.

## Access Rights

To avoid duplication of data the access rights are stored in another table called the PMT for access rights table. The first time a translation is loaded the access rights are looked-up from the PMT. A bit is set in the TLB entry indicating that the access rights are valid. On subsequent translations the access rights are not looked up, but instead they are read from values cached in the TLB.

## Thor2023 Hierarchical Page Table Setup

## Page Table Entries - PTE

For hierarchical tables the structure is like that of hashed page tables except that there is no need to store the virtual address. We know the virtual address because it is what is being translated and there is no chance of collisions unlike the hash table. The structure is 96 bits in size. This allows 1024 PTEs to fit into an 16kB page. ¼ of the 16kB page is not used. Note the size of pages in the table is a configuration parameter used to build the system.

There are two types of page table entries. The first type,  $T=0$ , is a pointer to a page of memory, the second type,  $T=1$ , is an entry that points to lower-level page tables. PTE's that point to lower-level page tables are sometimes called page table pointers, PTPs.

## Page Table Entry Format – PTE

V	LVL/BC <sub>5</sub>	RGN <sub>3</sub>	M	A	T	S	G	SW <sub>2</sub>	CACHE <sub>4</sub>	MRWX <sub>3</sub>	HRWX <sub>3</sub>	SRWX <sub>3</sub>	URWX <sub>3</sub>
PPN <sub>31..0</sub>													
PPN <sub>63..32</sub>													

## Small Page Table Entry Format – SPTE

The small PTE format is used when the physical address space is less than 46-bits in size. The small PTE occupies only 64-bits. 2048 SPTEs will fit into an 16kB page.

V	LVL/BC <sub>5</sub>	RGN <sub>3</sub>	M	A	T	S	G	SW <sub>2</sub>	CACHE <sub>4</sub>	MRWX <sub>3</sub>	HRWX <sub>3</sub>	SRWX <sub>3</sub>	URWX <sub>3</sub>
PPN <sub>31..0</sub>													

Field	Size	Purpose
PPN	64	Physical page number
URWX	3	User read-write-execute override
SRWX	3	Supervisor read-write-execute override
HRWX	3	Hypervisor read-write-execute override
MRWX	3	Machine read-write-execute override
CACHE	4	Cache-ability bits
A	1	1=accessed/used
M	1	1=modified
V	1	1 if entry is valid, otherwise 0
S	1	1=shared page
G	1	1=global, ignore ASID
T	1	0=page pointer, 1= table pointer
RGN	3	Region table index
LVL/BC	5	the page table level of the entry pointed to

## Super Pages

The hierarchical page table allows “super pages” to be defined. These pages bypass lower levels of page tables by using an entry at a high level to represent a block containing many pages.

Normally a PTE with LVL=0 is a pointer to an 16kB memory page. However, super-pages may be defined by specifying a page pointer with a LVL greater than zero. For instance, if T=0 and LVL=1 then the page pointed to is a super-page within an 16MB block of contiguous memory.

T=0, LVL=      Page Size

0	16 kB page
1	16 MB page
2	16 GB page
3	16 TB page
4	16 EB page
5	
6	
7	reserved

A super page pointer contains both a pointer to the block of pages and a super page length field. The length field is provided to restrict memory access to an address range between the super page pointer and the super page pointer plus the number of pages specified in the length. A typical use



For example, a system ROM is located 512 MB before the end of physical memory. The ROM is only 1MB in size. So, it is desired to setup a super page pointer to the ROM and restrict access to a single megabyte. The PTE for this would look like:

V	I <sub>5</sub>	RGN <sub>3</sub>	M	A	0	S	G	SW <sub>2</sub>	~ <sub>4</sub>	MRWX <sub>3</sub>	HRWX <sub>3</sub>	SRWX <sub>3</sub>	URWX <sub>3</sub>
PPN=0x3FFFE0 <sub>22</sub>										NPG=0x03F <sub>10</sub>			
PPN=0xFFFFFFFF <sub>63..32</sub>													

format for 16MB page

V	$1_5$	$RGN_3$	M	A	0	S	G	$SW_2$	$\sim_4$	$MRWX_3$	$HRWX_3$	$SRWX_3$	$URWX_3$
PPN <sub>31..10</sub>										NPG <sub>10</sub>			
PPN <sub>63..32</sub>													

	25	30
	PPN <sub>31</sub> .	

V	2 <sub>5</sub>	RGN <sub>3</sub>	M	A	0	S	G	SW <sub>2</sub>	~ <sub>4</sub>	MRWX <sub>3</sub>	HRWX <sub>3</sub>	SRWX <sub>3</sub>	URWX <sub>3</sub>
PPN <sub>31..20</sub>					NPG <sub>20</sub>								
PPN <sub>63..32</sub>													

# TLB – Translation Lookaside Buffer

## Overview

A simple page map is limited in the translations it can perform because of its size. The solution to allowing more memory to be mapped is to use main memory to store the translations tables.

However, if every memory access required two or three additional accesses to map the address to a final target access, memory access would be quite slow, slowed down by a factor of two or three, possibly more. To improve performance, the memory mapping translations are stored in another unit called the TLB standing for Translation Lookaside Buffer. This is sometimes also called an address translation cache ATC. The TLB offers a means of address virtualization and memory protection. A TLB works by caching address mappings between a real physical address and a virtual address used by software. The TLB deals with memory organized as pages.

Typically, software manages a paging table whose entries are loaded into the TLB as translations are required.

The TLB is a cache specialized for address translations. Thor2023's TLB is quite large being six-way associative with 1024 entries per way. This choice of size was based on the minimum number of block RAMs that could be used to implement the TLB. On a TLB miss the page table is searched for a translation and if found the translation is stored in one of the ways of the TLB. The way selected is determined either randomly or in a least-recently-used fashion as one of the first four ways. The last way may not be updated automatically by a page table search, it must be updated by software.

## Size / Organization

The TLB has 1024 entries per set. The size was chosen as it is the size of one block ram for 32-bit data in the FPGA. This is quite a large TLB. Many systems use smaller TLBs. Typically, systems vary between 64 and 1024 entries. There is not really a need for such a large one, however it is available.

The TLB is organized as a six-way set associative cache. The last way may only be updated by software. The last way allows translations to be stored that will not be overwritten. The first four ways may use hardware LRU replacement in addition to fixed or random replacement.

Way	Page size
0	16kB pages
1	16kB pages
2	16kB pages
3	16kB pages
4	16MB pages
5	16kB pages

Note that 16MB pages do not need multiple ways as there are sufficient TLB entries to allow distinct entries for each 16MB page if the virtual address space is 34-bits or less.

## TLB Entries - TLBE

Closely related to page table entries are translation look-aside buffer, TLB, entries. TLB entries have additional fields to match against the virtual address. The count field is used to invalidate the entire TLB. Note that the least significant 10-bits of the virtual address are not stored as these bits are used as an index for the TLB entry.

Count <sub>6</sub>	LRU <sub>3</sub>
--------------------	------------------

V	LVL/BC <sub>5</sub>	RGN <sub>3</sub>	M	A	T	S	G	SW <sub>2</sub>	CACHE <sub>4</sub>	MRWX <sub>3</sub>	HRWX <sub>3</sub>	SRWX <sub>3</sub>	URWX <sub>3</sub>
PPN <sub>31..0</sub>													
PPN <sub>63..32</sub>													

VPN <sub>41..10</sub>													
VPN <sub>73..42</sub>													
~4	ASID <sub>11..0</sub>							~5	VPN <sub>83..73</sub>				

## Small TLB Entries - TLBE

The small TLB is used for the test system which contains only 512MB of physical RAM to conserve hardware resources. The address ranges are more limited, 40-bits for the physical address and 70-bits for the virtual address.

Count <sub>6</sub>	LRU <sub>3</sub>
--------------------	------------------

V	LVL/BC <sub>5</sub>	RGN <sub>3</sub>	M	A	T	S	G	SW <sub>2</sub>	CACHE <sub>4</sub>	MRWX <sub>3</sub>	HRWX <sub>3</sub>	SRWX <sub>3</sub>	URWX <sub>3</sub>
~6		PPN <sub>25..0</sub>											

VPN <sub>41..10</sub>													
~4	ASID <sub>11..0</sub>							PS	~	VPN <sub>55..42</sub>			

## What is Translated?

The TLB processes addresses including both instruction and data addresses for all modes of operation. It is known as a *unified* TLB.

## Page Size

Because the TLB caches address translations it can get away with a much smaller page size than the page map can for a larger memory system. 4kB is a common size for many systems. There are some indications in contemporary documentation that a larger page size would be better. In this case the TLB uses 16kB. For a 512MB system (the size of the memory in the test system) there are 32768 16kB pages.

## Ways

The first four ways in the TLB are reserved for 16kB page translations. The next way, 4 is reserved for 16MB page translations. The last way is reserved for fixed translations of 16kB pages.

## Management

The TLB unit may be updated by either software or hardware. This is selected in the page table base register. If software miss handling is selected when a translation miss occurs, an exception is generated to allow software to update the TLB. It is left up to software to decide how to update the TLB. There may be a set of hierarchical page tables in memory, or there could be a hash table used to store translations.

## Accessing the TLB

A TLB entry contains too much information to be updated with a single register write. Since the information must also be updated atomically to ensure correct operation, the TLB update occurs in an indirect fashion. First holding registers are loaded with the desired values, then all the holding registers are written to the TLB in a single atomic cycle. The TLB is addressed in the physical memory space in the address range \$F...FE000xx. There are eight buckets which must be filled with TLB info using store instructions. Then address \$F...FE0007E is written to causing the TLB to be updated.

The low order bits of the bucket six determine which way to update in the TLB if the algorithm is a fixed way algorithm. Otherwise, if LRU is selected the LRU entry will be updated, otherwise a way to update will be selected randomly. The data is octa-byte aligned.

00	TLBE (PTE <sub>63..0</sub> )									
08						TLBE (PTE <sub>95..64</sub> )				
10	TLBE (VPN <sub>63..0</sub> )									
18						TLBE (VPN <sub>95..64</sub> )				
20	TLB Miss Address <sub>63..0</sub>									
28	~ <sub>4</sub>	Miss ASID <sub>12</sub>	~ <sub>16</sub>			TLB Miss Address <sub>95..64</sub>				
30 to 68										
70						AL <sub>2</sub>	0	Entry Num <sub>10</sub>	~	Way <sub>4</sub>
78	RWTRIG	WTRIG	RTRIG	~ <sub>8</sub>		~ <sub>32</sub>				

### ADR

7C	No operation
7D	Read TLBE
7E	Write TLBE
7F	Read and Write TLBE

## ?RWX<sub>3</sub>

If RWX3 attributes are specified non-zero, then they will override the attributes coming from the region table. Otherwise RWX attributes are determined by the region table.

## CACHE<sub>4</sub>

The cache<sub>4</sub> field is combined with the cache attributes specified in the region table. The region table takes precedence; however, if the cache<sub>4</sub> field indicates non-cache-ability then the data will not be cached.

### Example TLB Update Routine

```

_TLBMap:
    ldo          a0,0[sp]
    ldo          a1,8[sp]
    ldo          a2,16[sp]
    ldo          a3,24[sp]
    ; <lock TLB update semaphore>
    sto          a0,0xFFE00000          # TLBE value
    sto          a1,0xFFE00008          # TLBE value
    sto          a2,0xFFE00010          # TLBE value
    sto          a3,0xFFE00070          # control
    stb          a0,0xFFE0007E          # triggers a TLB update
    ; <unlock TLB update semaphore>
    add          sp,sp,32
    rts

```

## TLB Entry Replacement Policies

The TLB supports three algorithms for replacement of entries with new entries on a TLB miss. These are fixed replacement (0), least recently used replacement (1) and random replacement (2). The replacement method is stored in the AL<sub>2</sub> bits of the page table base register.

For fixed replacement, the way to update must be specified by a software instruction. Least recently used replacement, LRU, selects the least recently used address translation to be overwritten. Random replacement chooses a way to replace at random.

## Flushing the TLB

The TLB maintains the address space (ASID) associated with a virtual address. This allows the TLB translations to be used without having to flush old translations from the TLB during a task switch.

## Reset

On a reset the TLB is preloaded with translations that allow access to the system ROM.

### Global Bit

In addition to the ASID the TLB entries contain a bit that indicates that the translation is a global translation and should be present in every address space.

## Card Table

### Overview

Also present in the memory system is the Card table. The card table is a telescopic memory which reflects with increasing detail where in the memory system a pointer write has occurred. This is for the benefit of garbage collection systems. Card table is updated using a write barrier when a pointer value is stored to memory, or it may be updated automatically using the STPTR instruction.

### Organization

At the lowest level memory is divided into 256-byte card memory pages. Each card has a single byte recording whether a pointer store has taken place in the corresponding memory area. To cover a 512MB memory system 2MB card memory is required at the outermost layer. A byte is used rather than a bit to allow byte store operations to update the table directly without having to resort to multiple instructions to perform a bit-field update.

To improve the performance of scanning a hardware card table, HCT, is present which divides memory at an upper level into 8192-byte pages. The hardware card table indicates if a pointer store operation has taken place in one of the 8192-byte pages. It is then necessary to scan only cards representing the 8192-byte page rather than having to scan the entire 2MB card table. Note that this memory is organized as 2048 32-bit words. Allowing 32-bits at a time to be tested.

To further improve performance a master card table, MCT, is present which divides memory at the uppermost layer into 16-MB pages.

Layer	Resolving Power	
0	2 MB	256B pages
1	64k bits	8kB pages
2	32 bits	16 MB pages

There is only a single card memory in the system, used by all tasks.

## Location

Card memory must be based at physical address zero, extending up to the amount of card memory required. This is so that the address calculation of the memory update may be done with a simple right-shift operation.

## Operation

As a program progresses it writes pointer values to memory using the write barrier. Storing a pointer triggers an update to all the layers of card memory corresponding to the main memory location written. A bit or byte is set in each layer of the card memory system corresponding to the memory location of the pointer store.

The garbage collection system can very quickly determine where pointer stores have occurred and skip over memory that has not been modified.

## Sample Write Barrier

```
; Milli-code routine for garbage collect write barrier.
; This sequence is short enough to be used in-line.
; Three level card memory.
; a2 is a register pointing to the card table.
; STPTR will cause an update of the master card table, and hardware card table.
;
```

GCWriteBarrier:

```
STPTR    a0,[a1]           ; store the pointer value to memory at a1
LSR      t0,a1,#8          ; compute card address
STB      r0,[a2+t0]        ; clear byte in card memory
```

## System Memory Map

There are several components to the system which use tables in memory. These tables are statically allocated at the time the system is built. The table sizes depend on the size of main memory. The card memory table must be located at address zero. So, it is probably best to group the tables together at the low end of memory.

Address	Usage	
\$00000000 to \$001FFFFFFF	Card Table (2 MB)	
\$00210000 to \$0022FFFF	PAM (128kB 2 copies)	
\$00280000 to \$0029FFFF	Key memory (128 kB)	



# Instruction Set

## Overview

Thor was a variable length instruction set with instructions varying in length from one to eight bytes. Thor2023 is primarily a fixed length instruction with provision for additional instruction words used for constants. Reducing the variety of instruction sizes makes implementation of decoders more economical.

## Predicated Instruction Execution

Some processors include the ability to execute virtually any instruction conditionally, for example the ARM processor or INTEL Itanium IA64. It's a powerful means of removing branches from the instruction stream. Sequences of instructions executed with predicates rather than branching around the instructions should be kept short. The issue is the amount of time spent fetching the instructions and treating them as NOPs versus the time it would take to branch around the instructions. A compiler can optimize this and choose the best means. One of the problems of predicates is that they use up bits in the instruction regardless of whether they are useful. For instance, the Itanium has a six-bit field in virtually every instruction. The result is that a wider instruction format of 41 bits is used. A second problem with predicates is that they act like a second instruction being executed at the same time as the instruction they are associated with. The predicate operation requires a predicate register read, and a predicate evaluation operation. This adds complexity to the processor. Predicate registers are another form of register that must be present and bypassed in an overlapped or superscalar design.

The first Thor processing core features uses a whole byte for predicates, but gains back some of the opcode space by using redundant forms of the predicates as single byte instructions. The most recent version of Thor has two means of predication. A vector mask register may be specified for a scalar operation in which case the scalar operation takes place only if the mask register is equal to one. The second means of predication is via an instruction modifier. An instruction modifier precedes the instruction to add to or modify its operation. Since predicates are used infrequently the use of a modifier is an efficient manner to encode the operation.

# Instruction Descriptions

## Scalar Instructions Layout

39	9	8	7	5	4	0
Payload <sub>31</sub>	0	Size	Opcode			

## Vector Instruction Layout

A vector instruction is identical to its scalar counterpart except that one of the vector bits of the instruction is set and there may be an additional field present to specify the mask register. This field adds one byte to the instruction.

47	46	40	39	31	30	29	9	8	7	5	4	0
~	Mask	Payload <sub>9</sub>	Vb	Payload <sub>21</sub>	V	Size	Opcode					

## Register-Register Vector Decode

Vc, Vb, V	Rb	Ra	Rt	Mask
000	scalar	scalar	scalar	No
001	scalar	scalar	scalar	Yes
010	scalar	vector	vector	No
011	scalar	vector	vector	Yes
100	vector	vector	vector	No
101	vector	vector	vector	Yes

## Register-Immediate Vector Decode

Vc, V	Ra	Rt	Mask
00	scalar	scalar	No
01	scalar	scalar	Yes
10	vector	vector	No
11	vector	vector	Yes

## Opcode Maps

## Major Opcode

	0	1	2	3	4	5	6	7
0x	0 TRAP	1	2 {R2}	3 {CSR}	4 ADDI	5 CMPI	6 MULI	7 DIVI
	8 ANDI	9 ORI	10 EORI	11 CHK	12 {FLT2}	13 {BIT}	14 {SHIFT}	15 FMA
1x	16 LOAD	17 LOADZ	18 STORE	19 BMAP	20 FADDI	21 FCMPI	22 FMULI	23 FDIVI
	24 JSR, JMP	25 CMPXCHG	26 {AMO}	27 Bcc	28 LBcc	29 DBcc	30	31 PFX / NOP / REGS

## {R2} Operations

	0	1	2	3	4	5	6	7
0x	0 CNTLZ	1	2 CNTPOP	3 ABS	4 ADD	5 CMP	6 MUL	7 DIV
	8 AND	9 OR	10 EOR	11 SEQ SNE	12	13 CHRNDX	14 CLMUL	15 SQRT
1x	16 DIF	17 PTRDIF	18 REVBIT	19 BMAP	20 SGE SLT	21 SGT SLE	22 SM4ED	23 SM4KS
	24 JMP / JSR	25	26 AES64DS	27 AES64DSM	28 AES64ES	29 AES64ESM	30 AES64KS1I	31 AES64KS2
2x	32 PRED	33 CARRY	34 REP	35 ATOM	36 ROUND	37	38	39
	40 V2BITS	41 BITS2V	42 VEX	43 VEINS	44 VGNDX	45	46 SGEU	47 SGTU
3x	48 MIN	49 MAX	50 BMM	51 MUX	52	53 AES64IM	54 SM3P0	55 SM3P1
	56 SHA256 SIG0	57 SHA256 SIG1	58 SHA256 SUM0	59 SHA256 SUM1	60 SHA512 SIG0	61 SHA512 SIG1	62 SHA512 SUM0	63 SHA512 SUM1

## {CSR – Func3}

	0	1	2	3	4	5	6	7
0x	0 CSR RD	1 CSR WR	2 CSR RS	3 CSR RC	4 CSR EOR	5	6	7

## {BIT – Func3}

	0	1	2	3	4	5	6	7
0x	0 CLR	1 SET	2 COM	3 SBX	4 EXTU	5 EXTS	6	7 {BITRR}

**{SHIFT – Func5}**

	0	1	2	3	4	5	6	7
0x	0 ASL	1 ASR	2 LSL	3 LSR	4 ROL	5 ROR	6	7
	8 ASLI	9 ASRI	10 LSLI	11 LSRI	12 ROLI	13 RORI	14	15
1x	16 VSHLV	17 VSHRV	18	19	20	21	22	23
	24 VSHLVI	25 VSHRVI	26	27	28	29	30	31

**{FLT2} Operations**

	0	1	2	3	4	5	6	7
0x	0 FSCALEB	1 {FLT1}	2 FMIN	3 FMAX	4 FADD	5 FCMP	6 FMUL	7 FDIV
	8 FScc	9	10	11	12	13	14 FNXT	15 FREM
1x	16							
	24							

**{FLT1} Operations**

	0	1	2	3	4	5	6	7
0x	0	1	2 FOTI	3 ITOF	4	5	6 FSIGN	7 FSIG
	8 FSQRT	9 FS2D	10 FS2Q	11 FD2Q	12	13	14 ISNAN	15 FINITE
1x	16	17	18	19	20	21 FTRUNC	22	23 FRES
	24	25 FD2S	26 FQ2S	27 FQ2D	28	29	30 FCLASS	31
2x	32 FABS	33	34 FNEG	35	36	37	38	39
	40							
3x	48							
	56							

## {AMO} Operations

	0	1	2	3	4	5	6	7
0x	0 SWAP	1	2 MIN	3 MAX	4 ADD	5	6 ASL	7 LSR
	8 AND	9 OR	10 EOR	11	12 MINU	13 MAXU	14	15 CAS
1x	16 SWAPI	17	18 MIN	19 MAX	20 ADDI	21	22 ASLI	23 LSRI
	24 ANDI	25 ORI	26 EORI	27	28 MINU	29 MAXU	30	31 CAS

## {PFX – Sz3}

	0	1	2	3	4	5	6	7
0x	0 PFX32	1 PFX64	2 PFX128	3	4	5 REGS	6	7 NOP

## Operand Swapping

Many instructions allow first and second source operands to be swapped. This is indicated by the swap 'S' bit in the instruction. This is particularly useful for instructions that are non-commutative like SUB and DIV.

### Operand Swap

Operand Order	S
Normal	0
1 <sup>st</sup> and 2 <sup>nd</sup> Swapped	1

## Operand Sizes


Many instructions support five different operand sizes: byte, wyde, tetra, octa and hexi. The operand size is selected by suffixing the mnemonic with 'b' for byte, 'w' for wyde, 't' for tetra, 'o' for octa and 'h' for hexi. Size code 6 selects decimal arithmetic mode.

Sz <sub>3</sub>	Ext.	Operand
0	.b	8-bit Byte
1	.w	16-bit Wyde
2	.t	32-bit Tetra
3	.o	64-bit Octa
4	.h	128-bit Hexi
5		Reserved
6	.d	128-bit decimal
7		reserved


# Arithmetic Operations

## Representations


long




int




short




char



byte



decimal



Decimal integers use densely packed decimal format which provide 38 digits of precision.

## Arithmetic Operations

Arithmetic operations include addition, subtraction, multiplication and division. These are available with the ADD, SUB, CMP, MUL, and DIV instructions. There are several variations of the instructions to deal with signed and unsigned values. The format of the typical immediate mode instruction is shown below:

### ADD Rt,Ra,Imm<sub>16</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	Vc	Op	Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	4 <sub>5</sub>						

Immediate instructions may have the constant extended via the use of postfix immediates.

### ADD Rt,Ra,Imm<sub>15</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
~7	Vc	Op	~8	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	4 <sub>5</sub>						
Immediate <sub>31..0</sub>										0 <sub>3</sub>	31 <sub>5</sub>					

There may seem to be significant wasted space in the instruction when an instruction postfix is used. However, the use of a postfix is the rare case which occurs when a fifteen-bit immediate value is not sufficient. Having the postfix begin with bit 0 to 31 encoded is to allow for instructions that do not have space for an immediate field in the instruction. The postfix is kept consistent between all instructions to make decoding easier to handle and smaller resource wise.

Note that all arithmetic instructions can use an immediate value via a postfix immediate. Not all arithmetic instructions support a fifteen-bit immediate field. Instead, when a postfix is used it will override the value coming from register Rb. The following instruction ignores the Rb register value and multiplies by a postfix immediate.

### MULSU Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
6 <sub>6</sub>	1	Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						
Immediate <sub>31..0</sub>										0 <sub>3</sub>	31 <sub>5</sub>								

There are both signed and unsigned versions of the arithmetic operations. However, note there is no signed or unsigned compare operation as a single compare instruction produces results for both signed and unsigned comparisons. Signed and unsigned ADD and SUB currently work the same way. Two separate versions have been reserved to support the overflow exception in the future.

## Logical Operations

Thor logic operations include only the basic 'and', 'or' and 'xor' operations. Other variations of the instructions like 'nand', or 'nor' are possible by inverting registers. The assembler recognizes



some of the possible combinations of register inversion as distinct instructions like ‘nand’ or ‘nor’. Unlike Table888 there are more immediate forms available even for rarely used instructions.

## Immediate Operate Functions

Most instructions at the root level are immediate operate instructions, the ones that are not are heavily bordered in the table below. Note that in some cases non-immediate format instructions may use an immediate via a postfix instruction.

### Major Opcode

	0	1	2	3	4	5	6	7
0x	0 TRAP	1	2 {R2}	3 {CSR}	4 ADDI	5 CMPI	6 MULI	7 DIVI
	8 ANDI	9 ORI	10 EORI	11 CHK	12 {FLT2}	13 {BIT}	14 {SHIFT}	15 FMA
1x	16 LOAD	17 LOADZ	18 STORE	19 BMAP	20 FADDI	21 FCMPI	22 FMULI	23 FDIVI
	24 JSR, JMP	25 CMPXCHG	26 {AMO}	27 Bcc	28 FBcc	29 DBcc	30	31 PFX / NOP

# ABS – Absolute Value

## Description:

This instruction computes the absolute value of the contents of the source operand and places the result in Rt.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

**Integer Instruction Format: R2**

**ABS Rt, Ra**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
3 <sub>6</sub>	~ <sub>2</sub>	0	0	0	0 <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>							

**Clock Cycles: 5**

## Operation:

If  $Ra < 0$   
      $Rt = -Ra$   
 else  
      $Rt = Ra$

**Execution Units:** Integer ALU #0

**Clock Cycles: 1**

**Exceptions:** none

**Notes:**

# ADD - Addition

## Description:

Add two source operands and place the sum in the target register. All registers are treated as integer registers. Arithmetic is signed twos-complement values unless decimal mode is selected ( $SZ_3=6$ ) in which case values are treated as BCD numbers. This instruction may be used with the [CARRY](#) modifier to perform extended precision addition.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

## Operation:

$$Rt = Ra + Rb \text{ or } Rt = Ra + Imm$$

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

**Instruction Formats: R2, RI**

### ADD Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
4 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 5**

### ADD Rt,Ra,Imm<sub>15</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	Vc	0	Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	4 <sub>5</sub>						

**Clock Cycles: 5**

### ADD Rt,Ra,Imm<sub>32</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	Vc	0	Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	4 <sub>5</sub>						

**Clock Cycles: 5**

## Cycles

IFETCH	*
DECODE	*
OFETCH	*
EXECUTE	*
WRITEBACK	*
Total	5

# AND – Bitwise And

## Description:

Bitwise ‘and’ two source operands and place the result in the target register. The one’s complement of operands may be used by setting the appropriate ‘S’ bit in the instruction.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

**Clock Cycles:** 1

## Operation:

$Rt = Ra \& Rb$  or  $Rt = Ra \& Imm$

**Instruction Formats:** R2, RI

## AND Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
8 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 5

## AND Rt,Ra,Imm<sub>15</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>			Vc	0	Imm <sub>7..0</sub>		Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	8 <sub>5</sub>			

**Clock Cycles:** 5

**Execution Units:** All Integer ALU’s

**Exceptions:** none

## Notes:

Cycles

IFETCH	*
DECODE	*
OFETCH	*
EXECUTE	*
WRITEBACK	*
Total	5

# BMAP – Byte Map

## Description:

First the target register is cleared, then bytes are mapped from the 16-byte source Ra into bytes in the target register. This instruction may be used to permute the bytes in register Ra and store the result in Rt. This instruction may also pack bytes, wydes or tetras. The map is determined by the low order 64-bits of register Rb or a 64-bit immediate constant. Bytes which are not mapped will end up as zero in the target register.

## Instruction Formats: R2, RI

### BMAP Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
19 <sub>6</sub>	~	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 5**

### BMAP Rt,Ra,Imm<sub>48</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
~7	Vc	0	~8	Sa	Ra6	St	Rt6	V	Sz3	195						
Immediate31..0										03	315					
Immediate63..32										13	315					

**Clock Cycles: 5**

## Operation:

### Vector Operation

**Execution Units:** First Integer ALU

**Clock Cycles: 1**

**Exceptions:** none

## Notes:

Cycles

IFETCH	*
DECODE	*
OFETCH	*
EXECUTE	*
WRITEBACK	*
Total	5

# CHK – Check Register Against Bounds

## Description:

A register is compared to two values. If the register is outside of the bounds defined by Rb and an immediate value then an exception will occur. Ra must be greater than or equal to Rb and Ra must be less than the immediate.

## Instruction Formats:

### CHK Ra, Rb, Cn – Register direct

39	38	37	32	31	30	29	28	23	22	21	16	15	14	12	11	9	8	7	5	4	0
~	Vc	Rc <sub>6</sub>	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	~	~ <sub>3</sub>	Cn <sub>3</sub>	V	Sz <sub>3</sub>	11 <sub>5</sub>							

**Clock Cycles: 1**

cn <sub>3</sub>	exception when not
0	Ra >= Rb and Ra < Rc
1	Ra >= Rb and Ra <= Rc
2	Ra > Rb and Ra < Rc
3	Ra > Rb and Ra <= Rc
4	not (Ra >= Rb and Ra < Rc)
5	not (Ra >= Rb and Ra <= Rc)
6	not (Ra > Rb and Ra < Rc)
7	not (Ra > Rb and Ra <= Rc)

### CHKI Ra, Imm, Cn

39	33	32	31	30	29	28	23	22	21	16	15	12	11	9	8	7	5	4	0
Imm <sub>10..4</sub>		Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	Imm <sub>3..0</sub>		Cn <sub>3</sub>	V	Sz <sub>3</sub>	11 <sub>5</sub>					

**Clock Cycles: 1**

cn <sub>3</sub>	exception when not
0	Ra >= Rb and Ra < Imm
1	Ra >= Rb and Ra <= Imm
2	Ra > Rb and Ra < Imm
3	Ra > Rb and Ra <= Imm
4	not (Ra >= Rb and Ra < Imm)
5	not (Ra >= Rb and Ra <= Imm)
6	not (Ra > Rb and Ra < Imm)
7	not (Ra > Rb and Ra <= Imm)

**Clock Cycles: 1**

**Execution Units:** Integer ALU

**Exceptions:** bounds check

**Notes:**

The system exception handler will typically transfer processing back to a local exception handler.

# CLMUL – Carry-less Multiply

## Description:

Compute the low order product bits of a carry-less multiply.

## Instruction Formats:

### CLMUL Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
14 <sub>6</sub>	~	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 7 + n, where n = number of bits

### CLMUL Rt,Ra,Imm<sub>8</sub>

39	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
14 <sub>6</sub>	~	Vc	1	Imm <sub>8</sub>			Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>				

**Clock Cycles:** 7 + n, where n = number of bits

**Exceptions:** none

**Execution Units:** First Integer ALU

Operations

$$Rt = Ra * Rb$$

## Vector Operation

for x = 0 to VL - 1

if (Vm[x]) Vt[x] = Va[x] \* Vb[x]

else if (z) Vt[x] = 0

else Vt[x] = Vt[x]

**Exceptions:** none



# CMP - Comparison

## Description:

Compare two source operands and place the result in the target register. The result is a vector identifying the relationship between the two source operands as signed and unsigned integers.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

## Operation:

$Rt = Ra \text{ ? } Rb \text{ or } Rt = Ra \text{ ? } Imm \text{ or } Rt = Imm \text{ ? } Ra$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

**Instruction Formats: R2, RI**

**CMP Rt, Ra, Rb**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
5 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 4/6

**CMP Rt,Ra,Imm<sub>15</sub>**

39		33	32	31	30		23	22	21	16	15	14	9	8	7	5	4	0
	Imm <sub>14..8</sub>	Vc	S		Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	5 <sub>5</sub>						

**Clock Cycles:** 4/6

Rt Bit	Mnem.	Meaning	Test
		<b>Integer Compare Results</b>	
0	EQ	= equal	$a == b$
1	NE	< > not equal	$a <> b$
2	LT	< less than	$a < b$
3	LE	<= less than or equal	$a <= b$
4	GE	>= greater than or equal	$a >= b$
5	GT	> greater than	$a > b$
6	BC	Bit clear	$!a[b]$
7	BS	Bit set	$a[b]$
8			
9			
10	LO / CS	< unsigned less than	$a < b$
11	LS	<= unsigned less than or equal	$a <= b$
12	HS / CC	unsigned greater than or equal	$a >= b$
13	HI	unsigned greater than	$a > b$
14			

15			
----	--	--	--

## CMPS.B – Signed Byte Comparison

### Description:

Compare two source operands and place the result in the target register. The result is a vector identifying the relationship between the two source operands as signed integers.

**Supported Operand Sizes:** .b

### Operation:

$Rt = Ra \text{ ? } Rb \text{ or } Rt = Ra \text{ ? } Imm \text{ or } Rt = Imm \text{ ? } Ra$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

### Instruction Formats:

**CMPS.B Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
5 <sub>6</sub>	1	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	0 <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 4/6

**CMPS.B Rt,Ra,Imm<sub>15</sub>**

39	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
~ <sub>6</sub>	1	Vc	S	Imm <sub>7..0</sub>			Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	0 <sub>3</sub>	5 <sub>5</sub>				

**Clock Cycles:** 4/6

Rt bit	Mnem.	Meaning	Test
		<b>Integer Compare Results</b>	
0	EQ	= equal	
1	NE	< > not equal	
2	LT	< less than	
3	LE	<= less than or equal	
4	GE	>= greater than or equal	
5	GT	> greater than	
6			
7			

## CMPU.B – Unsigned Byte Comparison

### Description:

Compare two source operands and place the result in the target register. The result is a vector identifying the relationship between the two source operands as unsigned integers.

**Supported Operand Sizes:** .b, .w, .t, .o, .c, .p, .n

### Operation:

$Rt = Ra ? Rb$  or  $Rt = Ra ? Imm$  or  $Rt = Imm ? Ra$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

### Instruction Formats:

**CMPU.B Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
5 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	0 <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

**CMPU.B Rt,Ra,Imm<sub>8</sub>**

39	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
~ <sub>6</sub>	0	Vc	S	Imm <sub>7..0</sub>			Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	0 <sub>3</sub>	5 <sub>5</sub>				

**Clock Cycles:** 1

Rt bit	Mnem.	Meaning	Test
		<b>Integer Compare Results</b>	
0	EQ	= equal	
1	NE	< > not equal	
2	LTU	< less than	
3	LEU	<= less than or equal	
4	GEU	>= greater than or equal	
5	GTU	> greater than	
6			
7			

## CNTLZ – Count Leading Zeros

### Description:

This instruction counts the number of consecutive zero bits beginning at the most significant bit towards the least significant bit.

**Supported Operand Sizes:** .b, .w, .t, .o

**Integer Instruction Format: R1**

**CNTLZ Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0 <sub>6</sub>	~	Vc	0	Vb	0	0 <sub>6</sub>	0	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

**Operation:**

**Execution Units:** Integer ALU #0

**Clock Cycles: 1**

**Exceptions:** none

**Notes:**

## CNTLO – Count Leading Ones

### Description:

This instruction counts the number of consecutive one bits beginning at the most significant bit towards the least significant bit.

**Supported Operand Sizes:** .b, .w, .t, .o

**Integer Instruction Format: R1**

**CNTLO Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0 <sub>6</sub>	~	Vc	0	Vb	0	0 <sub>6</sub>	1	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

**Operation:**

**Execution Units:** Integer ALU #0

**Clock Cycles: 1**

**Exceptions:** none

**Notes:**

# CNTPOP – Count Population

## Description:

This instruction counts the number of bits set in a register.

**Supported Operand Sizes:** .b, .w, .t, .o

**Integer Instruction Format: R1**

**CNTPOP Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
2 <sub>6</sub>	~	Vc	0	Vb	0	0 <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

**Operation:**

**Execution Units:** Integer ALU #0

**Clock Cycles: 1**

**Exceptions:** none

**Notes:**

# CSR – Control and Special Registers Operations

## Description:

Perform an operation on a CSR.

Operation	Op <sub>3</sub>	
Read CSR	0	
Write CSR	1	
Or to CSR (set bits)	2	
And complement to CSR (clear bits)	3	
Exclusive Or to CSR (flip bits)	4	

Supported Operand Sizes: N/A

Regno		
\$000	reserved	Not used
\$002	sr	Status register (privileged)
\$120	Tick	Tick count (read only)
\$121	Coreno	Core number ( read only) (privileged)
\$127		

## Instruction Formats:

**OR Rt, Ra, CSR**

**ANDC Rt, Ra, CSR**

**EOR Rt, Ra, CSR**

**CSR Rt,Ra,#Regno<sub>12</sub>**

39	38	37	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Regno <sub>13..8</sub>	S	Regno <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Op <sub>3</sub>	3 <sub>5</sub>						

**Clock Cycles: 1**

**CSR Rt, #imm,#Regno<sub>12</sub>**

39	38	37	32	31	30	23	22		16	15	14	9	8	7	5	4	0
1	Vc	Regno <sub>13..8</sub>	S	Regno <sub>7..0</sub>	Imm <sub>7</sub>			St	Rt <sub>6</sub>	V	Op <sub>3</sub>	3 <sub>5</sub>					

# DIVS – Signed Division

## Description:

Divide source dividend operand by divisor operand and place the quotient in the target register.  
All registers are integer registers. Arithmetic is signed twos-complement values.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$Rt = Ra / Rb$  or  $Rt = Ra / Imm$  or  $Rt = Imm / Ra$

## Instruction Formats:

### DIVS Rt, Ra, Rb – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
7 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 100**

### DIVS Rt,Ra,Imm<sub>16</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>			Vc	S	Imm <sub>7..0</sub>		Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	7 <sub>5</sub>			

**Clock Cycles: 100**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# DIVU – Unsigned Division

## Description:

Divide source dividend operand by divisor operand and place the sum in the target register. All registers are integer registers. Arithmetic is unsigned twos-complement values.

15-bit immediate mode is not available for this instruction.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$Rt = Ra / Rb$  or  $Rt = Ra / Imm$  or  $Rt = Imm / Ra$

## Instruction Formats:

**DIVU Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
7 <sub>6</sub>	1	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 100**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**



# EOR – Bitwise Exclusive Or

## Description:

Bitwise exclusive ‘or’ two source operands and place the sum in the target register. All registers are integer registers. Arithmetic is signed twos-complement values.

**Supported Operand Sizes:** .b, .w, .t, .o, .c, .p, .n

## Operation:

$$Rt = Ra \wedge Rb \text{ or } Rt = Ra \wedge Imm$$

## Instruction Formats:

### EOR Rt, Ra, Rb – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
10 <sub>6</sub>	~	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### EOR Rt,Ra,Imm<sub>16</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>				Vc	S	Imm <sub>7..0</sub>		Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	10 <sub>5</sub>		

**Clock Cycles: 1**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# ENOR – Bitwise Exclusive Nor

## Description:

Bitwise exclusive ‘nor’ two source operands and place the result in the target register.

**Supported Operand Sizes:** .b, .w, .t, .o, .c, .p, .n

## Operation:

$$Rt = \sim(Ra \wedge Rb) \text{ or } Rt = \sim(Ra \wedge Imm)$$

## Instruction Formats:

### ENOR Rt, Ra, Rb – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
10 <sub>6</sub>	~	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### ENOR Rt,Ra,Imm<sub>16</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>			Vc	S	Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>	10 <sub>5</sub>				

**Clock Cycles: 1**

**Clock Cycles: 1**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

## MODS – Signed Modulus

### Description:

Divide source dividend operand by divisor operand and place the remainder in the target register.  
All registers are integer registers. Arithmetic is signed twos-complement values.

15-bit Immediate mode is not available for this instruction.

**Supported Operand Sizes:** .b, .w, .t, .o

### Operation:

$Rt = Ra / Rb$  or  $Rt = Ra / Imm$  or  $Rt = Imm / Ra$

### Instruction Formats:

**MODS Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
$7_6$	0	Vc	1	Vb	Sb	$Rb_6$	Sa	$Ra_6$	St	$Rt_6$	V	$Sz_3$	$2_5$						

**Clock Cycles: 100**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

## MODU – Unsigned Modulus

### Description:

Divide source dividend operand by divisor operand and place the remainder in the target register. All registers are integer registers. Arithmetic is unsigned twos-complement values.

15-bit Immediate mode is not available for this instruction.

**Supported Operand Sizes:** .b, .w, .t, .o

### Operation:

$Rt = Ra / Rb$  or  $Rt = Ra / Imm$  or  $Rt = Imm / Ra$

### Instruction Formats:

**MODU Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
7 <sub>6</sub>	1	Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 100**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# MULS – Multiply Signed

## Description:

Multiply two source operands and place the product in the target register. All registers are treated as integer registers. Arithmetic is signed twos-complement values. The ‘S’ flag indicates to perform an unsigned multiply.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$$Rt = Ra * Rb \text{ or } Rt = Ra * Imm$$

## Instruction Formats:

### MULS Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
6 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 7+n, where n is the number of bits.

### MULS Rt,Ra,Imm<sub>15</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>			Vc	S	Imm <sub>7..0</sub>		Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	6 <sub>5</sub>			

**Clock Cycles:** 7+n

**Clock Cycles:** 12

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# MULSU – Signed-Unsigned Multiplication

## Description:

Multiply two source operands and place the product in the target register. The first operand is a signed value, the second operand is unsigned. All registers are treated as integer registers. Arithmetic is twos-complement values.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

## Operation:

$$Rt = Ra * Rb \text{ or } Rt = Ra * Imm$$

## Instruction Formats:

### MULSU Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
6 <sub>6</sub>	1	Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 7+n, where n is the number of bits.

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# MULU – Unsigned Multiplication

## Description:

Multiply two source operands and place the product in the target register. All registers are treated as integer registers. Arithmetic is signed twos-complement values. The ‘S’ flag indicates to perform an unsigned multiply. Unsigned multiply can be used during index calculations.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

## Operation:

$$Rt = Ra * Rb \text{ or } Rt = Ra * Imm$$

## Instruction Formats:

### MULU Rt, Ra, Rb – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
6 <sub>6</sub>	0	Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 7+n, where n is the number of bits.

### MULU Rt,Ra,Imm<sub>16</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>			Vc	1	Imm <sub>7..0</sub>		Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	6 <sub>5</sub>			

**Clock Cycles:** 7+n

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# NAND – Bitwise And and Invert

## Description:

Bitwise ‘nand’ two source operands and place the result in the target register.

**Supported Operand Sizes:** .b, .w, .t, .o, .c, .p, .n

**Clock Cycles:** 1

## Operation:

$$Rt = \sim(Ra \& Rb)$$

## Instruction Formats:

**NAND Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
8 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

**NAND Rt,Ra,Imm<sub>15</sub>**

39		33	32	31	30		23	22	21	16	15	14	9	8	7	5	4	0
	Imm <sub>14..8</sub>	Vc	0		Imm <sub>7..0</sub>	Sa		Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>					8 <sub>5</sub>	

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**



# NOR – Bitwise Or and Invert

## Description:

Bitwise ‘or’ two source operands invert the result and place the result in the target register. All registers are integer registers.

**Supported Operand Sizes:** .b, .w, .t, .o, .c, .p, .n

## Operation:

$$Rt = \sim(Ra \mid Rb)$$

## Instruction Formats:

### NOR Rt, Ra, Rb – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
9 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### NOR Rt,Ra,Imm<sub>15</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>		Vc	0	Imm <sub>7..0</sub>		Sa	Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>	9 <sub>5</sub>				

**Clock Cycles: 1**

**Execution Units:** All Integer ALU’s

**Exceptions:** none

**Notes:**

# OR – Bitwise Or

## Description:

Bitwise ‘or’ two source operands and place the sum in the target register. All registers are integer registers. Arithmetic is signed twos-complement values.

**Supported Operand Sizes:** .b, .w, .t, .o, .c, .p, .n

## Operation:

$$Rt = Ra \mid Rb \text{ or } Rt = Ra \mid Imm$$

## Instruction Formats:

### OR Rt, Ra, Rb – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
9 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### OR Rt,Ra,Imm<sub>15</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>			Vc	0	Imm <sub>7..0</sub>		Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	9 <sub>5</sub>			

**Clock Cycles: 1**

**Clock Cycles: 2**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# PFX – Constant Postfix

## Description:

The PFX instruction postfix is used to provide large constants for use in the preceding instruction as the immediate constant for the instruction. There are four postfix instructions providing 16, 32, 64, and 128-bit constants.

Postfixes are normally caught at the decode stage and do not progress further in the pipeline. They are treated as a NOP instruction.

**Supported Operand Sizes:** N/A

## Instruction Format: PFX0

This format provides a 16-bit constant, and sign extends the value to the width of the machine. For float operations the constant is converted from a half precision float value to a quad precision float value.

24	8	7	5	4	0
Immediate <sub>16</sub>				0 <sub>3</sub>	31 <sub>5</sub>

## Instruction Format: PFX1

This format provides a 32-bit constant, and sign extends the value to the width of the machine. For float operations the constant is converted from a single precision float value to a quad precision float value.

39	8	7	5	4	0
Immediate <sub>32</sub>				1 <sub>3</sub>	31 <sub>5</sub>

## Instruction Format: PFX2

This format provides a 64-bit constant, and sign extends the value to the width of the machine. For float operations the constant is converted from a double precision float value to a quad precision float value.

71	8	7	5	4	0
Immediate <sub>64</sub>				2 <sub>3</sub>	31 <sub>5</sub>

## Instruction Format: PFX3

This format provides a 128-bit constant.

135	8	7	5	4	0
Immediate <sub>127..0</sub>				4 <sub>3</sub>	31 <sub>5</sub>



## REGS – Registers List

### Description:

This instruction postfix specifies additional operands for the preceding instruction. When applied to a load or store operation it causes a multiple register load or store to be performed. This postfix should be the last postfix applied to an instruction coming after a postfix for the displacement.

### Instruction Format:

71	8	7	5	4	0
Reglist <sub>64</sub>	5 <sub>3</sub>			31 <sub>5</sub>	

Bit	Reg	Bit	Reg	Bit	Reg
0	A0	25	A3	50	VM7
1	A1	26	A4	51	LR1
2	A2	27	A5	52	LR2
3	T0	28	A6	53	LR3
4	T1	29	A7	54	TP
5	T2	30	A8	55	GP1
6	T3	31	A9	56	CTA
7	T4	32	A10	57	R48
8	T5	33	T8	58	R49
9	T6	34	T9	59	R50
10	T7	35	T10	60	R51
11	S0	36	T11	61	ASP
12	S1	37	S9	62	SSP
13	S2	38	S10	63	HSP
14	S3	39	S11		
15	S4	40	S12		
16	S5	41	S13		
17	S6	42	S14		
18	S7	43	S15		
19	S8	44	VM1		
20	VM0	45	VM2		
21	LC	46	VM3		
22	LR0	47	VM4		
23	GP0	48	VM5		
24	FP	49	VM6		

### Assembler Syntax:

### Example:



## REVBIT – Reverse Bit Order

### Description:

This instruction reverses the order of bits in Ra and stores the result in Rt.

### Integer Instruction Format: R2

#### REVBIT Rt, Ra – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
18 <sub>6</sub>	~	Vc	0	Vb	0	0 <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

Clock Cycles: 1

### Operation:

### Execution Units: I

### Clock Cycles: 1

### Exceptions: none

### Notes:

# SEQ – Set if Equal

## Description:

Compare two source operands for equality and place the result in the target register. The result is a Boolean true or false.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

## Operation:

$Rt = Ra == Rb$  or  $Rt = Ra == Imm$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

## Instruction Formats:

### SEQ Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
11 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

### SEQ Rt, Ra, Imm

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
11 <sub>6</sub>	0	Vc	0	Vb	~	~ <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						
Immediate <sub>31..0</sub>												0 <sub>3</sub>	31 <sub>5</sub>						

**Clock Cycles:** 1



# SGE – Set if Greater Than or Equal

## Description:

Compare two source operands for greater than or equal and place the result in the target register. The result is a Boolean true or false. This is the same instruction as [SLT](#) except that the result is inverted.

**Supported Operand Sizes:** .b, .w, .t, .o, .h, .d

## Operation:

$Rt = Ra < Rb$  or  $Rt = Ra < Imm$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

**Instruction Formats:**

**SGE Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
21 <sub>6</sub>	0	Vc	S	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

## SGT – Set if Greater Than

### Description:

Compare two source operands for greater than and place the result in the target register. The result is a Boolean true or false. This is the same instruction as [SLE](#) except that the result is complemented.

**Supported Operand Sizes:** .b, .w, .t, .o, .h, .d

### Operation:

$Rt = Ra > Rb$  or  $Rt = Ra > Imm$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

**Instruction Formats:**

**SLE Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
20 <sub>6</sub>	0	Vc	S	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

## SLE – Set if Less Than or Equal

### Description:

Compare two source operands for less than or equal and place the result in the target register. The result is a Boolean true or false.

**Supported Operand Sizes:** .b, .w, .t, .o, .h, .d

### Operation:

$Rt = Ra \leq Rb$  or  $Rt = Ra \leq Imm$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

### Instruction Formats:

**SLE Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
20 <sub>6</sub>	0	Vc	S	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	0	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

## SLT – Set if Less Than

### Description:

Compare two source operands for less than and place the result in the target register. The result is a Boolean true or false.

**Supported Operand Sizes:** .b, .w, .t, .o, .h, .d

### Operation:

$Rt = Ra < Rb$  or  $Rt = Ra < Imm$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

### Instruction Formats:

**SLT Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
21 <sub>6</sub>	0	Vc	S	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	0	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

## SNE – Set if Not Equal

### Description:

Compare two source operands for inequality and place the result in the target register. The result is a Boolean true or false.

**Supported Operand Sizes:** .b, .w, .t, .o, .h, .d

### Operation:

$Rt = Ra == Rb$  or  $Rt = Ra == Imm$

**Clock Cycles:** 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

### Instruction Formats:

**SNE Rt, Ra, Rb – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
11 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	1	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

# SQRT – Square Root

## Description:

This instruction computes the square root value of the contents of the source operand and places the result in Rt.

**Supported Operand Sizes:** .b, .w, .t, .o

**Integer Instruction Format: R2**

**SQRT Rt, Ra – Register direct**

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
15 <sub>6</sub>	~	Vc	0	Vb	0	0 <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

## Operation:

$Rt = \text{SQRT}(Ra)$

**Execution Units:** Integer ALU #0

**Clock Cycles: 1**

**Exceptions:** none

**Notes:**

# SUB - Subtraction

## Description:

Subtract two source operands and place the difference in the target register. All registers are treated as integer registers. Arithmetic is signed twos-complement values unless decimal mode is selected (SZ<sub>3</sub>=6) in which case values are treated as BCD numbers. This instruction may be used with the [CARRY](#) modifier to perform extended precision subtraction.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

## Operation:

$$Rt = Ra + -Rb \text{ or } Rt = Ra + -Imm$$

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

## Instruction Formats:

### SUB Rt, Ra, Rb – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
4 <sub>6</sub>	0	Vc	0	Vb	1	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### SUB Rt,Ra,Imm<sub>15</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	Vc	0	Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	4 <sub>5</sub>						

**Clock Cycles: 1**





# Vector Arithmetic Operations

Vector arithmetic operations are identical to scalar ones except that they may operate on vector registers. An extra register specification field may be present in the instruction to allow a mask register to be specified. Vector instructions are either 40 or 48 bits in length depending on the presence of a mask register.

The instruction is prefixed with the letter 'V' to indicate a vector form of the instruction. The assembler will recognize a vector instruction if a vector register is specified as one of the operands. The convention of prepending a 'V' to the instruction is a visual aid.

# VADD - Addition

## Description:

Add two source operands and place the sum in the target register. All registers are treated as integer registers. Arithmetic is signed twos-complement values unless the decimal mode flag is set in which case values are treated as densely packed BCD numbers. This instruction may be used with the [CARRY](#) modifier to perform extended precision addition. The following image shows the addition of vectors and how the vector mask register comes into play.

1	0	1	1	1	0	0	1	vmask
7	21	5	9	15	33	15435	6700	Source op A
+	+	+	+	+	+	+	+	
4	435	956	17	45	1021	876	52	Source op B
=	=	=	=	=	=	=	=	
11	12123	961	26	60	87	555	6752	Dest op T

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$$Rt = Ra + Rb \text{ or } Rt = Ra + Imm$$

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

## Instruction Formats:

**VADD Rt, Ra, Rb, Vm – Register direct**

47	46	45	40	39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
~	Svm	Vm <sub>6</sub>	4 <sub>6</sub>	~		0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	1	Sz <sub>3</sub>	2 <sub>5</sub>							

**Clock Cycles: 1**

**VADD Rt,Ra,Imm<sub>16</sub>,Vm**

47	46	45	40	39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
~	Svm	Vm <sub>6</sub>	Imm <sub>14..8</sub>	Vc	0	Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	1	Sz <sub>3</sub>	4 <sub>5</sub>							

**Clock Cycles: 1**

# VAND – Bitwise And

## Description:

Bitwise ‘and’ two source operands and place the result in the target register. The one’s complement of operands may be used by setting the appropriate ‘Sx’ bit in the instruction.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

**Clock Cycles:** 1

## Operation:

$Rt = Ra \& Rb$  or  $Rt = Ra \& Imm$

## Instruction Formats:

### VAND Rt, Ra, Rb – Register direct

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
8 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

### VAND Rt,Ra,Imm<sub>15</sub>

39		33	32	31	30		23	22	21	16	15	14	9	8	7	5	4	0
	Imm <sub>14..8</sub>	Vc	0		Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	8 <sub>5</sub>						

**Clock Cycles:** 1

**Execution Units:** All Integer ALU’s


**Exceptions:** none

**Notes:**

# VSHLV – Shift Vector Left

## Description

Elements of the vector are transferred upwards to the next element position. The first is loaded with the value zero. The highest element is lost. This is also called a slide operation. Elements may be moved a variable number of elements to the left. The image depicts just a single element shift.



## Instruction Formats:

### VSHLV Rt, Ra, Rb

39	35	3433	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
16 <sub>5</sub>	~ <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>						

Clock Cycles: 1

### VSHLV Rt, Ra, Imm<sub>7</sub>

39	35	3433	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
24 <sub>5</sub>	~ <sub>2</sub>	Vc	0	Vb	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>						

Clock Cycles: 1

## Operation

Amt = Rb

For x = VL-1 to Amt

$$Vt[x] = Va[x-amt]$$

For x = Amt-1 to 0


$$Vt[x] = 0$$

**Exceptions:** none

# VSHRV – Shift Vector Right

## Description

Elements of the vector are transferred downwards to the next element position. The last is loaded with the value zero. This is also called a slide operation. Elements may be moved a variable number of elements to the right. The image depicts just a single element shift.



## VSHRV Rt, Ra, Rb

39	35	3433	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
17 <sub>5</sub>	~ <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>						

Clock Cycles: 1

## VSHRV Rt, Ra, Imm<sub>7</sub>

39	35	3433	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
25 <sub>5</sub>	~ <sub>2</sub>	Vc	0	Vb	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>						

Clock Cycles: 1

## Operation

Amt = Rb

For x = 0 to VL-Amt

$$Vt[x] = Va[x+amt]$$

For x = VL-Amt + 1 to VL-1

$$Vt[x] = 0$$

**Exceptions:** none

# Floating-Point Operations

## Precision

Floating point operations are always performed at the greatest precision available. Lower precision formats are available for storage.

For decimal floating-point three storage formats are supported. 96-bit triple precision, 64-bit double precision, and 32-bit single precision values.


## Representations

### Binary Floats


Triple Precision, Float:128

The core uses a 128-bit quad precision binary floating-point representation.


Quad Precision, long double




Double Precision, double



Single Precision, float



Half Precision, short float



## Decimal Floats

The core uses a 128-bit densely packed decimal triple precision floating-point representation.

127	126	122	121	110	109	0
S	Combo <sub>5</sub>	Exponent <sub>12</sub>	Significand <sub>110</sub>			

The significand stores 34 densely packed decimal digits. One whole digit before the decimal point.

The exponent is a power of ten as a binary number with an offset of 1535. Range is  $10^{-1535}$  to  $10^{1536}$

64-bit double precision decimal floating point:

63	62	58	57	50	49	0
S	Combo <sub>5</sub>	Exponent <sub>8</sub>	Significand <sub>50</sub>			

The significand stores 16 DPD digits. One whole digit before the decimal point.

32-bit single precision decimal floating point:

31	30	26	25	20	19	0
S	Combo <sub>5</sub>	Exponent <sub>6</sub>	Significand <sub>20</sub>			

The significand store 7 DPD digits. One whole digit before the decimal point.

## Rounding Modes

### Binary Float Rounding Modes

Rm3	Rounding Mode
000	Round to nearest ties to even
001	Round to zero (truncate)
010	Round towards plus infinity
011	Round towards minus infinity
100	Round to nearest ties away from zero
101	Reserved
110	Reserved
111	Use rounding mode in float control register

### Decimal Float Rounding Modes

Rm3	Rounding Mode
000	Round ceiling
001	Round floor
010	Round half up
011	Round half even
100	Round down
101	Reserved

110	Reserved
111	Use rounding mode in float control register

## Operand Sizes

Sz <sub>3</sub>	Ext.	Operand
0		Reserved
1	.h	16-bit half
2	.s	32-bit single
3	.d	64-bit double
4	.q	128-bit quad
5		reserved
6		128-bit decimal
7		reserved



## FABS – Absolute Value

### Description:

This instruction computes the absolute value of the contents of the source operand and places the result in Rt. The sign bit of the value is cleared. No rounding occurs.

### Integer Instruction Format: R1

#### FABS Rt, Ra, Rb – Register direct

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
1 <sub>5</sub>	~ <sub>3</sub>	Vc	Vb	0	32 <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 1**

### Operation:

$$FPt = \text{Abs}(FPa)$$

**Execution Units:** FPU #0

**Clock Cycles:** 1

**Exceptions:** none

**Notes:**

# FADD –Float Addition

## Description:

Add two source operands and place the sum in the target register. All registers values are treated as quad precision floating-point values. An immediate value is converted to quad precision value from half, single, or double precision.

## Supported Operand Sizes:

## Operation:

$$Rt = Ra + Rb \text{ or } Rt = Ra + Imm$$

## Clock Cycles: 8

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

## Instruction Formats:

### FADD Rt, Ra, Rb – Register direct

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
4 <sub>5</sub>	Rm <sub>3</sub>	Vc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

### FADD Rt,Ra,Imm<sub>16</sub>

39	32	31	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>15..9</sub>	Vc	Imm <sub>8..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	20 <sub>5</sub>					

### FADD Rt,Ra,Imm<sub>32</sub>

39	33	32	31	23	22	21	16	15	14	9	8	7	5	4	0
~7	Vc	~9	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	20 <sub>5</sub>						
Immediate <sub>32</sub>												0 <sub>3</sub>	31 <sub>5</sub>		

### FADD Rt,Ra,Imm<sub>64</sub>

39	32	31	23	22	21	16	15	14	9	8	7	5	4	0
~7	Vc	~9	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	20 <sub>5</sub>					
Immediate <sub>31..0</sub>									0 <sub>3</sub>	31 <sub>5</sub>				
Immediate <sub>63..32</sub>									1 <sub>3</sub>	31 <sub>5</sub>				

### FADD Rt,Ra,Imm<sub>128</sub>

39	32	31	31	23	22	21	16	15	14	9	8	7	5	4	0
~7	Vc	~9	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	20 <sub>5</sub>						
Immediate <sub>31..0</sub>									0 <sub>3</sub>	31 <sub>5</sub>					
Immediate <sub>63..32</sub>									1 <sub>3</sub>	31 <sub>5</sub>					
Immediate <sub>95..64</sub>									2 <sub>3</sub>	31 <sub>5</sub>					
Immediate <sub>127..96</sub>									2 <sub>3</sub>	31 <sub>5</sub>					

# FCMP - Comparison

## Description:

Compare two source operands and place the result in the target register. The result is a vector identifying the relationship between the two source operands as floating-point values. This instruction may compare against lower precision immediate values to conserve code space. Note that result inversion is not available as results are already available in normal and inverted forms in the bit vector. Instead the 'S' bit indicates to swap the operands.

## Supported Operand Sizes:

## Operation:

$R_t = R_a \text{ ? } R_b$  or  $R_t = R_a \text{ ? } Imm$  or  $R_t = Imm \text{ ? } R_a$

## Clock Cycles: 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

## Instruction Formats:

### FCMP Rt, Ra, Rb – Register direct

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
5 <sub>5</sub>	0 <sub>3</sub>	V <sub>c</sub>	V <sub>b</sub>	S <sub>b</sub>	R <sub>b</sub> <sub>6</sub>	S <sub>a</sub>	R <sub>a</sub> <sub>6</sub>	~	R <sub>t</sub> <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 1**

### FCMP Rt,Ra,Imm<sub>16</sub>

39	33	32	31	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>15..9</sub>	V <sub>c</sub>	Imm <sub>8..0</sub>	S <sub>a</sub>	R <sub>a</sub> <sub>6</sub>	S	R <sub>t</sub> <sub>6</sub>	V	Sz <sub>3</sub>	21 <sub>5</sub>						

**Clock Cycles: 1**

Rt bit	Mnem.	Meaning	Test
<b>Float Compare Results</b>			
0	EQ	equal	!nan & eq
1	NE	not equal	!eq
2	GT	greater than	!nan & !eq & !lt & !inf
3	UGT	Unordered or greater than	Nan    (!eq & !lt & !inf)
4	GE	greater than or equal	Eq    (!nan & !lt & !inf)
5	UGE	Unordered or greater than or equal	Nan    (!lt    eq)
6	LT	Less than	Lt & (!nan & !inf & !eq)
7	ULT	Unordered or less than	Nan   (!eq & lt)
8	LE	Less than or equal	Eq   (lt & !nan)
9	ULE	unordered less than or equal	Nan   (eq   lt)
10	GL	Greater than or less than	!nan & (!eq & !inf)
11	UGL	Unordered or greater than or less than	Nan   !eq

---

12	ORD	Greater than less than or equal / ordered	!nan
13	UN	Unordered	Nan
14		Reserved	
15		reserved	

## FDIV –Float Division

### Description:

Divide two source operands and place the quotient in the target register. All registers values are treated as 96-bit floating-point values.

### Supported Operand Sizes:

### Operation:

$Rt = Ra / Rb$  or  $Rt = Ra / Imm$  or  $Rt = Imm / Ra$

### Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

### Notes:

### Instruction Formats:

#### FDIV Rt, Ra, Rb – Register direct

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
7 <sub>5</sub>	Rm <sub>3</sub>	Vc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 150**

#### FDIV Rt,Ra,Imm<sub>16</sub>

39	33	32	31	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>15..9</sub>	Vc	Imm <sub>8..0</sub>	Sa	Ra <sub>6</sub>	S	Rt <sub>6</sub>	V	Sz <sub>3</sub>	23 <sub>5</sub>						

**Clock Cycles: 150**

## FSEQ – Float Set if Equal

## FSNE – Float Set if Not Equal

### Description:

Compares two source operands for equality and places the result in the target register. The result is a Boolean true or false. Positive and negative zero are considered equal. This instruction does not support a 16-bit immediate. 32, 64, and 128-bit immediates are supported. For FSEQ is either operand is a NaN zero the result is zero.

### Supported Operand Sizes:

### Operation:

$$Rt = Ra == Rb \text{ or } Rt = Ra == \text{Imm}$$

### Clock Cycles: 1

**Execution Units:** All FPU's

**Exceptions:** none

### Notes:

### Instruction Formats:

#### FSEQ Rt, Ra, Rb

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
8	0 <sub>3</sub>	V <sub>C</sub>	V <sub>b</sub>	S <sub>b</sub>	R <sub>b</sub> <sub>6</sub>	S <sub>a</sub>	R <sub>a</sub> <sub>6</sub>	S <sub>t</sub>	R <sub>t</sub> <sub>6</sub>	V	S <sub>z</sub> <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 1**

#### FSNE Rt, Ra, Rb

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
8 <sub>5</sub>	0 <sub>3</sub>	V <sub>C</sub>	V <sub>b</sub>	S <sub>b</sub>	R <sub>b</sub> <sub>6</sub>	S <sub>a</sub>	R <sub>a</sub> <sub>6</sub>	1	R <sub>t</sub> <sub>6</sub>	V	S <sub>z</sub> <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 1**

# FSGE – Float Set if Greater Than or Equal

## Description:

Compares two source operands for greater than or equal and places the result in the target register. The result is a Boolean true or false.

## Supported Operand Sizes:

## Operation:

$Rt = Ra \geq Rb$  or  $Rt = Ra \geq Imm$

**Clock Cycles:** 1

**Execution Units:** All FPU's

**Exceptions:** none

**Notes:**

## Instruction Formats:

### FSGE Rt, Ra, Rb

39 35	34 32	31	30	29	28 23	22	21 16	15	14	9	8	7 5	4 0
8 <sub>5</sub>	4 <sub>3</sub>	V <sub>C</sub>	V <sub>b</sub>	S <sub>b</sub>	R <sub>b</sub> <sub>6</sub>	S <sub>a</sub>	R <sub>a</sub> <sub>6</sub>	St	R <sub>t</sub> <sub>6</sub>	V	S <sub>Z</sub> <sub>3</sub>	12 <sub>5</sub>	

# FSGT – Float Set if Greater Than

## Description:

Compares two source operands for greater than and places the result in the target register. The result is a Boolean true or false.

## Supported Operand Sizes:

## Operation:

$Rt = Ra > Rb$  or  $Rt = Ra > Imm$

## Clock Cycles: 1

## Execution Units: All FPU's

## Exceptions: none

## Notes:

## Instruction Formats:

### FSLT Rt, Ra, Rb

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
8 <sub>5</sub>	5 <sub>3</sub>	V <sub>C</sub>	V <sub>b</sub>	S <sub>b</sub>	R <sub>b</sub> <sub>6</sub>	S <sub>a</sub>	R <sub>a</sub> <sub>6</sub>	S <sub>t</sub>	R <sub>t</sub> <sub>6</sub>	V	S <sub>Z</sub> <sub>3</sub>	12 <sub>5</sub>							

## Clock Cycles: 1





# FSLE – Float Set if Less Than or Equal

## Description:

Compares two source operands for less than or equal and places the result in the target register.  
The result is a Boolean true or false.

## Supported Operand Sizes:

## Operation:

$Rt = Ra \leq Rb$  or  $Rt = Ra \leq Imm$  or  $Rt = Imm \leq Ra$

## Clock Cycles: 1

## Execution Units: All FPU's

## Exceptions: none

## Notes:

## Instruction Formats:

## FSLE Rt, Ra, Rb

39 35	34 32	31	30	29	28 23	22	21 16	15	14	9	8	7 5	4 0
8 <sub>5</sub>	2 <sub>3</sub>	Vc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>	

# FSLT – Float Set if Less Than

## Description:

Compares two source operands for less than and places the result in the target register. The result is a Boolean true or false.

## Supported Operand Sizes:

## Operation:

$$Rt = Ra < Rb \text{ or } Rt = Ra < \text{Imm} \text{ or } Rt = \text{Imm} < Ra$$

## Clock Cycles: 1

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

## Instruction Formats:

### FSLT Rt, Ra, Rb

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
8 <sub>5</sub>	3 <sub>3</sub>	V <sub>C</sub>	V <sub>b</sub>	S <sub>b</sub>	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 1**

# FMUL –Float Multiplication

## Description:

Multiply two source operands and place the product in the target register. All registers values are treated as 128-bit floating-point values.

## Supported Operand Sizes:

## Operation:

$$Rt = Ra * Rb \text{ or } Rt = Ra * Imm$$

## Clock Cycles:

**Execution Units:** All FPU's

**Exceptions:** none

## Notes:

## Instruction Formats:

### FMUL Rt, Ra, Rb

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
6 <sub>5</sub>	Rm <sub>3</sub>	Vc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 8**

### FMUL Rt,Ra,Imm<sub>13</sub>

39	33	32	31	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>15..9</sub>	Vc	Imm <sub>8..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	22 <sub>5</sub>						

**Clock Cycles: 8**

# FNEG – Negate Value

## Description:

This instruction computes the negative value of the contents of the source operand and places the result in Rt. The sign bit of the value is inverted. No rounding occurs. Note that in most cases a FNEG operation can be absorbed into a previous instruction by negating the result.

## Integer Instruction Format: R1

### FNEG Rt, Ra

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
1 <sub>5</sub>	~ <sub>3</sub>	Vc	Vb	0	34 <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 1**

## Operation:

$$Rt = -Ra$$

**Execution Units:** FPU #0

**Clock Cycles:** 1

**Exceptions:** none

**Notes:**

## FSCALEB –Scale Exponent

# FSUB –Float Subtraction

## Description:

Subtract two source operands and place the difference in the target register. All registers values are treated as 128-bit floating-point values. This is an alternate mnemonic for the [FADD](#) instruction where the second source operand, Rb is assumed negated.

## Supported Operand Sizes:

## Operation:

$$Rt = Ra + -Rb \text{ or } Rt = Ra + -Imm$$

## Clock Cycles: 8

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

## Instruction Formats:

### FSUB Rt, Ra, Rb – Register direct

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
4 <sub>5</sub>	Rm <sub>3</sub>	Vc	Vb	1	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

### FSUB Rt,Ra,Imm<sub>16</sub>

39	32	31	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>15..9</sub>	Vc	Imm <sub>8..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	20 <sub>5</sub>					

**Clock Cycles: 8**

# FTRUNC – Truncate Fraction

## Description:

This instruction truncates off the fractional portion of the number leaving only the integer portion.  
No rounding occurs.

## Integer Instruction Format: R1

### FTRUNC Rt, Ra

39	35	34	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
1 <sub>5</sub>	~ <sub>3</sub>	Vc	Vb	0	21 <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	12 <sub>5</sub>							

**Clock Cycles: 1**

## Operation:

$$Rt = \text{Trunc}(Ra)$$

**Execution Units:** FPU #0

**Clock Cycles:** 1

**Exceptions:** none

**Notes:**



## ORF – Bitwise Or to Float

### Description:

Convert the immediate constant to quad precision format and bitwise ‘or’ with source operand Ra and place the result in the target register. The immediate constant may be a half, single, double, or quad precision value. This instruction is provided mainly for loading a floating-point value into a register. The value may be compressed into the minimum size format for representation without loss of precision. [FADD](#) could also be used to load a float constant into a register but it has a longer latency.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

### Operation:

$$Rt = Ra \mid \text{Convert(Imm)}$$

### Instruction Formats:

**ORF Rt,Ra,Imm<sub>15</sub>**

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>				Vc	1	Imm <sub>7..0</sub>				Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	9 <sub>5</sub>

**Clock Cycles: 1**

**Clock Cycles: 1**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# String Operations

## Representations

### Strings

95	92	91	64	63	0
Typ	Length <sub>28</sub>	Pointer <sub>64</sub>			

### UTF8 Chars

127	0
16 characters	

### UTF32 Chars

127	0
4 characters	

# CHRNDX – Character Index

## Description:

This instruction searches Ra, which is treated as an array of characters, for a character value specified by Rb and places the index of the character into the target register Rt. If the character is not found -1 is placed in the target register. A common use would be to search for a null byte. The index result may vary from -1 to +15 for UTF8 characters or -1 to +3 for UTF32 characters. The index of the first found byte is returned (closest to zero).

**Supported Operand Sizes:** .b, .t

## Instruction Formats:

### CHRNDX Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
13 <sub>6</sub>	0	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### CHRNDX Rt, Ra, Imm

39	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
13 <sub>6</sub>	1	Vc	0		Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>					

**Clock Cycles: 1**

## Operation:

Rt = Index of (Rb in Ra)

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# Block Instructions

## BCMP – Block Compare

### Description:

This instruction compares data from the memory location addressed by Ra to the memory location addressed by Rb until the loop counter LC reaches zero or until a mismatch occurs. Ra and Rb increment by independently specified amounts. This instruction is interruptible.

### Instruction Format: RINDL

### BCMP [Ra], [Rb]

39	35	3433	32	31	30	29	28	23	22	21	16	15	1413	12	11	10	9	8	7	5	4	0
0 <sub>5</sub>	Cb <sub>2</sub>	Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	~	Ca <sub>2</sub>	Ai <sub>2</sub>	Bi <sub>2</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>							

### Clock Cycles:

Bi <sub>2</sub>	Rb	Ai <sub>2</sub>	Ra
0	No Change	0	No change
1	Increment by Sz <sub>3</sub>	1	Increment by Sz <sub>3</sub>
2	Decrement by Sz <sub>3</sub>	2	Decrement by Sz <sub>3</sub>
3	reserved	3	reserved

### Assembler Example

```
LDI LC,200
BCMP.O [Ra]+,[Rb]+
SUBF LC,LC,200      ; get index of difference
```

### Execution Units: Memory

### Operation:

```
temp = 0
while LC <> 0 and mem[Rb] = mem[Ra]
    Ra = Ra + amt
    Rb = Rb + amt
    LC = LC - 1
```

# BFND – Block Find

## Description:

This instruction compares data from the memory location addressed by Rb to the data in register Ra until the loop counter LC reaches zero or until a match occurs. This instruction is interruptible.

## Instruction Format: RINDL

### BFND Ra, [Rb], Imm

39	35	3433	32	31	30	29	28	23	22	21	16	15	11	10	9	8	7	5	4	0
1 <sub>5</sub>	Cb <sub>2</sub>	Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	~ <sub>5</sub>	Bi <sub>2</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>							

## Clock Cycles:

Bi <sub>2</sub>	Rb
0	No Change
1	Increment by Sz <sub>3</sub>
2	Decrement by Sz <sub>3</sub>
3	reserved

## Execution Units: Memory

## Operation:

```
temp = 0
while LC <> 0
    if (mem[Rb] = Ra)
        stop
    Rb = Rb + amt
    LC = LC - 1
```

# BMOV –Block Move

## Description:

This instruction moves a data from the memory location addressed by Ra to the memory location addressed by Rb until the loop counter LC reaches zero. Ra and Rb are adjusted by a specified amount after the move. This instruction is interruptible.

This instruction may be used to stream data to or from an I/O port by keeping one of the addresses fixed.

## Instruction Format: RINDL

### BMOV [Ra], [Rb], Imm

39	35	3433	32	31	30	29	28	23	22	21	16	15	1413	1211	10	9	8	7	5	4	0
2 <sub>5</sub>	Cb <sub>2</sub>	Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	~	Ca <sub>2</sub>	Ai <sub>2</sub>	Bi <sub>2</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

**Ca2** = load cache-ability

**Cb2** = store cache-ability

Bi <sub>2</sub>	Rb	Ai <sub>2</sub>	Ra
0	No Change	0	No change
1	Increment by Sz <sub>3</sub>	1	Increment by Sz <sub>3</sub>
2	Decrement by Sz <sub>3</sub>	2	Decrement by Sz <sub>3</sub>
3	reserved	3	reserved

## Assembler Example

```
LDI LC,200
BMOVB [Ra++],[Rb++]
```

## Execution Units: Memory

## Operation:

```
temp = 0
while LC <> 0
    t0 = mem[Ra]
    mem[Rb] = t0
    Ra = Ra + amt
    Rb = Rb + amt
    LC = LC - 1
```



# BSET – Block Set

## Description:

This instruction stores data contained in register Ra to consecutive memory locations beginning at the address in Rb until the loop counter reaches zero. Rb is updated by the number of bytes written.

## Instruction Format: RINDL

### BSET Ra, [Rb], Imm

39	35	3433	32	31	30	29	28	23	22	21	16	1513	12	9	8	7	5	4	0
3 <sub>5</sub>	Ca <sub>2</sub>	Vc	1	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	~ <sub>3</sub>	Bi <sub>4</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

## Clock Cycles:

Bi <sub>4</sub>	Adjustment Amount
0	0
1	1
2	2
3	4
4	8
5	16
15	-1
14	-2
13	-4
12	-8
11	-16
others	reserved

## Execution Units: Memory

## Operation:

```

if LC <> 0
    mem[Rb] = Ra
    Rb = Rb + amt
    LC = LC - 1

```

## Assembler Example

```

LDI LC,200
BSETB Ra,[Rb],1

```





# Bit Manipulation Operations

Many CPUs do not have direct support for bit-field manipulation. Instead, they rely on ordinary logical and shift operations. The benefit of having bit-field operations is that they are more code dense than performing the operations using other ALU ops.

Bitfield operations repurpose the size field for use as an opcode extension.

The beginning and end of a bitfield may be specified as either a pair of immediate constants or in a pair of registers.

# CLR – Clear Bit Field

## Description:

A bit field in the source operand is cleared and the result placed in the target register. The specified bit to clear is modulo the operand size.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

**Flag Updates:** none

## Operation:

$Rt = Ra \ \& \sim \text{bit } Rb$  or  $Ra = Ra \ \& \sim \text{bit } imm$

## Instruction Formats:

### CLR Rt, Ra, Rb, Rc

39	38	37	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Rc <sub>6</sub>	~	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	0 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles: 4**

### CLR Rt, Ra, Offs<sub>7</sub>, Wid<sub>7</sub>

39	38	37	36	30	29	23	22	21	16	15	14	9	8	7	5	4	0
1	Vc	~	end <sub>7</sub>	begin <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	0 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles: 4**

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

# COM – Complement Bit Field

## Description:

A bit-field in the source operand is changed and placed in the target register. The specified bit to change is modulo the operand size.

**Supported Operand Sizes:** .b, .w, .t, .o

**Flag Updates:** none

## Operation:

$$Rt[Rb] = \sim Ra[Rb] \text{ or } Rt[Imm] = \sim Ra[Imm]$$

## Instruction Formats:

### Instruction Formats:

#### CLR Rt, Ra, Rb

39	38	37	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Rc <sub>6</sub>	~	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	2 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles: 4**

#### CLR Rt, Ra, Offs<sub>7</sub>, Wid<sub>7</sub>

39	38	37	36	30	29	23	22	21	16	15	14	9	8	7	5	4	0
1	Vc	~	end <sub>7</sub>	begin <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	2 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles: 4**

**Clock Cycles: 1**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# DEP – Deposit Bit Field

## Description:

A source operand is transferred to a bitfield in the target register.

**Supported Operand Sizes:** .b, .w, .t, .o

**Flag Updates:** none

## Operation:

MB = offset

ME = offset + width

Rt[ME:MB] = Ra

## Instruction Formats:

### DEP Rt, Ra, Rb, Rc

39	38	37	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Rc <sub>6</sub>	~	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	6 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles:** 1

### DEP Rt, Ra, Offs<sub>7</sub>, Wid<sub>7</sub>

39	38	37	36	30	29	23	22	21	16	15	14	9	8	7	5	4	0
1	Vc	~	end <sub>7</sub>	begin <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	6 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles:** 1

**Clock Cycles:**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# EXTS – Extract Signed Bit Field

## Description:

Extract a bit field from the source operand and place the bit field in the target register. The field is sign extended.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

## Operation:

$R_t = R_a[R_b]$  or  $R_t = R_a[Imm]$

## Instruction Formats:

### EXTS $R_t, R_a, R_b$

39	38	37	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Rc <sub>6</sub>	~	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	5 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles: 1**

### EXTS $R_t, R_a, Offs_7, Wid_7$

39	38	37	36	30	29	23	22	21	16	15	14	9	8	7	5	4	0
1	Vc	~	end <sub>7</sub>	begin <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	5 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles: 1**

**Clock Cycles: 1**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# EXTU – Extract Bit Field

## Description:

Extract a bit field from the source operand and place the bit field in the target register. The field is zero extended.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$R_t = R_a[R_b]$  or  $R_t = R_a[Imm]$

## Instruction Formats:

### EXTU Rt, Ra, Rb, Rc

39	38	37	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Rc <sub>6</sub>	~	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	4 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles: 1**

### EXTU Rt, Ra, Offs<sub>7</sub>, Wid<sub>7</sub>

39	38	37	36	30	29	23	22	21	16	15	14	9	8	7	5	4	0
1	Vc	~	end <sub>7</sub>	begin <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	4 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles: 1**

**Clock Cycles: 1**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

## SBX – Sign Bit Extend

### Description:

Sign extend a value beginning at a specified bit to the width specified and place the result in the target register. All registers are integer registers.

**Supported Operand Sizes:** .b, .w, .t, .o

### Operation:

### Instruction Formats:

#### SBX Rt, Ra, Rb, Rc

39	38	37	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Rc <sub>6</sub>	~	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	3 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles:** 1

#### SBX Rt, Ra, Offs<sub>7</sub>, Wid<sub>7</sub>

39	38	37	36	30	29	23	22	21	16	15	14	9	8	7	5	4	0
1	Vc	~	end <sub>7</sub>	begin <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	3 <sub>3</sub>	13 <sub>5</sub>						

**Clock Cycles:** 1

### Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

### Notes:



# SET – Set Bit Field

## Description:

A bit in the source operand is set and placed in the target register.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$Rt = Ra \mid \text{bit } Rb \text{ or } Rt = Ra \text{ or Bit}[\text{Imm}]$

## Instruction Formats:

### SET Rt, Ra, Rb, Rc

39	38	37	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Rc <sub>6</sub>	~	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	l <sub>3</sub>	l <sub>5</sub>						

**Clock Cycles: 1**

### SET Rt, Ra, Offs<sub>7</sub>, Wid<sub>7</sub>

39	38	37	36	30	29	23	22	21	16	15	14	9	8	7	5	4	0
1	Vc	~	end <sub>7</sub>	begin <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	l <sub>3</sub>	l <sub>5</sub>						

**Clock Cycles: 1**

**Clock Cycles: 1**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# Shift and Rotate Operations

Shift instructions can take the place of some multiplication and division instructions. Some architectures provide shifts that shift only by a single bit. Others use counted shifts, the original 80x88 used multiple clock cycles to shift by an amount stored in the CX register. Table 8.88 and Thor use a barrel shifter to allow shifting by an arbitrary amount in a single clock cycle. Shifts are infrequently used, and a barrel (or funnel) shifter is relatively expensive in terms of hardware resources.

# ASL – Arithmetic Shift Left

## Description:

Shift the first source operand to the left by the number of bits specified by the second source operand and place the result in the target register. All registers are integer registers. Arithmetic is signed twos-complement values. The least significant bit is filled with the value of 'N' specified in the instruction.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$R_t = R_a \ll R_b$  or  $R_t = R_a \ll \text{Imm}$

## Instruction Formats:

### ASL Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
0 <sub>5</sub>	~ <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### ASL Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
8 <sub>5</sub>	~ <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

**Clock Cycles:**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# ASR – Arithmetic Shift Right

## Description:

Shift the first source operand to the right, preserving the sign bit, by the number of bits specified by the second source operand and place the result in the target register. All registers are integer registers. Arithmetic is signed twos-complement values.

**Supported Operand Sizes:** .b, .w, .l

## Operation:

$Rt = Ra \gg Rb$  or  $Rt = Ra \gg Imm$

## Instruction Formats:

### ASR Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
l <sub>5</sub>	~ <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### ASR Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
9 <sub>5</sub>	~ <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

# LSL – Logical Shift Left

## Description:

Shift the first source operand to the left by the number of bits specified by the second source operand and place the result in the target register. All registers are integer registers. Arithmetic is signed twos-complement values. Fill the least significant bit with the value specified by ‘N’ in the instruction.

This instruction may be used to generate a bitmask by setting N to one, and shifting a zero.

**Supported Operand Sizes:** .b, .w, .l

## Operation:

$R_t = R_a \ll R_b$  or  $R_t = R_a \ll \text{Imm}$

## Instruction Formats:

### LSL Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
2 <sub>5</sub>	0 <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### LSL Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
10 <sub>5</sub>	0 <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

# LSLAND – Logical Shift Left and And

## Description:

Shift the first source operand to the left by the number of bits specified by the second source operand and bitwise ‘and’ the result to the target register. All registers are integer registers. Arithmetic is signed twos-complement values. Fill the least significant bit with the value specified by ‘N’ in the instruction.

This instruction may be used to isolate a bitfield in a target register.

**Supported Operand Sizes:** .b, .w, .l

## Operation:

$Rt = Rt \& (Ra \ll Rb)$  or  $Rt = Rt \& (Ra \ll Imm)$

## Instruction Formats:

### LSL.AND Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
2 <sub>5</sub>	1 <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### LSL.AND Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
10 <sub>5</sub>	1 <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

# LSLOR – Logical Shift Left and Or

## Description:

Shift the first source operand to the left by the number of bits specified by the second source operand and bitwise ‘or’ the result to the target register. All registers are integer registers. Arithmetic is signed twos-complement values. Fill the least significant bit with the value specified by ‘N’ in the instruction.

This instruction may be used to insert a bitfield into a target register.

**Supported Operand Sizes:** .b, .w, .l

## Operation:

$$Rt = Rt \mid (Ra \ll Rb) \text{ or } Rt = Rt \mid (Ra \ll Imm)$$

## Instruction Formats:

### LSLOR Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
2 <sub>5</sub>	2 <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### LSLOR Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
10 <sub>5</sub>	2 <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

# LSLXOR – Logical Shift Left and Exclusive Or

## Description:

Shift the first source operand to the left by the number of bits specified by the second source operand and bitwise exclusive 'or' the result to the target register. All registers are integer registers. Arithmetic is signed twos-complement values. Fill the least significant bit with the value specified by 'N' in the instruction.

This instruction may be used to insert or invert a bitfield in a target register.

**Supported Operand Sizes:** .b, .w, .l

## Operation:

$$Rt = Rt \wedge (Ra \ll Rb) \text{ or } Rt = Rt \wedge (Ra \ll Imm)$$

## Instruction Formats:

### LSLXOR Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
2 <sub>5</sub>	3 <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### LSLXOR Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
10 <sub>5</sub>	3 <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:



# LSR – Logical Shift Right

## Description:

Shift the first source operand to the right by the number of bits specified by the second source operand and place the result in the target register. All registers are integer registers. Arithmetic is signed twos-complement values. Fill the least significant bit with the value specified by 'N' in the instruction.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$R_t = R_a \gg R_b$  or  $R_t = R_a \gg \text{Imm}$

## Instruction Formats:

### LSL Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
3 <sub>5</sub>	0 <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### LSL Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
11 <sub>5</sub>	0 <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

**Clock Cycles:**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# ROL – Rotate Left

## Description:

Rotate the first source operand to the left by the number of bits specified by the second source operand and place the result in the target register. All registers are integer registers. Arithmetic is signed twos-complement values. The least significant bit is set to the value of the most significant bit exclusively or'd with the value 'N' from the instruction.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$R_t = R_a \ll R_b$  or  $R_t = R_a \ll \text{Imm}$

## Instruction Formats:

### ROL Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
4 <sub>5</sub>	0 <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### ROL Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
12 <sub>5</sub>	0 <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

**Clock Cycles:**

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

# ROR – Rotate Right

## Description:

Rotate the first source operand through the carry to the right by the number of bits specified by the second source operand and place the result in the target register. All registers are integer registers. Arithmetic is signed twos-complement values. The most significant bit is set to the value of the least significant bit exclusively or'd with the value 'N' from the instruction.

**Supported Operand Sizes:** .b, .w, .l

## Operation:

$Rt = Ra \gg Rb$  or  $Rt = Ra \gg Imm$

## Instruction Formats:

### ROR Rt, Ra, Rb

39	34	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
5 <sub>5</sub>	0 <sub>2</sub>	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

### ROR Rt, Ra, Imm<sub>7</sub>

39	34	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
13 <sub>5</sub>	0 <sub>2</sub>	Vc	0	0	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>							

**Clock Cycles: 1**

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

# Flow Control Instructions

## The Branch Set

One of the first things the author looks at when evaluating an ISA is the branch set. Is it semi-sensible or non-sense ? Branches may represent up to one quarter of instruction executed. Branches are one item that should be well done in an architecture. What conditions will the processor branch on ? Is it a simple branch on zero / non-zero test or are there more complex conditions available ? What the branch set supports impacts what other instructions need to be available in the architecture. If branching only supports a zero / non-zero test, then other instructions must be present to setup the branch test. In the DLX architecture for instance, there are a set of 'set' instructions that set a register to a one or zero based on a condition. After a set instruction is done, then a conditional branch may occur. Many architectures include a compare instruction(s). For instance, the MMIX architecture includes both signed (CMP) and unsigned compare (CMPU) instructions that set the value of a register to -1, 0, or 1 for less than, equal, or greater than another register. The same paradigm was used for the Raptor64 processor. For the Thor processor there is a fairly standard set of branches. Because the instruction set is wide enough, two registers may be compared during the branch.

## Mnemonics

There are mnemonics for specifying the comparison method. Floating-point comparisons prefix the branch mnemonic with 'F' as in FBEQ. Decimal-floating point comparisons prefix the branch mnemonic with 'DF' as in DFBEQ. And finally posit comparisons prefix the branch mnemonic with a 'P' as in 'PBEQ'.

## Conditions

Conditional branches branch to the target address only if the condition is true. The condition is determined by the comparison of two general-purpose registers or the comparison of a general-purpose register and an immediate value.

*The original Thor machine used instruction predicates to implement conditional branching. Another instruction was required to set the predicate before branching. Combining compare and branch in a single instruction may reduce the dynamic instruction count. An issue with comparing and branching in a single instruction is that it may lead to a wider instruction format.*

The comparison used is determined by a three-bit field in the instruction. There are four comparison types that may be performed as outlined in the table below.

Cm <sub>2</sub>	Comparison Type
0	integer comparisons
1	posit comparison
2	quad float comparison
3	quad decimal float comparison

## Conditional Branch Format

Branches use 40 or 48-bit opcodes.

*A 32-bit opcode would not leave a large enough target field and would end up using two or more instructions to implement most branches. With the prospect of using two instructions to perform compare then branches as many architectures do, it is more space efficient to simply use a wider instruction format.*

39		23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>			Cm <sub>2</sub>	Rn <sub>6</sub>		Rm <sub>6</sub>		Cond <sub>4</sub>		27 <sub>5</sub>	

47		23	2221	20	15	14	9	8	5	4	0
Disp <sub>24..0</sub>			Cm <sub>2</sub>	Rn <sub>6</sub>		Rm <sub>6</sub>		Cond <sub>4</sub>		28 <sub>5</sub>	

## Branch Conditions

The branch condition field combined with the opcode and comparison method determines the condition under which the branch will execute.

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>		Cm <sub>2</sub>	Rn <sub>6</sub>		Rm <sub>6</sub>		Cond <sub>4</sub>		27 <sub>5</sub>	

0h	Equal	Equal	Equal
1h	Not equal	Not equal	Not equal
2h	Signed less than	Greater than	less than
3h	Signed less than or equal	Unordered greater than	Less than or equal
4h	Signed greater than or equal	Greater than or equal	Greater than or equal
5h	Signed greater than	Unordered greater than or equal	Greater than
6h	Bit clear	Less than	
7h	Bit set	Unordered less than	
8h	Bit clear immediate	Less than or equal	
9h	Bit set immediate	Unordered less than or equal	
Ah	Unsigned less than	Greater than or less than	
Bh	Unsigned less than or equal	Unordered greater than or less than	
Ch	Unsigned greater than or eq	Greater than, less than, or equal	
Dh	Unsigned greater than	Unordered	
Eh	Branch always		
Fh	Branch Subroutine		

## Linkage

Unconditional branches may specify a linkage register which is updated with the address of the next instruction. This allows position independent subroutines to be called. There are four link registers in the architecture.

39	23	22	11	109	8	5	4	0
Disp <sub>16..0</sub>		Disp <sub>28..17</sub>		Rt <sub>2</sub>	15 <sub>4</sub>		27 <sub>5</sub>	

Rt <sub>2</sub>	Meaning
0	Use Register 56
1	Use Register 57

2	Use Register 58
3	Use Register 59

## Branch Target

For conditional branches, the target address is formed as the sum of the instruction pointer and a 17 or 25-bit constant specified in the instruction. Branches are PC relative with a range of  $\pm 64\text{kB}$  or  $\pm 16\text{MB}$ .

*The target displacement field is recommended to be at least 16-bits. It is possible to get by with a displacement as small as 12-bits before a significant percentage of branches must be implemented as two or more instructions.*



39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>		Cm <sub>2</sub>	Rn <sub>6</sub>		Rm <sub>6</sub>		Cond <sub>4</sub>		27 <sub>5</sub>	

## Bcc – Conditional Branch

Bcc Rm, Rn, label

### Description:

Branch if the condition is met. The condition is a relationship between either two registers or a register and an immediate value. The displacement is relative to the address of the branch instruction. The branch range is +/- 64kB.

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16,0</sub>				Cm <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	Cond <sub>4</sub>	27 <sub>5</sub>		

Cond <sub>4</sub>	Mnem.	Meaning	Test
<b>Integer Compare Results</b>			
0	EQ	= equal	a == b
1	NE	< > not equal	a <> b
2	LT	< less than	a < b
3	LE	<= less than or equal	a <= b
4	GE	>= greater than or equal	a >= b
5	GT	> greater than	a > b
6	BC	Bit clear	!a[b]
7	BS	Bit set	a[b]
8	BCI	Bit clear immediate	!a[b]
9	BSI	Bit set immediate	a[b]
10	LO / CS	< unsigned less than	a < b
11	LS	<= unsigned less than or equal	a <= b
12	HS / CC	unsigned greater than or equal	a >= b
13	HI	unsigned greater than	a > b
14	RA	Branch always	1
15	SR	Branch subroutine	1

**Clock Cycles: 4**



## BBC – Branch if Bit Clear

### Description:

This instruction branches to the target address if bit Rb of Ra is clear, otherwise program execution continues with the next instruction. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ . For a further description see Branch Instructions. Testing a bit beyond bit 63 requires an immediate postfix.

### Instruction Format: B

BBC t0,t1,.label

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	6 <sub>4</sub>	27 <sub>5</sub>		

### Instruction Format: B

BBC t0,3,.label

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Imm <sub>6</sub>	Rm <sub>6</sub>	8 <sub>4</sub>	27 <sub>5</sub>		

### Operation:

If (Ra.bit[Rb] == 0)  
PC = PC + Constant

**Execution Units:** Branch

**Exceptions:** none

**Notes:**

## BBS – Branch if Bit Set

### Description:

This instruction branches to the target address if bit Rb of Ra is set, otherwise program execution continues with the next instruction. For a further description see Branch Instructions. Testing a bit beyond bit 63 requires an immediate postfix.

### Instruction Format: B

BBS t0,t1,.label

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	7 <sub>4</sub>	27 <sub>5</sub>		

### Instruction Format: B

BBS t0,3,.label

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Imm <sub>6</sub>	Rm <sub>6</sub>	9 <sub>4</sub>	27 <sub>5</sub>		

### Operation:

If (Ra.bit[Rb] == 1)  
PC = PC + Constant

**Execution Units:** Branch

**Exceptions:** none

**Notes:**

## BCC –Branch if Carry Clear

BCC Rm, Rn, label

### Description:

Branch if the carry would be set when comparing the first source operand to the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	12 <sub>4</sub>	27 <sub>5</sub>		

Clock Cycles: 4

## BCS –Branch if Carry Set

BCS Rm, Rn, label

### Description:

This is an alternate mnemonic for the [BLO](#) instruction. Branch if the carry would be set because of the comparison of the first operand to the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	10 <sub>4</sub>	27 <sub>5</sub>		

**Clock Cycles: 4**

# BEQ –Branch if Equal

BEQ Rm, Rn, label

## Description:

Branch if two source operands are equal. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

## Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	0 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**

## BGE –Branch if Greater Than or Equal

BGE Rm, Rn, label

### Description:

Branch if the first source operand is greater than or equal to the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as signed integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	4 <sub>4</sub>	27 <sub>5</sub>		

**Clock Cycles: 4**

## BGEU –Branch if Unsigned Greater Than or Equal

BGEU Rm, Rn, label

### Description:

Branch if the first source operand is greater than or equal to the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	12 <sub>4</sub>	27 <sub>5</sub>		

**Clock Cycles: 4**

## BGT –Branch if Greater Than

BGT Rm, Rn, label

### Description:

Branch if the first source operand is greater than the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as signed integer values.

The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	5 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**



## BGTU –Branch if Unsigned Greater Than

BGTU Rm, Rn, label

### Description:

Branch if the first source operand is greater than the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	13 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**

## BHI –Branch if Higher

BHI Rm, Rn, label

### Description:

This is an alternate mnemonic for the [BGU](#) instruction. Branch if the first source operand is greater than the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	13 <sub>4</sub>	27 <sub>5</sub>		

**Clock Cycles: 4**

## BHS –Branch if Higher or Same

BHS Rm, Rn, label

### Description:

Branch if the first source operand is greater than or equal to the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	12 <sub>4</sub>	27 <sub>5</sub>		

**Clock Cycles: 4**

## BLE –Branch if Less Than or Equal

BLE Rm, Rn, label

### Description:

Branch if the first source operand is less than or equal to the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as signed integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>			0 <sub>2</sub>	Rn <sub>6</sub>		Rm <sub>6</sub>		3 <sub>4</sub>	27 <sub>5</sub>	

**Clock Cycles: 4**

## BLEU –Branch if Unsigned Less Than or Equal

BLEU Rm, Rn, label

### Description:

Branch if the first source operand is less than or equal to the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	11 <sub>4</sub>	27 <sub>5</sub>		

**Clock Cycles: 4**

## BLO –Branch if Lower

BLO Rm, Rn, label

### Description:

This is an alternate mnemonic for the [BLTU](#) instruction. Branch if the first source operand is less than the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	10 <sub>4</sub>	27 <sub>5</sub>		

Clock Cycles: 4

## BLS –Branch if Lower or Same

BLS Rm, Rn, label

### Description:

This is an alternate mnemonic for the [BLTU](#) instruction. Branch if the first source operand is less than or equal to the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	11 <sub>4</sub>	27 <sub>5</sub>		

**Clock Cycles: 4**

## BLT –Branch if Less Than

BLT Rm, Rn, label

### Description:

Branch if the first source operand is less than the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as signed integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	2 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**



## BLTU –Branch if Unsigned Less Than

BLTU Rm, Rn, label

### Description:

Branch if the first source operand is less than the second. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as unsigned integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	10 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**

## BNE –Branch if Not Equal

BNE Rm, Rn, label

### Description:

Branch if two source operands are not equal. The first operand is in a register, the second in a register or an immediate value. Both operands are treated as integer values. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

### Instruction Format: B

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	1 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**

# BRA – Unconditional Branch

## Description:

Unconditionally branch to a new program address. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 256\text{MB}$  or  $\pm 64\text{GB}$ .

## Instruction Format: BL2

39	23	22	11	109	8	5	4	0
Disp <sub>16..0</sub>	Disp <sub>28..17</sub>	0 <sub>2</sub>	14 <sub>4</sub>	27 <sub>5</sub>				

## Instruction Format: LBL2

47	23	22	11	109	8	5	4	0
Disp <sub>24..0</sub>	Disp <sub>36..25</sub>	0 <sub>2</sub>	14 <sub>4</sub>	28 <sub>5</sub>				

## Instruction Format: JSR

### BRA [(Ra)],Limit

This form of the instruction branches to an address formed by adding the program counter to a value loaded from a table following the instruction inline in memory. The address loaded from must be no more than limit bytes after the address of the instruction or an exception will occur.

## Operation:

table address = next pc + (Ra \* scale)

if table address > next pc + limit then table limit exception

39	33	32	31	30	28	27	23	22	21	16	15	14	9	8	7	5	4	0
~7	0	1	~3	Limit <sub>4..0</sub>	1	Ra <sub>6</sub>	Op	0 <sub>6</sub>	0	Sz <sub>3</sub>	24 <sub>5</sub>							

**Clock Cycles: 1**

Op	Operation
0	Load PC LSBs
1	Add to PC

**Clock Cycles: 3**

## BRK – Breakpoint

### Description:

Execute the breakpoint exception. This is a form of the TRAP instruction.

### Instruction Format:

39	5	4	0
0	0 <sub>5</sub>		

### Operation:

The program counter and the status register are pushed on an internal stack. Next the BRK vector is fetched from the exception vector table and jumped to.

## BSR – Branch to Subroutine

### Description:

Branch to a subroutine placing the address of the next instruction in a register. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 256\text{MB}$  or  $\pm 64\text{GB}$ .

### Instruction Format: BL2

39	23	22	11	109	8	5	4	0
Disp <sub>16..0</sub>	Disp <sub>28..17</sub>	Rt <sub>2</sub>	15 <sub>4</sub>	27 <sub>5</sub>				

### Instruction Format: LBL2

47	23	22	11	109	8	5	4	0
Disp <sub>24..0</sub>	Disp <sub>36..25</sub>	Rt <sub>2</sub>	15 <sub>4</sub>	28 <sub>5</sub>				

Rt <sub>2</sub>	Meaning
0	Use Register 56
1	Use Register 57
2	Use Register 58
3	Use Register 59

### Instruction Format: JSR

#### BSR [(Ra)],Limit

This form of the instruction branches to an address formed by adding the program counter to a value loaded from a table following the instruction inline in memory. The address loaded from must be no more than limit bytes after the address of the instruction or an exception will occur.

### Operation:

table address = next pc + (Ra \* scale)

if table address > next pc + limit then table limit exception

39	33	32	31	30	28	27	23	22	21	16	15	14	9	8	7	5	4	0
~7	0	1	~3	Limit <sub>4..0</sub>	1	Ra <sub>6</sub>	Op	Rt <sub>6</sub>	0	Sz <sub>3</sub>	24 <sub>5</sub>							

**Clock Cycles: 4**

Op	Operation
0	Load PC LSBs
1	Add to PC

## DBcc – Decrement and Branch

DBcc Rm, Rn, label

### Description:

Decrement the loop counter and branch if the condition is false and the loop counter is not equal to minus one. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

### Instruction Format:

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	Cm <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	Cond <sub>4</sub>	29 <sub>5</sub>					

## FBcc – Conditional Branch

FBcc Rm, Rn, label

### Description:

Branch if the condition is met. The condition is a relationship between either two registers or a register and an immediate value. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

### Instruction Format: RR

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	2 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	Cond <sub>4</sub>	27 <sub>5</sub>					

Cond <sub>4</sub>	Mnem.	Meaning	Test
<b>Float Compare Results</b>			
0	FBEQ	equal	!nan & eq
1	FBNE	not equal	!eq
2	FBGT	greater than	!nan & !eq & !lt & !inf
3	FBUGT	Unordered or greater than	Nan    (!eq & !lt & !inf)
4	FBGE	greater than or equal	Eq    (!nan & !lt & !inf)
5	FBUGE	Unordered or greater than or equal	Nan    (!lt    eq)
6	FBLT	Less than	Lt & (!nan & !inf & !eq)
7	FBULT	Unordered or less than	Nan   (!eq & lt)
8	FBLE	Less than or equal	Eq   (lt & !nan)
9	FBULE	unordered less than or equal	Nan   (eq   lt)
10	FBGL	Greater than or less than	!nan & (!eq & !inf)
11	FBUGL	Unordered or greater than or less than	Nan   !eq
12	FBORD	Greater than less than or equal / ordered	!nan
13	FBUN	Unordered	Nan
Cond <sub>5</sub>	Mnem.	Meaning	Test
14			
15			

**Clock Cycles: 4**



## FBEQ – Float Branch if Equal

FBEQ Rm, Rn, label

### Description:

Branch if two source operands are equal. The condition is a relationship between either two registers or a register and an immediate value. Values are treated as quad precision floating-point values. Negative and positive zero are treated as equal. If either operand is a NaN the branch will not be taken. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

### Instruction Format: RR

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	2 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	0 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**

## FBGE – Float Branch if Greater Than or Equal

FBGE Rm, Rn, label

### Description:

Branch if the first source operand is greater than or equal to the second. The condition is a relationship between either two registers or a register and an immediate value. Values are treated as quad precision floating-point values. Negative and positive zero are treated as equal. If either operand is a NaN the branch will not be taken. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

### Instruction Format: RR

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16.0</sub>	2 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	4 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**

## FBGT – Float Branch if Greater Than

FBGT Rm, Rn, label

### Description:

Branch if the first source operand is greater than the second. The condition is a relationship between either two registers or a register and an immediate value. Values are treated as quad precision floating-point values. Negative and positive zero are treated as equal. If either operand is a NaN the branch will not be taken. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

### Instruction Format: RR

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16.0</sub>	2 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	2 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**

## FBNE – Float Branch if Not Equal

FBEQ Rm, Rn, label

### Description:

Branch if two source operands are not equal. The condition is a relationship between either two registers or a register and an immediate value. Values are treated as quad precision floating-point values. Negative and positive zero are treated as equal. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

### Instruction Format: RR

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>	2 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	1 <sub>4</sub>	27 <sub>5</sub>					

**Clock Cycles: 4**

## FBUGE – Float Branch if Unordered Greater Than or Equal

FBUGE Rm, Rn, label

### Description:

Branch if the first source operand is greater than or equal to the second or if the operands are unordered. The condition is a relationship between either two registers or a register and an immediate value. Values are treated as quad precision floating-point values. Negative and positive zero are treated as equal. If either operand is a NaN the branch will be taken. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

### Instruction Format: RR

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>		2 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>		5 <sub>4</sub>		27 <sub>5</sub>		

Clock Cycles: 4

## FBUGT – Float Branch if Unordered or Greater Than

FBUGT Rm, Rn, label

### Description:

Branch if the first source operand is greater than the second or if the comparison is unordered. The condition is a relationship between either two registers or a register and an immediate value. Values are treated as quad precision floating-point values. Negative and positive zero are treated as equal. If either operand is a NaN the branch will be taken. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 64\text{kB}$ .

### Instruction Format: RR

39	23	2221	20	15	14	9	8	5	4	0
Disp <sub>16..0</sub>		2 <sub>2</sub>	Rn <sub>6</sub>		Rm <sub>6</sub>		3 <sub>4</sub>	27 <sub>5</sub>		

**Clock Cycles: 4**

# JMP – Jump to Address

## Description:

Compute the effective address and jump to it. If Ra=53 then the program counter is used. If the indirection bit 'I' of the instruction is set then load the address from memory specified by the effective address and jump to it.

## Operation:

$PC = Ra + Rb$  or  $PC = Ra + Imm$

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

**Notes:**

## Instruction Formats:

### JMP d(Ra, Rb)

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
24 <sub>6</sub>	~	0	0	0	Sb	Rb <sub>6</sub>	0	Ra <sub>6</sub>	0	0 <sub>6</sub>	0	0 <sub>6</sub>	0	Sc <sub>3</sub>	2 <sub>5</sub>				

**Clock Cycles: 1**

### JMP Imm<sub>15</sub> (Ra)

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	0	0		Imm <sub>7..0</sub>	0		Ra <sub>6</sub>	0		0 <sub>6</sub>	0	Sz <sub>3</sub>		24 <sub>5</sub>		

**Clock Cycles: 1**

### JMP [d(Ra, Rb)]

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
24 <sub>6</sub>	~	0	1	0	Sb	Rb <sub>6</sub>	0	Ra <sub>6</sub>	Op	0 <sub>6</sub>	0	Sc <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### JMP [Imm<sub>15</sub> (Ra)]

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>15..8</sub>	0	1		Imm <sub>7..0</sub>	0		Ra <sub>6</sub>	Op		0 <sub>6</sub>	0	Sz <sub>3</sub>		24 <sub>5</sub>		

**Clock Cycles: 1**

**BRA [(Ra)],Limit**

This form of the instruction branches to an address formed by adding the program counter to a value loaded from a table following the instruction inline in memory. The address loaded from must be no more than limit bytes after the address of the instruction or an exception will occur.

**Operation:**

table address = next pc + (Ra \* scale)

if table address > next pc + limit then table limit exception

39	33	32	31	30	28	27	23	22	21	16	15	14	9	8	7	5	4	0
~7	0	1	~3	Limit <sub>4..0</sub>	1	Ra <sub>6</sub>	Op	Rt <sub>6</sub>	0	Sz <sub>3</sub>	24 <sub>5</sub>							

**Clock Cycles: 1**

Op	Operation
0	Load PC LSBs
1	Add to PC



# JSR – Jump to Subroutine

## Description:

Compute the effective address and jump to it. The address of the instruction is stored in a register. If Ra=53 then the program counter is used. If the indirection bit 'I' of the instruction is set then load the address from memory specified by the effective address and jump to it.

## Flag Updates:

None.

## Operation:

$R_t = PC$

$PC = Ra + R_b$  or  $PC = Ra + Imm$

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

## Instruction Formats:

### JSR d(Ra, Rb)

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
24 <sub>6</sub>	~	0	0	0	Sb	Rb <sub>6</sub>	0	Ra <sub>6</sub>	0	Rt <sub>6</sub>	0	Sc <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### JSR Imm<sub>15</sub> (Ra)

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	0	0	Imm <sub>7..0</sub>	0	Ra <sub>6</sub>	0	Rt <sub>6</sub>	0	Sz <sub>3</sub>	24 <sub>5</sub>						

**Clock Cycles: 1**

### JSR [d(Ra, Rb)]

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
24 <sub>6</sub>	~	0	1	0	Sb	Rb <sub>6</sub>	0	Ra <sub>6</sub>	Op	Rt <sub>6</sub>	0	Sc <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles: 1**

### JSR [Imm<sub>15</sub> (Ra)]

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>15..8</sub>	0	1	Imm <sub>7..0</sub>	0	Ra <sub>6</sub>	Op	Rt <sub>6</sub>	0	Sz <sub>3</sub>	24 <sub>5</sub>						

**Clock Cycles: 1**

**BSR [(Ra)],Limit**

This form of the instruction branches to an address formed by adding the program counter to a value loaded from a table following the instruction inline in memory. The address loaded from must be no more than limit bytes after the address of the instruction or an exception will occur.

**Operation:**

table address = next pc + (Ra \* scale)

if table address > next pc + limit then table limit exception

39	33	32	31	30	28	27	23	22	21	16	15	14	9	8	7	5	4	0
~7	0	1	~3	Limit <sub>4..0</sub>	1	Ra <sub>6</sub>	Op	Rt <sub>6</sub>	0	Sz <sub>3</sub>	24 <sub>5</sub>							

**Clock Cycles: 1**

Op	Operation
0	Load PC LSBs
1	Add to PC

# LBcc – Long Conditional Branch

LBcc Rm, Rn, label

## Description:

Branch if the condition is met. The condition is a relationship between either two registers or a register and an immediate value. The displacement is relative to the address of the branch instruction. The branch range is +/- 16MB.

A postfix instruction containing an immediate value may follow the branch instruction, in which case the immediate is used instead of Rn. Rn should be set to zero.

## Instruction Format: B

47	23	2221	20	15	14	9	8	5	4	0
Disp <sub>24..0</sub>				Cm <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	Cond <sub>4</sub>	28 <sub>5</sub>		

Cond <sub>4</sub>	Mnem.	Meaning	Test
<b>Integer Compare Results</b>			
0	EQ	= equal	a == b
1	NE	< > not equal	a <> b
2	LT	< less than	a < b
3	LE	<= less than or equal	a <= b
4	GE	>= greater than or equal	a >= b
5	GT	> greater than	a > b
6	BC	Bit clear	!a[b]
7	BS	Bit set	a[b]
8	BCI	Bit clear immediate	!a[b]
9	BSI	Bit set immediate	a[b]
10	LO / CS	< unsigned less than	a < b
11	LS	<= unsigned less than or equal	a <= b
12	HS / CC	unsigned greater than or equal	a >= b
13	HI	unsigned greater than	a > b
14	RA	Branch always	1
15	SR	Branch subroutine	1

**Clock Cycles: 4**

## LBSR – Long Branch to Subroutine

### Description:

Branch to a subroutine placing the address of the next instruction in a register. The displacement is relative to the address of the branch instruction. The branch range is +/- 64GB.

### Instruction Format: LBL2

47	23	22	11	109	8	5	4	0
Disp <sub>24..0</sub>			Disp <sub>36..25</sub>		Rt <sub>2</sub>	15 <sub>4</sub>		27 <sub>5</sub>

# LBBC – Long Branch if Bit Clear

## Description:

This instruction branches to the target address if bit Rb of Ra is clear, otherwise program execution continues with the next instruction. The displacement is relative to the address of the branch instruction. The branch range is  $\pm 16\text{MB}$ . For a further description see Branch Instructions. Testing a bit beyond bit 63 requires an immediate postfix.

## Instruction Format: B

BBC t0,t1,.label

47	23	2221	20	15	14	9	8	5	4	0
Disp <sub>24..0</sub>				0 <sub>2</sub>	Rn <sub>6</sub>	Rm <sub>6</sub>	6 <sub>4</sub>	28 <sub>5</sub>		

## Instruction Format: B

BBC t0,3,.label

47	23	2221	20	15	14	9	8	5	4	0
Disp <sub>24..0</sub>				0 <sub>2</sub>	Imm <sub>6</sub>	Rm <sub>6</sub>	8 <sub>4</sub>	28 <sub>5</sub>		

## Operation:

If (Ra.bit[Rb] == 0)  
PC = PC + Constant

**Execution Units:** Branch

**Exceptions:** none

**Notes:**

# NOP – No Operation

NOP

## Description:

This instruction does not perform any operation. Ty<sub>3</sub> 0 to 2 indicates a postfix instruction, and these codes should not be used for other NOPs. The value 3 to 6 for Ty<sub>3</sub> are reserved. The NOP operation is an opcode of all ones.

## Instruction Format:

39	8	7	5	4	0
0xFFFFFFFF <sub>32</sub>	7 <sub>3</sub>	31 <sub>5</sub>			

## RTD – Return from Subroutine, Deallocate

### Description:

Return from subroutine and deallocate stack. Add two source operands and place the sum in the target register. All registers are treated as integer registers. Arithmetic is signed twos-complement values. The program counter is loaded with the value of the specified link register.

**Supported Operand Sizes:** .b, .w, .t, .o

### Operation:

$$Rt = Ra + Rb \text{ or } Rt = Ra + Imm$$

$$PC = Lr$$

### Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

### Notes:

### Instruction Formats:

#### RTD Rt, Ra, Rb – Register direct

39	34	3332	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
4 <sub>6</sub>	Lr <sub>2</sub>	1	0	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	0	Sz <sub>3</sub>	2 <sub>5</sub>						

**Clock Cycles:** 1

#### RTD Rt,Ra,Imm<sub>16</sub>

39	33	32	31	30	27	2625	2423	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	0	1	Imm <sub>7..4</sub>	0 <sub>2</sub>	Lr <sub>2</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	0	Sz <sub>3</sub>	4 <sub>5</sub>						

**Clock Cycles:** 1

# RTS – Return from Subroutine

## Description:

Return from subroutine. Load the program counter with the contents of the specified link register.

**Supported Operand Sizes:** .b, .w, .t, .o

## Operation:

$Rt = Ra + Rb$  or  $Rt = Ra + Imm$

## Clock Cycles:

**Execution Units:** All Integer ALU's

**Exceptions:** none

## Notes:

## Instruction Formats:

### RTS Rt

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
4 <sub>6</sub>	Lr <sub>2</sub>	1	0	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	0	0 <sub>6</sub>	0	Sz <sub>3</sub>	2 <sub>5</sub>							

**Clock Cycles: 1**

### RTS Rt

39	33	32	31	30	27	26	25	24	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	0	1	Imm <sub>7..4</sub>	0 <sub>2</sub>	Lr <sub>2</sub>	Sa	Ra <sub>6</sub>	0	0 <sub>6</sub>	0	Sz <sub>3</sub>	4 <sub>5</sub>								

**Clock Cycles: 1**



# RTE – Return From Exception

## Instruction Formats:

### RTE Imm<sub>7</sub>

39	33	32	31	30	27	26	25	24	23	22	21	16	15	14	9	8	7	5	4	0
Inline <sub>6..0</sub>	0	1	Arg <sub>7..4</sub>	D <sub>2</sub>	~ <sub>2</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	0	Sz <sub>3</sub>	4 <sub>5</sub>								

**Clock Cycles: 1**

## Field Description:

The inline field is used for the number of bytes to skip past the return address. This is to allow inline subroutine arguments. Up to 128 bytes may be skipped over. For externally triggered interrupts this field should be zero.

D<sub>2</sub> specifies the number of internal stack entries to unstack. It may be used to perform a multi-level return. Legal values for D are 1 or 2. (0 is the RTD instruction). In most cases a single entry is unstacked. If two entries are unstacked a two-up level return will occur.

## Operation:

Optionally pop the status register, condition code group register, and program counter from the internal stack. Add inline bytes to the program counter, and Arg hexis to the stack pointer. If returning from an application trap the status register is not popped from the stack.

# TRAP – Trap

## Description:

Execute trap. The data field is loaded into the specified target register, Rt. The trap number to execute comes from the contents of register Ra or an immediate value encoded in the instruction. The trap number must be between 1 and 511. Trap numbers below 64 are reserved for the system. Trap numbers 64 and above may be used by applications.

Traps below 64 will use the vector base register to lookup the location of the service routine. Traps above 64 will use the application control register to lookup the location of the service routine.

Trap routines should return using an [RTE](#) instruction.

## Instruction Format:

### TRAP Rt,Ra,#Imm<sub>15</sub>

39	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
Imm <sub>14..8</sub>	0	0	Imm <sub>7..0</sub>	Sa	Ra <sub>6</sub>	0	Rt <sub>6</sub>	0	0 <sub>3</sub>	0 <sub>5</sub>						

Clock Cycles: 1

### TRAP Rt, #Vec, #Data

39	33	32	31	30	23	22	21	16	15	14	9	8	7	6	5	4	0
Imm <sub>14..8</sub>	0	1	Imm <sub>7..0</sub>	Vec <sub>7..0</sub>						Rt <sub>6</sub>	0	0 <sub>2</sub>	V <sub>9</sub>	0 <sub>5</sub>			

Clock Cycles: 1

## Operation:

The program counter and the status register are pushed on an internal stack. Next the vector is fetched from the exception vector table and jumped to.

# Memory Operations

## Addressing Modes

Instructions with a  $U_2$  field support four scaled indexed address modes outlined in the table below.

$U_2$	Mnemonic	Update
0	$d(Rb, Rc * Sc)$	None
1	$d(Rb, Rc++ * Sc)$	Auto post-increment
2	$d(Rb, --Rc * Sc)$	Auto pre-decrement
3	$[d(Rb, Rc * Sc)]$	None

Other instructions support register indirect with displacement or simple register indirect addressing.

## Cache Policy

### Load Operations

$Ca_2$	Policy	Qualifier	Comment
0	none	.io	Always read from main memory or I/O
1	Read	.rd	Read from cache if in cache, otherwise read main memory
2	Read, allocate	.rda	Allocate storage in cache, read from cache
3			Reserved

### Store Operations

$Ca_2$	Policy	Qualifier	Comment
0	Write through	.wt	Always write through to main memory
1	Writeback	.wb	Store to main memory only when data not in cache
2	Write through, write allocate	.wta	Write to main memory, and allocate in cache
3	Write back, write allocate	.wba	Allocate in cache, write to cache

# AMADD - Addition

## Description:

Atomically add source operand register Rb to value from memory and store the result back to memory. The original value of the memory cell is stored in register Rt. The memory address is contained in register Ra.

**Supported Operand Sizes:** .t, .o, .h

**Instruction Formats:** AMO

**AMADD Rt, Rb, [Ra]**

39	35	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
4 <sub>5</sub>	aq	rl	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

**AMADD Rt, imm, [Ra]**

39	35	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
20 <sub>5</sub>	aq	rl	Vc	0	Imm <sub>8</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

# AMAND – Bitwise And

## Description:

Bitwise ‘And’ source operand register Rb to value from memory and store the result back to memory. The original value of the memory cell is stored in register Rt. The memory address is contained in register Ra.

**Supported Operand Sizes:** .t, .o, .n

**Instruction Formats:** AMO

**AMAND Rt, Rb, [Ra]**

39	35	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
8 <sub>5</sub>	aq	rl	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

**AMAND Rt, imm, [Ra]**

39	35	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
24 <sub>5</sub>	aq	rl	Vc	0	Imm <sub>8</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

# AMASL – Arithmetic Shift Left

## Description:

Atomically shift left source operand from memory by Rb and store the result back to memory. The original value of the memory cell is stored in register Rt. The memory address is contained in register Ra.

**Supported Operand Sizes:** .t, .o, .n

**Instruction Formats:** AMO

**AMASL Rt, Rb, [Ra]**

39	35	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
6 <sub>5</sub>	aq	rl	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

**AMASL Rt, imm, [Ra]**

39	35	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
22 <sub>5</sub>	aq	rl	Vc	0	Imm <sub>8</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

# AMEOR – Bitwise Exclusive Or

## Description:

Bitwise exclusive ‘Or’ source operand register Rb to value from memory and store the result back to memory. The original value of the memory cell is stored in register Rt. The memory address is contained in register Ra.

**Supported Operand Sizes:** .t, .o, .n

**Instruction Formats:** AMO

**AMEOR Rt, Rb, [Ra]**

39	35	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
10 <sub>5</sub>	aq	rl	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

**AMEOR Rt, imm, [Ra]**

39	35	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
26 <sub>5</sub>	aq	rl	Vc	0	Imm <sub>8</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

# AMLSR – Logical Shift Right

## Description:

Atomically shift right source operand from memory by Rb and store the result back to memory. The original value of the memory cell is stored in register Rt. The memory address is contained in register Ra.

**Supported Operand Sizes:** .t, .o, .n

**Instruction Formats:** AMO

**AMLSR Rt, Rb, [Ra]**

39	35	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
7 <sub>5</sub>	aq	rl	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

**AMLSR Rt, imm, [Ra]**

39	35	34	33	32	31	30	23	22	21	16	15	14	9	8	7	5	4	0
23 <sub>5</sub>	aq	rl	Vc	0	Imm <sub>8</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**



## AMMIN - Minimum

### Description:

If Rb is less than the value from memory, store Rb to memory. The original value of the memory cell is stored in register Rt. The memory address is contained in register Ra. Values are treated as signed two's complement integers. This operation is performed in an atomic fashion.

**Supported Operand Sizes:** .t, .o, .n

**Instruction Formats:** AMO

**AMMIN Rt, Rb, [Ra]**

39	35	34	33	3231	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
2 <sub>5</sub>	aq	rl	0 <sub>2</sub>	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

## AMMINU - Minimum

### Description:

If Rb is less than the value from memory, store Rb to memory. The original value of the memory cell is stored in register Rt. The memory address is contained in register Ra. Values are treated as unsigned integers. This operation is performed in an atomic fashion.

**Supported Operand Sizes:** .t, .o, .n

**Instruction Formats:** AMO

**AMMINU Rt, Rb, [Ra]**

39	35	34	33	3231	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
12 <sub>5</sub>	aq	rl	0 <sub>2</sub>	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

# AMOR – Bitwise Or

## Description:

Bitwise ‘Or’ source operand register Rb to value from memory and store the result back to memory. The original value of the memory cell is stored in register Rt. The memory address is contained in register Ra.

**Supported Operand Sizes:** .t, .o, .n

## Instruction Formats: AMO

### AMOR Rt, Rb, [Ra]

39	35	34	33	3231	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
9 <sub>5</sub>	aq	rl	0 <sub>2</sub>	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

### AMOR Rt, imm, [Ra]

39	35	34	33	3231	30	23	22	21	16	15	14	9	8	7	5	4	0
25 <sub>5</sub>	aq	rl	0	Imm <sub>8</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	26 <sub>5</sub>						

**Clock Cycles:**

# CACHE <cmd>,<ea>

## Description:

Issue command to cache controller.

**Supported Operand Sizes:** N/A

Sz <sub>3</sub>	Ext.	Operand
0		LEA
1		
2	.t	STPTR
3	.o	STPTR
4	.h	STPTR
5		CACHE
6		
7		Indexed

## Instruction Formats: RINDS

**CACHE cmd, d(Rb)**

39	3837	36	31	30	29	28	24	22	21	16	15	9	8	75	4	0
1	~ <sub>2</sub>	Disp <sub>12..7</sub>	Vb	Sb	Rb <sub>6</sub>	St	Cmd <sub>6</sub>	D <sub>6..0</sub>	V	5 <sub>3</sub>	17 <sub>5</sub>					

**Clock Cycles:**

## Instruction Formats: NDXS

**CACHE cmd, d(Rb,Rc\*Sc)**

39	38	37	32	31	30	29	28	24	22	21	16	15	12	11	9	8	75	4	0
1	Vc	Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	St	Cmd <sub>6</sub>	~ <sub>4</sub>	5 <sub>3</sub>	V	7 <sub>3</sub>	17 <sub>5</sub>						

**Clock Cycles:**

Notes:

Cmd <sub>6</sub>	Cache	
???000	Ins.	Invalidate cache
???001	Ins.	Invalidate line
???100	TLB	Invalidate TLB
???101	TLB	Invalidate TLB entry
000???	Data	Invalidate cache
001???	Data	Invalidate line
010???	Data	Turn cache off
011???	Data	Turn cache on

# CMPXCHG – Compare and Exchange

## Description:

If the contents of the addressed memory cell is equal to the contents of Rb then a value is stored to memory from the source register Rc. The original contents of the memory cell are loaded into register Rt. The memory address is contained in register Ra. The memory address must be properly aligned. If the operation was successful then Rt and Rb will be the same value. The compare and swap operation are an atomic operation; no other access is allowed between the load and potential store operation.

**Supported Operand Sizes:** .t, .o, .n

Sz <sub>3</sub>	Ext.	Operand
0	.b	8-bit Byte
1	.w	16-bit Wyde
2	.t	32-bit Tetra
3	.o	64-bit Octa
4	.c	24-bit
5	.p	40-bit
6	.n	96-bit
7		reserved

## Instruction Formats: CMPXCHG

**CMPXCHG Rt, Rb, Rc, [Ra]**

39	38	37	32	31	30	29	28	24	22	21	16	15	14	9	8	7	5	4	0
0	Vc	Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	25 <sub>5</sub>						

**Clock Cycles:**

Notes:

## FLOAD Rn,<ea>

### Description:

Load register Rt from floating-point source. The source value is converted to the machine width; 128-bit quad precision. No rounding needs to take place; the smaller source can always be guaranteed to fit into the target register.

**Supported Operand Sizes:** h, .s, .d, .q

Sz <sub>3</sub>	Ext.	Operand
0		reserved
1	.h	16-bit half
2	.s	32-bit single
3	.d	64-bit double
4	.q	128-bit quad
5		reserved
6		reserved
7		reserved

### Instruction Formats: RINDL

#### FLOAD Rt, d(Rb)

39	3938	37	31	30	29	28	24	22	21	16	15	9	8	7	5	4	0
F	Ca <sub>2</sub>	Disp <sub>13..7</sub>	Vb	Sb	Rb <sub>6</sub>	St	Rt <sub>6</sub>	D <sub>6..0</sub>	V	Sz <sub>3</sub>	16 <sub>5</sub>						

**Clock Cycles:**

### Instruction Formats: NDXL

#### FLOAD Rt, d(Rb,Rc\*Sc)

39	38	37	32	31	30	29	28	24	22	21	16	1514	1312	11	9	8	7	5	4	0
F	Vc	Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	St	Rt <sub>6</sub>	U <sub>2</sub>	Ca <sub>2</sub>	Sz <sub>3</sub>	V	7 <sub>3</sub>	16 <sub>5</sub>						

**Clock Cycles:**

Notes:

F	Load operand type
0	Integer
1	Floating point

Ca <sub>2</sub>	Policy	Qualifier	Comment
0	none	.io	Always read from main memory or I/O
1	Read	.rd	Read from cache if in cache, otherwise read main memory
2	Read, allocate	.rda	Allocate storage in cache, read from cache
3			Reserved

## FSTORE Ra,<ea>

### Description:

Store register Ra to destination. The register is converted from quad precision to the storage precision.

**Supported Operand Sizes:** .h, .s, .d, .t

Sz <sub>3</sub>	Ext.	Operand
0		Reserved
1	.h	16-bit half
2	.s	32-bit single
3	.d	64-bit double
4	.q	128-bit quad
5		
6		
7		Indexed

### Instruction Formats: RINDS

#### FSTORE Ra, d(Rb)

39	38	37	36	31	30	29	28	24	22	21	16	15	9	8	7	5	4	0
F	Ca <sub>2</sub>		Disp <sub>12..7</sub>	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>		D <sub>6..0</sub>	V	Sz <sub>3</sub>						

**Clock Cycles:**

### Instruction Formats: NDXS

#### FSTORE Ra, d(Rb,Rc\*Sc)

39	38	37	32	31	30	29	28	24	22	21	16	15	14	13	12	11	9	8	7	5	4	0	
F	Vc		Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>		2 <sub>2</sub>	Ca <sub>2</sub>		Sz <sub>3</sub>	V								

**Clock Cycles:**

Notes:

F	Store operand type
0	Integer
1	Floating point

Ca <sub>2</sub>	Policy	Comment
0	Write through	Always write through to main memory
1	Writeback	Store to main memory only when data not in cache
2	Write through, write allocate	Write to main memory, and allocate in cache
3	Write back, write allocate	Allocate in cache, write to cache

## LA Ra,<ea>

### Description:

Load address into target register. The address is calculated as if a memory operation were occurring, then it is loaded into the target register.

### Supported Operand Sizes: N/A

Sz <sub>3</sub>	Ext.	Operand
0		LA
1		
2	.t	STPTR
3	.o	STPTR
4	.h	STPTR
5		CACHE
6		
7		Indexed

### Instruction Formats: RINDS

#### LA Rt, d(Rb)

39	38	37	36	31	30	29	28	24	22	21	16	15	9	8	7	5	4	0
1	~ <sub>2</sub>		Disp <sub>12..7</sub>	Vb	Sb	Rb <sub>6</sub>	St	Rt <sub>6</sub>	D <sub>6..0</sub>	V	0 <sub>3</sub>						17 <sub>5</sub>	

Clock Cycles:

### Instruction Formats: NDXS

#### LA Rt, d(Rb,Rc\*Sc)

39	38	37	32	31	30	29	28	24	22	21	16	15	14	13	12	11	9	8	7	5	4	0
1	Vc		Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	St	Rt <sub>6</sub>	U <sub>2</sub>	~ <sub>2</sub>	0 <sub>3</sub>						V		7 <sub>3</sub>	17 <sub>5</sub>	

Clock Cycles:

Notes:

# LOAD Rn,<ea>

## Description:

Load register Rt from source. The source value is sign extended to the machine width. Loading register r53, the stack canary placeholder, will cause a check trap if the value loaded is not equal to the current value of the stack canary register.

The default cache policy is read-allocate, .rda, if not specified.

**Supported Operand Sizes:** .b, .w, .t, .o, .h

Sz <sub>3</sub>	Ext.	Operand
0	.b	8-bit Byte
1	.w	16-bit Wyde
2	.t	32-bit Tetra
3	.o	64-bit Octa
4	.h	128-bit Hexi
5		256-bit
6	.g	512-bit group
7		Indexed/group

## Instruction Formats: RINDL

### LOAD Rt, d(Rb)

39	3837	36	31	30	29	28	24	22	21	16	15	9	8	75	4	0
F	Ca <sub>2</sub>	Disp <sub>12..7</sub>	Vc	Sb	Rb <sub>6</sub>	St	Rt <sub>6</sub>	D <sub>6..0</sub>	V	Sz <sub>3</sub>	16 <sub>5</sub>					

**Clock Cycles:**

## Instruction Formats: NDXL

### LOAD Rt, d(Rb,Rc\*Sc)

39	38	37	32	31	30	29	28	24	22	21	16	1514	1312	11	9	8	75	4	0
F	Vc	Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	St	Rt <sub>6</sub>	U <sub>2</sub>	Ca <sub>2</sub>	Sz <sub>3</sub>	V	7 <sub>3</sub>	16 <sub>5</sub>					

**Clock Cycles:**

Notes:

F	Load operand type
0	Integer
1	Floating point

Ca <sub>2</sub>	Policy	Qualifier	Comment
0	none	.io	Always read from main memory or I/O
1	Read	.rd	Read from cache if in cache, otherwise read main memory
2	Read, allocate	.rda	Allocate storage in cache, read from cache
3			Reserved



# LOADG Gn,<ea>

## Description:

Load a group of four scalar registers from source. The load operation may be masked so that only specific registers of the group are loaded. The instruction will always read 512-bits from memory. This instruction makes use of the REGS modifier to specify additional target registers.

## Supported Operand Sizes: .g

## Instruction Formats: RINDL

### LOADG Rt, d(Rb)

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
37 <sub>6</sub>	DS <sub>1</sub>	Vc	DS <sub>0</sub>	Vb	Sf	Rf <sub>6</sub>	Se	Re <sub>6</sub>	Sd	Rd <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

39	3837	36	31	30	29	28	24	22	21	16	15	9	8	7	5	4	0
0	Ca <sub>2</sub>	Disp <sub>12..7</sub>	~	Sb	Rb <sub>6</sub>	~	Rt <sub>6</sub>	D <sub>6..0</sub>	0	6 <sub>3</sub>	16 <sub>5</sub>						

## Clock Cycles:

## Instruction Formats: NDXL

### LOADG Rt, d(Rb,Rc\*Sc)

39	38	37	32	31	30	29	28	24	22	21	16	1514	1312	11	9	8	7	5	4	0
0	~	Rc <sub>6</sub>	Sc	~	Sb	Rb <sub>6</sub>	~	Rt <sub>6</sub>	U <sub>2</sub>	Ca <sub>2</sub>	6 <sub>3</sub>	0	7 <sub>3</sub>	16 <sub>5</sub>						

## Clock Cycles:

### LOADG Rt, d(Rb)

47	40	39	3837	36	31	30	29	28	24	22	21	16	15	9	8	7	5	4	0
Mask <sub>8</sub>	0	Ca <sub>2</sub>	Disp <sub>12..7</sub>	~	Sb	Rb <sub>6</sub>	~	Rt <sub>6</sub>	D <sub>6..0</sub>	1	6 <sub>3</sub>	16 <sub>5</sub>							

## Clock Cycles:

## Instruction Formats: NDXL

### LOADG Rt, d(Rb,Rc\*Sc)

47	40	39	38	37	32	31	30	29	28	24	22	21	16	1514	1312	11	9	8	7	5	4	0
Mask <sub>8</sub>	0	~	Rc <sub>6</sub>	Sc	~	Sb	Rb <sub>6</sub>	~	Gt <sub>6</sub>	U <sub>2</sub>	Ca <sub>2</sub>	6 <sub>3</sub>	1	7 <sub>3</sub>	16 <sub>5</sub>							

## Clock Cycles:

Notes:

# LOADZ Rn,<ea>

## Description:

Load register Rt from source. The source value is zero extended to the machine width. Loading register r53, the stack canary placeholder, will cause a check trap if the value loaded is not equal to the current value of the stack canary register.

**Supported Operand Sizes:** .b, .w, .t, .o, .p, .n

Sz <sub>3</sub>	Ext.	Operand
0	.b	8-bit Byte
1	.w	16-bit Wyde
2	.t	32-bit Tetra
3	.o	64-bit Octa
4	.h	128-bit Hexi
5		
6		
7		indexed

## Instruction Formats: RINDL

### LOADZ Rt, d(Rb)

39	3837	36	31	30	29	28	24	22	21	16	15	9	8	75	4	0
F	Ca <sub>2</sub>	Disp <sub>12..7</sub>	Vc	Sb	Rb <sub>6</sub>	St	Rt <sub>6</sub>	D <sub>6..0</sub>	V	Sz <sub>3</sub>	17 <sub>5</sub>					

**Clock Cycles:**

## Instruction Formats: NDXL

### LOADZ Rt, d(Rb,Rc\*Sc)

39	38	37	32	31	30	29	28	24	22	21	16	1514	1312	11	9	8	75	4	0
F	Vc	Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	St	Rt <sub>6</sub>	O <sub>2</sub>	Ca <sub>2</sub>	Sz <sub>3</sub>	V	7 <sub>3</sub>	17 <sub>5</sub>					

**Clock Cycles:**

Notes:

F	Store operand type
0	Integer
1	LEA, CACHE, STPTR

Ca <sub>2</sub>	Policy	Qualifier	Comment
0	none	.io	Always read from main memory or I/O
1	Read	.rd	Read from cache if in cache, otherwise read main memory
2	Read, allocate	.rda	Allocate storage in cache, read from cache
3			Reserved

# STORE Ra,<ea>

## Description:

Store register Ra to destination. The default cache policy is write-through, .wt.

**Supported Operand Sizes:** .b, .w, .t, .o, .p, .n

Sz <sub>3</sub>	Ext.	Operand
0	.b	8-bit Byte
1	.w	16-bit Word
2	.t	32-bit Tetra
3	.o	64-bit Octa
4	.h	128-bit
5		256-bit
6	.g	512-bit group
7		indexed

## Instruction Formats: RINDS

### STORE Ra, d(Rb)

39	3837	36	31	30	29	28	24	22	21	16	15	9	8	75	4	0
F	Ca <sub>2</sub>	Disp <sub>12..7</sub>	Vc	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	D <sub>6..0</sub>	V	Sz <sub>3</sub>	18 <sub>5</sub>					

**Clock Cycles:**

## Instruction Formats: NDXS

### STORE Ra, d(Rb,Rc\*Sc)

39	38	37	32	31	30	29	28	24	22	21	16	1514	1312	11	9	8	75	4	0
F	Vc	Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	U <sub>2</sub>	Ca <sub>2</sub>	Sz <sub>3</sub>	V	7 <sub>3</sub>	18 <sub>5</sub>					

**Clock Cycles:**

Notes:

F	Store operand type
0	Integer
1	Floating point

Ca <sub>2</sub>	Policy	Qualifier	Comment
0	Write through	.wt	Always write through to main memory
1	Writeback	.wb	Store to main memory only when data not in cache
2	Write through, write allocate	.wta	Write to main memory, and allocate in cache
3	Write back, write allocate	.wba	Allocate in cache, write to cache

U <sub>2</sub>	Mnemonic	Update
0	d(Rb, Rc*Sc)	None
1	d(Rb, Rc++*Sc)	Auto post-increment

2	$d(Rb, --Rc*Sc)$	Auto pre-decrement
3	$[d(Rb, Rc*Sc)]$	None

Stores using write through will always write through to main memory, and will also update the cache if the data is in the cache. If allocating is specified, the write operation will allocate storage in the cache for data.

Stores using writeback will update memory only when there is a cache collision and new data needs to be stored in the cache. Otherwise, references will be to and from the cache.

# STOREPTR Ra,<ea>

## Description:

Store a pointer contained in register Ra to destination.

Supported Operand Sizes: N/A

Sz <sub>3</sub>	Ext.	Operand
0		LEA
1		
2		STPTR
3		
4		
5		CACHE
6		
7		indexed

## Instruction Formats: RINDS

### STOREPTR Ra, d(Rb)

39	38	37	36	31	30	29	28	24	22	21	16	15	9	8	7	5	4	0
1	~ <sub>2</sub>		Disp <sub>12..7</sub>	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	D <sub>6..0</sub>	V	2 <sub>3</sub>						17 <sub>5</sub>	

Clock Cycles:

## Instruction Formats: NDXS

### STOREPTR Ra, d(Rb,Rc\*Sc)

39	38	37	32	31	30	29	28	24	22	21	16	15	14	13	12	11	9	8	7	5	4	0
1	Vc		Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	U <sub>2</sub>	~ <sub>2</sub>	2 <sub>3</sub>				V				7 <sub>3</sub>		17 <sub>5</sub>

Clock Cycles:

Notes:

# STOREG Gt,<ea>

## Description:

Store a register group to destination. 512-bits are always stored. Zeros may be written instead of specific registers via a mask. One use is to initialize large blocks of memory to zero.

Lower Group (bits 0 to 63)				Upper Group (bits 64 to 127)		
Gn	Group	Registers		Gn	Group	Registers
0	AG0	R0 to R7		8	UAG0	R0 to R7
1	TG0	R8 to R15		9	UTG0	R8 to R15
2	SG0	R16 to R23		10	USG0	R16 to R23
3	SG1	R24 to R31		11	USG1	R24 to R31
4	VMG	R32 to R39		12	VMG	R32 to R39
5	G5	R40 to R47		13	UG5	R40 to R47
6	G6	R48 to R55		14	UG6	R48 to R55
7	G7	R56 to R63		15	UG7	R56 to R63

**Supported Operand Sizes:** .b, .w, .l

**Instruction Formats:** RINDL

**STOREG Ra, d(Rb)**

39	3837	36	31	30	29	28	24	22	21	16	15	9	8	75	4	0
0	Ca <sub>2</sub>	Disp <sub>12..7</sub>	~	Sb	Rb <sub>6</sub>	~	Ga <sub>6</sub>	D <sub>6..0</sub>	0	6 <sub>3</sub>	18 <sub>5</sub>					

**Clock Cycles:**

**Instruction Formats:** NDXS

**STOREG Ra, d(Rb,Rc\*Sc)**

39	38	37	32	31	30	29	28	24	22	21	16	1514	1312	11	9	8	75	4	0
0	~	Rc <sub>6</sub>	Sc	~	Sb	Rb <sub>6</sub>	~	Ga <sub>6</sub>	0 <sub>2</sub>	Ca <sub>2</sub>	6 <sub>3</sub>	0	7 <sub>3</sub>	18 <sub>5</sub>					

**Clock Cycles:**

**STOREG Ra, d(Rb), Mask**

47	40	39	3837	36	31	30	29	28	24	22	21	16	15	9	8	75	4	0
Mask <sub>8</sub>	0	Ca <sub>2</sub>	Disp <sub>12..7</sub>	~	Sb	Rb <sub>6</sub>	~	Ga <sub>6</sub>	D <sub>6..0</sub>	1	6 <sub>3</sub>	18 <sub>5</sub>						

**Clock Cycles:**

**Instruction Formats:** NDXS

**STOREG Ra, d(Rb,Rc\*Sc), Mask**

47	40	39	38	37	32	31	30	29	28	24	22	21	16	1514	1312	11	9	8	75	4	0
Mask <sub>8</sub>	0	~	Rc <sub>6</sub>	Sc	~	Sb	Rb <sub>6</sub>	~	Ga <sub>6</sub>	0 <sub>2</sub>	Ca <sub>2</sub>	6 <sub>3</sub>	1	7 <sub>3</sub>	18 <sub>5</sub>						

**Clock Cycles:**

Notes:

### Compare and Exchange

```
ATOM a0,“AAAAAA”  
LOAD a0,[a3]  
CMP t0,a0,a1  
PEQ t0,”TTF”  
STORE a2,[a3]  
LDI a0,1  
LDI a0,0
```

### Load add and store:

```
ATOM “AAA”  
LOAD a0,[a2]  
ADD t0,a0,a1  
STORE t0,[a2]
```

### Load or and store

```
ATOM “AAA”  
LOAD a0,[a2]  
OR t0,a0,a1  
STORE t0,[a2]
```

### Load and complement and store

```
ATOM “AAA”  
LOAD a0,[a2]  
AND t0,a0,~a1  
STORE t0,[a2]
```

## STORE\_PAIR Rb, Rc, d[Ra]

### Description:

Store register pair to destination.

**Supported Operand Sizes:** .b, .w, .t, .o, .p, .n

Sz <sub>3</sub>	Ext.	Operand
0	.b	8-bit Byte
1	.w	16-bit Word
2	.t	32-bit Tetra
3	.o	64-bit Octa
4	.c	24-bit
5	.p	40-bit
6	.n	96-bit
7		group

### Instruction Formats: dRa

#### STORE Rb, Rc, d(Ra)

39	38	37	32	31	30	29	28	24	22	21	16	15	9	8	7	5	4	0
1	Vc	Rc <sub>6</sub>	Sc	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	D <sub>6..0</sub>	V	Sz <sub>3</sub>	25 <sub>5</sub>						

**Clock Cycles:**

Notes:



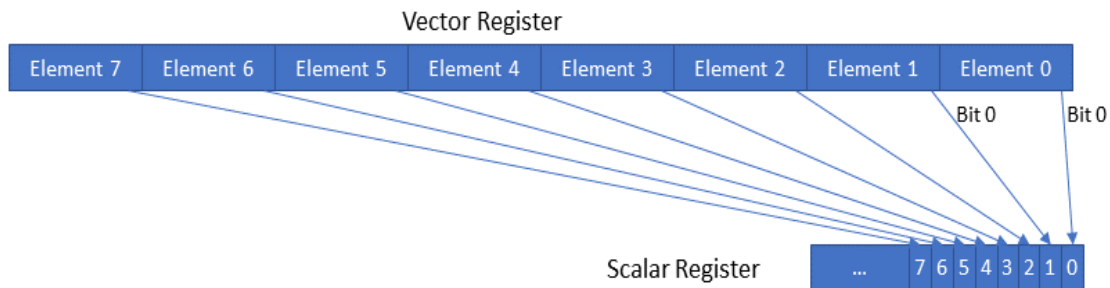
# Vector Specific Instructions

## V2BITS

### Description

Convert Boolean vector to bits. A bit specified by Rb or an immediate of each vector element is copied to the bit corresponding to the vector element in the target register. The target register is a scalar register. Usually, Rb would be zero so that the least significant bit of the vector is copied.

A typical use is in moving the result of a vector set operation into a mask register.



### Instruction Format: R2

#### V2BITS Rt, Ra, Rb

39	34	33	32	31	30	29	28	23	22	21	16	15	14	9	8	7	5	4	0
40 <sub>5</sub>	~	Vc	0	Vb	Sb	Rb <sub>6</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>						

Clock Cycles: 1

#### V2BITS Rt, Ra, Imm<sub>7</sub>

39	34	33	32	31	30	29	23	22	21	16	15	14	9	8	7	5	4	0
40 <sub>5</sub>	~	Vc	1	Vb	Imm <sub>7</sub>	Sa	Ra <sub>6</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	14 <sub>5</sub>						

Clock Cycles: 1

### Operation

For  $x = 0$  to  $VL-1$

$$Rt.bit[x] = Ra[x].bit[Rb]$$

**Exceptions:** none

### Example:

```
cmp v1,v2,v3      ; compare vectors v2 and v3
v2bits vm1,v1,#8   ; move NE status to bits in m1
vadd v4,v5,v6,vm1   ; perform some masked vector operations
vmuls v7,v8,v9,vm1
vadd v7,v7,v4,vm1
```

# Cryptographic Accelerator Instructions

## AES64DS – Final Round Decryption

### Description:

Perform the final round of decryption for the AES standard. Registers Rb, Ra represent the entire AES state.

### Integer Instruction Format: R3

47	41	49 38	37	36 35	34	29	28 27	26	21	20	15	14	9	8	7	0
50h <sub>7</sub>	m <sub>3</sub>	z	~ <sub>2</sub>	~ <sub>6</sub>	Tb <sub>2</sub>	Rb <sub>6</sub>	Ra <sub>6</sub>	Rt <sub>6</sub>	v	02h <sub>8</sub>						

1 clock cycle / N clock cycles (N = vector length)

### Operation:

$R_t = R_a \& R_b$

**Exceptions:** none

## AES64DSM – Middle Round Decryption

### Description:

Perform a middle round of decryption for the AES standard. Registers Rb, Ra represent the entire AES state.

### Integer Instruction Format: R3

47	41	49 38	37	36 35	34	29	28 27	26	21	20	15	14	9	8	7	0
51h <sub>7</sub>	m <sub>3</sub>	z	~ <sub>2</sub>	~ <sub>6</sub>	Tb <sub>2</sub>	Rb <sub>6</sub>	Ra <sub>6</sub>	Rt <sub>6</sub>	v	02h <sub>8</sub>						

1 clock cycle / N clock cycles (N = vector length)

### Operation:

$R_t = R_a \& R_b$

**Exceptions:** none

## AES64ES – Final Round Encryption

### Description:

Perform the final round of encryption for the AES standard. Registers Rb, Ra represent the entire AES state.

### Integer Instruction Format: R3

47	41	49 38	37	36 35	34	29	28 27	26	21	20	15	14	9	8	7	0
52h <sub>7</sub>	m <sub>3</sub>	z	~ <sub>2</sub>	~ <sub>6</sub>	Tb <sub>2</sub>	Rb <sub>6</sub>	Ra <sub>6</sub>	Rt <sub>6</sub>	v	02h <sub>8</sub>						

1 clock cycle / N clock cycles (N = vector length)

### Operation:

Rt = Ra & Rb

Exceptions: none

## AES64ESM – Middle Round Encryption

### Description:

Perform a middle round of encryption for the AES standard. Registers Rb, Ra represent the entire AES state.

### Integer Instruction Format: R3

47	41	49 38	37	36 35	34	29	28 27	26	21	20	15	14	9	8	7	0
53h <sub>7</sub>	m <sub>3</sub>	z	~ <sub>2</sub>	~ <sub>6</sub>	Tb <sub>2</sub>	Rb <sub>6</sub>	Ra <sub>6</sub>	Rt <sub>6</sub>	v	02h <sub>8</sub>						

1 clock cycle / N clock cycles (N = vector length)

### Operation:

Rt = Ra & Rb

Exceptions: none

## SHA256SIG0

### Description:

Implements the Sigma0 transformation function used in the SHA2-256 and SHA2-224 hash function. Only the low order 32 bits of Ra are operated on. The 32-bit result is sign extended to the machine width.

### Instruction Format: R2

#### SHA256SIG0 Rt, Ra – Register direct

39	34	33	32	31	30	29	24	23	22	21	16	15	14	13	8	7	5	4	0
56 <sub>6</sub>	~	0	0	0	0	0 <sub>6</sub>	Va	Sa	Ra <sub>6</sub>	Vt	St	Rt <sub>6</sub>	6 <sub>3</sub>	2 <sub>5</sub>					

**Clock Cycles: 1**

### Operation:

$$Rt = \text{sign extend}(\text{ror32}(Ra, 7) \wedge \text{ror32}(Ra, 18) \wedge (Ra_{32} \gg 3))$$

**Execution Units:** ALU #0

**Exceptions:** none

## SHA256SIG1

### Description:

Implements the Sigma1 transformation function used in the SHA2-256 and SHA2-224 hash function. Only the low order 32 bits of Ra are operated on. The 32-bit result is sign extended to the machine width.

### Instruction Format: R2

#### SHA256SIG1 Rt, Ra – Register direct

39	34	33	32	31	30	29	24	23	22	21	16	15	14	13	8	7	5	4	0
57 <sub>6</sub>	~	0	0	0	0	0 <sub>6</sub>	Va	Sa	Ra <sub>6</sub>	Vt	St	Rt <sub>6</sub>	6 <sub>3</sub>	2 <sub>5</sub>					

**Clock Cycles: 1**

### Operation:

$$Rt = \text{sign extend}(\text{ror32}(Ra, 17) \wedge \text{ror32}(Ra, 19) \wedge (Ra_{32} \gg 10))$$

**Execution Units:** ALU #0

**Exceptions:** none

## SHA256SUM0

### Description:

Implements the Sum0 transformation function used in the SHA2-256 and SHA2-224 hash function. Only the low order 32 bits of Ra are operated on. The 32-bit result is sign extended to the machine width.

### Instruction Format: R2

#### SHA256SUM0 Rt, Ra – Register direct

39	34	33	32	31	30	29	24	23	22	21	16	15	14	13	8	7	5	4	0
58 <sub>6</sub>	~	0	0	0	0	0 <sub>6</sub>	Va	Sa	Ra <sub>6</sub>	Vt	St	Rt <sub>6</sub>	6 <sub>3</sub>	2 <sub>5</sub>					

### Clock Cycles: 1

### Operation:

$$Rt = \text{sign extend}(\text{ror32}(Ra, 2) \wedge \text{ror32}(Ra, 13) \wedge \text{ror32}(Ra, 22))$$

### Execution Units: ALU #0

### Exceptions: none

## SHA256SUM1

### Description:

Implements the Sum1 transformation function used in the SHA2-256 and SHA2-224 hash function. Only the low order 32 bits of Ra are operated on. The 32-bit result is sign extended to the machine width.

### Instruction Format: R2

#### SHA256SUM1 Rt, Ra – Register direct

39	34	33	32	31	30	29	24	23	22	21	16	15	14	13	8	7	5	4	0
59 <sub>6</sub>	~	0	0	0	0	0 <sub>6</sub>	Va	Sa	Ra <sub>6</sub>	Vt	St	Rt <sub>6</sub>	6 <sub>3</sub>	2 <sub>5</sub>					

### Operation:

$$Rt = \text{sign extend}(\text{ror32}(Ra, 6) \wedge \text{ror32}(Ra, 11) \wedge \text{ror32}(Ra, 25))$$

### Execution Units: ALU #0

### Exceptions: none

## SHA512SIG0

### Description:

Implements the Sigma0 transformation function used in the SHA2-512 hash function.

### Instruction Format: R1

31	25	24	22	21	20	15	14	9	8	7	0
34h <sub>7</sub>	m <sub>3</sub>	z	Ra <sub>6</sub>	Rt <sub>6</sub>	v	01h <sub>8</sub>					

**Clock Cycles:** 1

### Operation:

$$Rt = \text{ror64}(Ra, 1) \wedge \text{ror64}(Ra, 8) \wedge (Ra \gg 7)$$

**Execution Units:** ALU #0

**Exceptions:** none

## SHA512SIG1

### Description:

Implements the Sigma1 transformation function used in the SHA2-512 hash function.

### Instruction Format: R1

31	25	24	22	21	20	15	14	9	8	7	0
35h <sub>7</sub>	m <sub>3</sub>	z	Ra <sub>6</sub>	Rt <sub>6</sub>	v	01h <sub>8</sub>					

**Clock Cycles:** 1

### Operation:

$$Rt = \text{ror64}(Ra, 19) \wedge \text{ror64}(Ra, 61) \wedge (Ra \gg 6)$$

**Execution Units:** ALU #0

**Exceptions:** none

## SHA512SUM0

Description:

Instruction Format: R1

31	25	24 22	21	20	15	14	9	8	7	0
36h <sub>7</sub>	m <sub>3</sub>	z	Ra <sub>6</sub>	Rt <sub>6</sub>	v	01h <sub>8</sub>				

## SHA512SUM1

Description:

Instruction Format: R1

31	25	24 22	21	20	15	14	9	8	7	0
37h <sub>7</sub>	m <sub>3</sub>	z	Ra <sub>6</sub>	Rt <sub>6</sub>	v	01h <sub>8</sub>				

## SM3P0

Description:

Instruction Format: R1

31	25	24 22	21	20	15	14	9	8	7	0
38h <sub>7</sub>	m <sub>3</sub>	z	Ra <sub>6</sub>	Rt <sub>6</sub>	v	01h <sub>8</sub>				

# SM3P1

Description:

Instruction Format: R1

31	25	24 22	21	20	15	14	9	8	7	0
39h <sub>7</sub>	m <sub>3</sub>	z	Ra <sub>6</sub>	Rt <sub>6</sub>	v	01h <sub>8</sub>				



## SM4ED

Description:

Instruction Format: R3

47	41	49 38	37	36 35	34	29	28 27	26	21	20	15	14	9	8	7	0
56h <sub>7</sub>	m <sub>3</sub>	z	Tc <sub>2</sub>	Rc <sub>6</sub>	Tb <sub>2</sub>		Rb <sub>6</sub>		Ra <sub>6</sub>		Rt <sub>6</sub>	v				02h <sub>8</sub>

## SM4KS

Description:

Instruction Format: R3

47	41	49 38	37	36 35	34	29	28 27	26	21	20	15	14	9	8	7	0
57h <sub>7</sub>	m <sub>3</sub>	z	Tc <sub>2</sub>	Rc <sub>6</sub>	Tb <sub>2</sub>	Rb <sub>6</sub>	Ra <sub>6</sub>	Rt <sub>6</sub>	v	02h <sub>8</sub>						

# Modifiers

## ATOM

### Description:

Treat the following sequence of instructions as an “atom”. Rt specifies the register results are to be written to.

Disable interrupts for the following instructions.

MASK Modifier Scope	Mask Bit	
	0,1	Instruction zero
	2,3	Instruction one
	4,5	Instruction two
	6,7	Instruction three
	8,9	Instruction four
	10,11	Instruction five
	12,13	Instruction six
	14,15	Instruction seven

Mask Bit	Meaning
00	No action
01	Disable interrupts
10	Disable interrupts and lock bus
11	Reserved

### Instruction Format:

39	34	3332	31	24	23	16	15	14	9	8	7	5	4	0
35 <sub>6</sub>	~2	Imm <sub>15..8</sub>	Imm <sub>7..0</sub>	St	Rt <sub>6</sub>	V	Sz <sub>3</sub>	2 <sub>5</sub>						

### Assembler Syntax:

#### Example:

```

ATOM “LLLLAA”
LOAD a0,[a3]
CMP t0,a0,a1
PEQ t0,”TTF”
STORE a2,[a3]
LDI a0,1
LDI a0,0

```

```

ATOM “LLLL”
LOAD a1,[a3]
ADD t0,a0,a1

```

```
MOV a0,a1  
STORE t0,[a3]
```

# CARRY

## Description:

Apply the carry modifier to following instructions according to a bit mask. This modifier may be used to perform extended precision addition. It may also be used to retrieve the high order multiplier bits or the divide remainder. Note that carry input is not available for the first instruction under the modifier's shadow. Generating carry output for the eight instruction is discarded. Note that postfixes do not count as instructions.

Carry Modifier Scope	Mask Bit	
	0,1	Instruction zero
	2,3	Instruction one
	4,5	Instruction two
	6,7	Instruction three
	8,9	Instruction four
	10,11	Instruction five
	12,13	Instruction six
	14,15	Instruction seven

Mask Bit	Letter	Meaning
00	N	No carry in or out
01	I	Use carry in
10	O	Generate carry out
11	C	Use carry in and generate carry out

## Instruction Format:

39	34	33	32	31	16	15	14	13	8	7	5	4	0
33 <sub>6</sub>	~	0	Imm <sub>15..0</sub>	~	0	Rn <sub>6</sub>	~ <sub>3</sub>	2 <sub>5</sub>					

## Assembler Syntax:

Specifying carry input / output capability for following instructions consists of a map using one of four characters: 'I' for input only, 'O' for output only, 'C' for both input and output and 'N' for neither input or output. A character is present in a string for each following instruction in sequence.

## Example:

CARRY "OCCCCINN" ; first generate carry out, second to fifth use carry in and out, sixth use carry in, seven and eight ignore carry.

ADD r6,r3,r7 ; 'O' gen carry  
 ADD r6,r6,#1234 ; 'C' carry in and carry out  
 ADD r6,r2,r1 ; 'C' carry in and carry out  
 ADD r6,r6,#456 ; 'C' carry in and carry out  
 ADD r7,r6,#456 ; 'C' carry in and carry out  
 ADD r8,r7,#987 ; 'I' carry in  
 MUL r8,r9,r10 ; 'N' no carry in or out



## VMASK (Superseded by mask byte)

### Description:

Apply the vector masking to following instructions according to a bit mask. Note that postfixes do not count as instructions. The mask register for the next four instructions may be specified. Note that the value in the mask register may be inverted. To have all mask bits enabled specify an inverted r0 as the mask register.

MASK Scope Modifier	Mask Bit	
	0 to 6	Instruction one
	7 to 13	Instruction two
	14 to 20	Instruction three
	21 to 27	Instruction four

### Instruction Format:

39	34	33	32	31	26	25	24	19	18	17	12	11	10	5	4	0
34 <sub>6</sub>	~	S3	Msk3	S2	Msk2	S1	Msk1	S0	Msk0	2 <sub>5</sub>						

### Assembler Syntax:

#### sExample:

```
VMASK s0,s1,s2,s3
ADD v6,v3,v7      ; vector mask reg s0
ADD v6,v6,#1234   ; vector mask reg s1
ADD v6,v2,v1      ; vector mask reg s2
ADD v6,v6,#456    ; vector mask reg s3
VMASK t0,t1,t2
ADD v7,v6,#456    ; vector mask reg t0
ADD v8,v7,#987    ; vector mask reg t1
MUL v8,v9,v10     ; vector mask reg t2
```

# PRED

## Description:

Apply the predicate to following instructions according to a bit mask. The predicate may be applied to a maximum of eight instructions. Note that postfixes do not count as instructions.

Pred Modifier Scope	Mask Bit	
	0,1	Instruction zero
	2,3	Instruction one
	4,5	Instruction two
	6,7	Instruction three
	8,9	Instruction four
	10,11	Instruction five
	12,13	Instruction six
	14,15	Instruction seven

Mask Bit	Meaning
00	Always execute (ignore predicate)
01	Execute only if predicate is true
10	Execute only if predicate is false
11	Always execute (ignore predicate)

## Instruction Format:

39	34	33	32	31	16	15	10	9	5	4	0
32 <sub>6</sub>	~	1	Imm <sub>15..0</sub>	Rn <sub>6</sub>	Cond <sub>5</sub>	2 <sub>5</sub>					

## Assembler Syntax:

The predicate condition is part of the mnemonic. 'PEQ' predicates logic if the equals flag in the register containing flags is set. Other conditions work in a similar fashion. After the instruction mnemonic the register containing the predicate flags is specified. Next a character string containing 'T' for True, 'F' for false, or 'I' for ignore for the next eight instructions is present.

## Example:

```
PEQ r2,"TTTTFFII" ; next three execute if true, three after execute if false, two after always execute
MUL r3,r4,r5      ; executes if True
ADD r6,r3,r7      ; executes if True
ADD r6,r6,#1234   ; executes if True
DIV r3,r4,r5      ; executes if FALSE
ADD r6,r2,r1      ; executes if FALSE
ADD r6,r6,#456    ; executes if FALSE
MUL r8,r9,r10     ; always executes
```

# REP

## Description:

This modifier indicates a short series of instructions to repeat while the loop counter condition is met. The repeat modifier includes instructions according to a count specified in the Cnt<sub>3</sub> field. The number of included instructions is one greater than Cnt. Up to eight instructions may be part of the repeat operation. The loop counter may be incremented or decremented for each repeat. Loop counter tests perform signed comparisons. The 15-bit immediate may be overridden with a constant postfix instruction. The constant postfix does not count as an instruction in the loop.

REP is limited to a 32-bit immediate value.

## Instruction Format:

39	34	33	32	31	30	16	15	14	12	11	9	8	7	5	4	0
34 <sub>6</sub>	0	0	0	Immed <sub>15</sub>			D	Cnt <sub>3</sub>		Cnd <sub>3</sub>		0	~ <sub>3</sub>	2 <sub>5</sub>		

D    Meaning

0	Decrement loop counter
1	Increment loop counter

Cnd<sub>3</sub>    Loop Counter Test

0h	Equal	LC == Imm
1h	Not equal	LC != Imm
2h	Signed less than	LC < Imm
3h	Signed less than or equal	LC <= Imm
4h	Signed greater than or equal	LC >= Imm
5h	Signed greater than	LC > Imm
6h	Bit clear	LC[imm] = 0
7h	Bit set	LC[imm] == 1

## Assembler Syntax:



# ROUND

## Description:

Set the rounding mode for following instructions according to a bit mask. Note that postfixes do not count as instructions.

ROUND Modifier Scope	Mask Bit	
	0 to 2	Instruction zero
	3 to 5	Instruction one
	6 to 8	Instruction two
	9 to 11	Instruction three
	12 to 14	Instruction four
	15 to 17	Instruction five
	18 to 20	Instruction six
	21 to 23	Instruction seven

## Instruction Format:

39	34	33	32	31		8	7	5	4	0
36 <sub>6</sub>	~	1		Imm <sub>23..0</sub>		~ <sub>3</sub>			2 <sub>5</sub>	

## Assembler Syntax:

## Example:

# MPU Hardware

## PIC – Programmable Interrupt Controller

### Overview

The programmable interrupt controller manages interrupt sources in the system and presents an interrupt signal to the cpu. The PIC may be used in a multi-CPU system as a shared interrupt controller. The PIC can guide the interrupt to the specified core. If two interrupts occur at the same time the controller resolves which interrupt the cpu sees. While the CPU's interrupt input is only level sensitive the PIC may process interrupts that are either level or edge sensitive. the PIC is a 32-bit I/O device.

### System Usage

There is just a single interrupt controller in the system. It supports 31 different interrupt sources plus a non-maskable interrupt source.

The PIC is located at an address determined by BAR0 in the configuration space.

### Priority Resolution

Interrupts have a fixed priority relationship with interrupt #1 having the highest priority and interrupt #31 the lowest. Note that interrupt priorities are only effective when two interrupts occur at the same time.

### Config Space

A 256-byte config space is supported. Most of the config space is unused. The only configuration is for the I/O address of the register set.

Regno	Width	R/W	Moniker	Description		
000	32	RO	REG_ID	Vendor and device ID		
004	32	R/W				
008	32	RO				
00C	32	R/W				
010	32	R/W	REG_BAR0	Base Address Register		
014	32	R/W	REG_BAR1	Base Address Register		
018	32	R/W	REG_BAR2	Base Address Register		
01C	32	R/W	REG_BAR3	Base Address Register		
020	32	R/W	REG_BAR4	Base Address Register		
024	32	R/W	REG_BAR5	Base Address Register		
028	32	R/W				
02C	32	RO		Subsystem ID		
030	32	R/W		Expansion ROM address		
034	32	RO				
038	32	R/W		Reserved		
03C	32	R/W		Interrupt		

040 to 0FF	32	R/W		Capabilities area		
---------------	----	-----	--	-------------------	--	--

REG\_BAR0 defaults to \$FEE20001 which is used to specify the address of the controller's registers in the I/O address space.

The controller will respond with a memory size request of 0MB (0xFFFFFFFF) when BAR0 is written with all ones. The controller contains its own dedicated memory and does not require memory allocated from the system.

#### Parameters

CFG\_BUS defaults to zero

CFG\_DEVICE defaults to six

CFG\_FUNC defaults to zero

Config parameters must be set correctly. CFG device and vendors default to zero.

## Registers

The PIC contains 40 registers spread out through a 256 byte I/O region. All registers are 32-bit and only 32-bit accessible. There are two different means to control interrupt sources. One is a set of registers that works with bit masks enabling control of multiple interrupt sources at the same time using single I/O accesses. The other is a set of control registers, one for each interrupt source, allowing control of interrupts on a source-by-source basis.

Regno	Access	Moniker	Purpose	
00	R	CAUSE	interrupt cause code for currently interrupting source	
04	RW	RE	request enable, a 1 bit indicates interrupt requesting is enabled for that interrupt, a 0 bit indicates the interrupt request is disabled.	
08	W	ID	Disables interrupt identified by low order five data bits.	
0C	W	IE	enables interrupt identified by low order five data bits	
10			reserved	
14	W	RSTE	resets the edge-sense circuit for edge sensitive interrupts, 1 bit for each interrupt source. This register has no effect on level sensitive sources. This register automatically resets to zero.	
18	W	TRIG	software trigger of the interrupt specified by the low order five data bits.	
20	W	ESL	The low bit for edge sensitivity selection. ESL and ESH combine to form a two bit select of the edge sensitivity.	
			ESH,EHL	Sensitivity
			00	level sensitive interrupt
			01	positive edge sensitive
			10	negative edge sensitive
			11	either edge sensitive
24	W	ESH	The high bit for edge sensitivity selection	
80	RW	CTRL0	control register for interrupt #0	
84	RW	CTRL1	control register for interrupt #1	
...		...		

FC	RW	CTRL31	control register for interrupt #31
----	----	--------	------------------------------------

## Control Register

All the control registers are identical for all interrupt sources, so only the first control register is described here.

Bits

0 to 7	CAUSE	The cause code associated with the interrupt; this register is copied to the cause register when the interrupt is selected.
8 to 10	IRQ	This register determines which signal lines of the cpu are activated for the interrupt. Signal lines are typically used to resolve priority.
16	IE	This is the interrupt enable bit, 1 enables the interrupt, 0 disables it. This is the same bit reflected in the RE register.
17	ES	This bit controls edge sensitivity for the interrupt 0 = level, 1 = pos. edge sensitive. This same bit is present in the ESL register.
18		reserved
19	IRQAR	Respond to an IRQ Ack cycle
20 to 23		reserved
24 to 29	CORE	Core number to select for interrupt processing
30 to 31		reserved

# PIT – Programmable Interval Timer

## Overview

Many systems have at least one timer. The timing device may be built into the cpu, but it is frequently a separate component on its own. The programmable interval timer has many potential uses in the system. It can perform several different timing operations including pulse and waveform generation, along with measurements. While it is possible to manage timing events strictly through software it is quite challenging to perform in that manner. A hardware timer comes into play for the difficult to manage timing events. A hardware timer can supply precise timing. In the test system there are two groups of four timers. Timers are often grouped together in a single component. The PIT is a 64-bit peripheral. The PIT while powerful turns out to be one of the simpler peripherals in the system.

## System Usage

One programmable timer component, which may include up to 32 timers, is used to generate the system time slice interrupt and timing controls for system garbage collection. The second timer component is used to aid the paged memory management unit. There are free timing channels on the second timer component.

Each PIT is given a 64kB-byte memory range to respond to for I/O access. As is typical for I/O devices part of the address range is not decoded to conserve hardware.

PIT#1 is located at \$FFFFFFFFFEE4xxxx

PIT#2 is located at \$FFFFFFFFFEE5xxxx

## Config Space

A 256-byte config space is supported. Most of the config space is unused. The only configuration is for the I/O address of the register set and the interrupt line used.

Regno	Width	R/W	Moniker	Description		
000	32	RO	REG_ID	Vendor and device ID		
004	32	R/W				
008	32	RO				
00C	32	R/W				
010	32	R/W	REG_BAR0	Base Address Register		
014	32	R/W	REG_BAR1	Base Address Register		
018	32	R/W	REG_BAR2	Base Address Register		
01C	32	R/W	REG_BAR3	Base Address Register		
020	32	R/W	REG_BAR4	Base Address Register		
024	32	R/W	REG_BAR5	Base Address Register		
028	32	R/W				
02C	32	RO		Subsystem ID		
030	32	R/W		Expansion ROM address		
034	32	RO				
038	32	R/W		Reserved		

03C	32	R/W		Interrupt		
040 to 0FF	32	R/W		Capabilities area		

REG\_BAR0 defaults to \$FEE40001 which is used to specify the address of the controller's registers in the I/O address space. Note for additional groups of timers the REG\_BAR0 must be changed to point to a different I/O address range. Note the core uses only bits determined by the address mask in the address range comparison. It is assumed that the I/O address select input, cs\_io, will have bits 24 and above in its decode and that a 64kB page is required for the device, matching the MMU page size.

The controller will respond with a mask of 0x00FF0000 when BAR0 is written with all ones.

#### Parameters

CFG\_BUS defaults to zero

CFG\_DEVICE defaults to four

CFG\_FUNC defaults to zero

CFG\_ADDR\_MASK defaults to 0x00FF0000

CFG\_IRQ\_LINE defaults to 29

Config parameters must be set correctly. CFG device and vendors default to zero.

## Parameters

**NTIMER:** This parameter controls the number of timers present. The default is eight. The maximum is 32.

**BITS:** This parameter controls the number of bits in the counters. The default is 48 bits. The maximum is 64.

**PIT\_ADDR:** This parameter sets the I/O address that the PIT responds to. The default is \$FEE40001.

**PIT\_ADDR\_ALLOC:** This parameter determines which bits of the address are significant during decoding. The default is \$00FF0000 for an allocation of 64kB. To compute the address range allocation required, 'or' the value from the register with \$FF000000, complement it then add 1.

## Registers

The PIT has 134 registers addressed as 64-bit I/O cells. It occupies 2048 consecutive I/O locations. All registers are read-write except for the current counts which are read-only. All registers are 64-bit accessible; all 64 bits must be read or written. Values written to registers do not take effect until the synchronization register is written.

Note the core may be configured to implement fewer timers in which case timers that are not implemented will read as zero and ignore writes. The core may also be configured to support fewer bits per count register in which case the unimplemented bits will read as zero and ignore writes.

Regno	Access	Moniker	Purpose
00	R	CC0	Current Count
08	RW	MC0	Max count
10	RW	OT0	On Time
18	RW	CTRL0	Control
20 to 7F8	...	...	Groups of four registers for timer #1 to #63
800	RW	USTAT	Underflow status
808	RZW	SYNC	Synchronization register
810	RW	IE	Interrupt enable
818	RW	TMP	Temporary register
820	RO	OSTAT	Output status
828	RW	GATE	Gate register
830	RZW	GATEON	Gate on register
838	RZW	GATEOFF	Gate off register

## Control Register

This register contains bits controlling the overall operation of the timer.

Bit		Purpose
0	LD	setting this bit will load max count into current count, this bit automatically resets to zero.
1	CE	count enable, if 1 counting will be enabled, if 0 counting is disabled and the current count register holds its value. On counter underflow this bit will be reset to zero causing the count to halt unless auto-reload is set.
2	AR	auto-reload, if 1 the max count will automatically be reloaded into the current count register when it underflows.
3	XC	external clock, if 1 the counter is clocked by an external clock source. The external clock source must be of lower frequency than the clock supplied to the PIT. The PIT contains edge detectors on the external clock source and counting occurs on the detection of a positive edge on the clock source. This bit is forced to 0 for timers 4 to 31.
4	GE	gating enable, if 1 an external gate signal will also be required to be active high for the counter to count, otherwise if 0 the external gate is ignored. Gating the counter using the external gate may allow pulse-width measurement. This bit is forced to 0 for timers 4 to 31.
5 to 63	~	not used, reserved

## Current Count

This register reflects the current count value for the timer. The value in this register will change by counting downwards whenever a count signal is active. The current count may be automatically reloaded at underflow if the auto reload bit (bit #2) of the control byte is set. The current count may also be force loaded to the max count by setting the load bit (bit #0) of the counter control byte.

## Max Count

This register holds onto the maximum count for the timer. It is loaded by software and otherwise does not change. When the counter underflows the current count may be automatically reloaded from the max count register.

## On Time

The on-time register determines the output pulse width of the timer. The timer output is low until the on-time value is reached, at which point the timer output switches high. The timer output remains high until the counter reaches zero at which point the timer output is reset back to zero. So, the on time reflects the length of time the timer output is high. The timer output is low for max count minus the on-time clock cycles.

## Underflow Status

The underflow status register contains a record of which timers underflowed.



Writing the underflow register clears the underflows and disable further interrupts where bits are set in the incoming data. Interrupt processing should read the underflow register to determine which timers underflowed, then write back the value to the underflow register.

## Synchronization Register

The synchronization register allows all the timers to be updated simultaneously. Values written to timer registers do not take effect until the synchronization register is written. The synchronization register must be written with a '1' bit in the bit position corresponding to the timer to update. For instance, writing all one's to the sync register will cause all timers to be updated. The synchronization register is write-only and reads as zero.

## Interrupt Enable Register

Each bit of the interrupt enable register enables the interrupt for the corresponding timer. Interrupts must also be globally enabled by the interrupt enable bit in the config space for interrupts to occur. A '1' bit enables the interrupt, a '0' bit value disables it.

## Temporary Register

This is merely a register that may be used to hold values temporarily.

## Output Status

The output status register reflects the current status of the timers output (high or low). This register is read-only.

## Gate Register

The internal gate register is used to temporarily halt or resume counting for the timer corresponding to the bit position of this register. Writing a value to this register will turn on all timers where there is a '1' bit in the value and turn off all timers where there is a '0' bit in the value.

## Gate On Register

The internal gate 'on' register is used to resume counting for the timer corresponding to the bit position of this register. Writing a value to this register will turn on all timers where there is a '1' bit in the value. Where there is a '0' in the value the timer will not be affected. This register reads as zero.

## Gate Off Register

The internal gate 'off' register is used to halt counting for the timer corresponding to the bit position of this register. Writing a value to this register will turn off all timers where there is a '1' bit in the value. Where there is a '0' in the value the timer will not be affected. This register reads as zero.

# Programming

The PIT is a memory mapped i/o device. The PIT is programmed using 64-bit load and store instructions (LDO and STO). Byte loads and stores (LDB, STB) may be used for control register access. It must reside in the non-cached address space of the system.

## Interrupts

The core is configured use interrupt signal #29 by default. This may be changed with the CFG\_IRQ\_LINE parameter. Interrupts may be globally disabled by writing the interrupt disable bit in the config space with a '1'. Individual interrupts may be enabled or disabled by the setting of the interrupt enable register in the I/O space.

## Testing and Debugging

This section seems short for the amount of testing I do. 90% of the work is in the testing. But this is a book about implementing or developing a processor, not a book about testing. Whole books could easily be written about testing. The key to avoiding backtracking and wasted time down the road is lots of testing along the way. Every bug fix is a test. When one bug is fixed, the next one shows up. Sometimes they seem like a two-headed hydra. Good testing skills are a requirement for developing and debugging a processor. Once you've managed to get such a thing working you're probably an ace at testing. Sometimes the processor and programming cannot help you to find a bug in the processor itself. You must be able to think in terms of 'what test can I do ?' to fix the bug. There are usually a least several wow-zzy bugs. For example, I had a bug where a register exchange instruction only failed on a cache miss, when the instruction was at the end of a cache line. Many programs worked fine, and the processor seemed not to work intermittently. It took quite a while to find. I finally noticed the instruction failed when the cache was turned off. So, one thing to try for testing is turning the cache on or off.

## Test Benches

If you're going to build it there must be some way to perform testing. I'd recommend writing a test-bench first and trying the code in a simulator before trying out the code in an FPGA. A test bench is an artificial environment setup specifically to test a component. Inputs simulating a real environment are sent to the component then the output of the component is monitored for correctness. In the test bench usually so-called corner cases are tested, which are cases testing the extremes to which the component should work. If the component works in the extremes of the test bench it'll certainly work when it's put to real use is the general idea. A simulator is a tool built specifically for running test benches. The simulator has features to aid in debugging logic. One may set breakpoints, points which force the logic to stop at a particular place, and view the outputs of a component.

A simple test bench for the Thor divider circuit is shown below. Note that most test bench files don't have any input or output ports. Instead, signals are selected in the simulator for viewing.

In this case parameters for the divider were manually altered in the test bench to check for specific cases.

```
module Thor_divider_tb();
parameter WID=64;
reg rst;
reg clk;
```

```

reg ld;
wire done;
wire [WID-1:0] qo,ro;

initial begin
    clk = 1;
    rst = 0;
    #100 rst = 1;
    #100 rst = 0;
    #100 ld = 1;
    #150 ld = 0;
end

always #10 clk = ~clk; // 50 MHz

Thor_divider #(WID) u1
(
    .rst(rst),
    .clk(clk),
    .ld(ld),
    .sgn(1'b1),
    .isDivi(1'b0),
    .a(64'd10005),
    .b(64'd27),
    .imm(64'd123),
    .qo(qo),
    .ro(ro),
    .dvByZr(),
    .done(done)
);

```

### endmodule

Note that it is possible to automate test cases and even use file I/O in some tools. Test benches can become quite complex. Test benches for the float components often use a test input file containing the operands for the design under test, DUT, and output the results along with the input operands in a results output file. The output file can then be studied at leisure for issues to correct. Having a file output allows different revisions of the core to be compared and may make regression testing easier.

It is extremely unlikely that one would get the HDL code perfect the first time. The processor is not likely to be working, so how do you fix it up ? One needs debugging dumps of course, and those are only available from a simulator. Judiciously placed debug output can be real aid to getting the cpu working. Unless a fix-up is minor and well-known, I run simulator traces before attempting to run the code in an FPGA.

As a first test running software code in the FPGA try something simple like turning an LED on or off. One of the first lines of code Table888 executes is:

```
start
    sei                                ; disable interrupts
    ld      r1,$FF
    st      r1,LEDS
```

which turns on all the LEDs on the board.

This idea is popular for debugging hardware. The IBM PC had a “post-code” which was a byte value periodically written to an I/O port during startup for debugging. Depending on the display of the byte one could tell where in start-up it failed. Something like a missing or bad display adapter would end up with a specific code.

Another suggestion for test-benches is to use the actual system being loaded into the FPGA device as a component of the test-bench. If one keeps the system simple enough to start with then it’s possible to debug using the test-bench.

## Emulators

An invaluable tool for debugging software prior to the processor being finished is the software emulator. A software emulator is an emulation of the device or system written as a software program to run on a workstation. Software emulators are often significantly slower than the real hardware. It’s also a tool where events applied to the system can be generated by user input. The code for the software emulation of a system mirrors the code for processor implementation itself. The code is just written in a different language. Having an emulator available allows for consistency checks between the emulation and the “real” device. Ideally the emulator should produce the same results as the real device would, except that it’s in a virtual environment of the emulator. The emulator can help resolve software problems that would be too difficult to do using the logic simulator.

Emulators can be cycle-exact, meaning they emulate what happens during each cycle of the processor’s clock. Cycle-exact emulators are often slower than non-cycle exact ones. An emulator that is not cycle exact may only emulate running software, interpreting object code, rather than performing all the internal operations that the CPU does.

## Bootstrap Code vs the “Real Code”

The next thing to do after getting simpler I/O tests working is more complex I/O like a video display. Being able to display things on-screen can be invaluable (a character LCD display or LED display works well too). Many low-cost FPGA boards come with a numeric LED displays for output and buttons for input. It’s slightly more challenging to drive a numeric display and may make a good second test. Also being able to get a keystroke can be valuable too. One of the first routines my processors execute is the clear-screen routine. If it can’t clear the screen I know something’s seriously wrong in the start-up. While the blue screen-of-death may be a bad sign, it’s a good sign at least the processor is working that much. When setting the processor software up (bootstrapping) don’t go for the most complex algorithms to begin with. Go with simple

things. I have two versions of keyboard routines. The one that ‘works the right way’ and the one I use for bootstrapping. The bootstrapping routine goes directly to the keyboard port to read a character. It’s very simple, and pauses the whole machine waiting for a character.

## Data Alignment

Are your variables mysteriously getting over-written ? There could be a problem with address generation in the processor, or perhaps a problem with the external address decoding.

One approach to aligning data structures in memory is to ensure that the structures don’t have partially overlapping addresses. This may help if there are memory addressing problems. For instance, if data structure addresses all end in xxx000, then if there is an address decoding problem, all the structures may get overwritten by values intended for other variables. If the variable addresses are somewhat mangled for example 0xxxx004,xx1018, xx2036 (ending in different LSB’s) then it may be less likely for data to be corrupted. This is a temporary debugging approach. One would want to have the var’s properly listed in a program.

## Get Rid of Complexity

One of the best ways to be able to debug something is to get rid of all the extra complexities involved with it. Many is the time that the author has backtracked on a project and removed features in favor of getting something to work. Add one feature at a time, make it a component that can be easily disabled or removed from the design. Disable the complex features of the design. It’s great to be able to do a complex design. But all the complicated stuff started out small and simple. One doesn’t need caches, interrupts, branch predictors, and so on to have a working design. It’s very rewarding to have even the simplest design working.

## Disabling Interrupts

This bit only applies if you’ve managed to get some sort of interrupt facility working. Several smaller, simpler systems don’t make use of interrupts. The original Apple computer did not use interrupts. Interrupts aren’t something that one must get working right away. They would be part of a longer-term project goal (if at all). Start small and simple and expand from there. There are alternatives to interrupts the main one being polling in a loop.

When working with the real hardware having a set of switches available can be invaluable. The switches can be wired to key signals in the design to offer a manual override option. There may be times when one desires to disable a feature under development while other aspects of the project are taking place. For instance, eventually at some point in time one might want to venture into the world of interrupt processing. Interrupts are a challenge to get working. It’s nice to be able to disable interrupts using an external switch. Also, there are times when one wants to know if the processor is capable of executing a linear sequence of instructions, without the interference of interrupts. Debugging the processor with interrupts enabled can be tricky. Development of an interrupt system is something for a later stage of development. Get the processor running longer sequences of code successfully first before trying to deal with interrupts.

## IRQ Live Indicator

The IRQ live indicator is one of the first debug techniques the author uses once the core can run some code. An indicator that IRQ's are happening seems like a friendly image. It can be useful to see that IRQ's are happening on a regular basis. An IRQ indicator can let one know if the machine is just busy, or really, really stuck. This can be accomplished by incrementing a character at a fixed location on-screen. If that character stops flipping around one knows there's real trouble. Another common approach is to use an LED to indicate the presence of IRQ's. Turning a LED on and off at a low frequency can be handy to visually detect the presence of IRQs.

## Disable Caching

This tip applies only if a cache is present. Implementing a cache isn't priority number one. The first few projects I did, did not include any caching. It was too complex to add a cache to begin with. As mentioned before, it sometimes necessary to disable the cache. Nice-to-have instructions are a cache-on and cache-off instruction. The processor should end up with the same results regardless of whether caching is enabled. If results seem flaky try disabling the cache.

## Clock Frequency

Be conservative when choosing a clock frequency. Don't try to run at the fastest possible frequency until the design is thoroughly debugged. Sometimes changing the clock frequency will provide clues to timing or synchronization problems. If the problem varies with a change in clock frequency, then maybe it's a timing problem. If the problem is consistent regardless of the clock frequency, it's likely some other problem. Note we are dealing with debugging probabilities here. Just because a problem is consistent at different clock frequencies doesn't mean it's not a timing problem.

Another nice aspect of a conservative clock frequency is that the tools used for building the system often work much faster if it's easy for the tools to meet the timing requirements. A conservative clock frequency is a way to speed up the development cycle.

## More Advanced Debugging Options

The following debugging mechanisms fall under the category of being more sophisticated in nature and more difficult to do, but they can sometime prove invaluable. They require interrupts or exceptions.

## Debug Registers

One option that aids primarily software debugging is the presence and use of debug registers. Adding debug registers to the core may make software debugging easier to do. Typically, there are one or more address matching registers that cause an interrupt or exception when the processor's program counter or data address matches the one in the debug register. One must have a working interrupt system for this to be usable.

## Trace / Program Counter History

One of the debug facilities that I've added to cores is the capability to capture the history of the program counter. While the processor is running at full speed, the program counter is stored in a small history table which is usually some sort of shift register. When an exceptional condition occurs in the processor core the history capture is turned off. In the exception processing routine, the program counter history can then be dumped to the screen showing where the program went awry.

The technique is called "trace". A good trace history will often be able to be triggered perhaps at a specific address or via debug match register. The trace may record all instructions, but it is common to record only the branch history, and then a few of the instruction addresses for synchronization purposes. Since branches are either taken or not taken a single bit can be used to record the history making trace very compact. With only a couple of block RAMs a trace history of thousands of instructions is possible.

## Stuck on a Bug ?

This is a brain trick. Try changing the code around in the area of the bug. Sometimes just by changing the code, refactoring without really changing operation, you will be able to spot a bug that wasn't readily apparent. It's a bit like moving your eyes around on the horizon to try and spot an enemy. The action of changing or simply moving the code causes a bug to pop out, out of the shadows.

## The Rare Chance

There is a rare chance that it's a problem in the toolset. A problem like this can make things really difficult, especially if it's a free toolset with no technical support. In about 20 years or so, of using toolsets I've found a few bugs. The toolsets, generally speaking are superb, so the chance of it being a bug in a toolset is extremely remote but not impossible. The one bug I ran into was in extending a complement of a single bit value. The toolset returned a binary "10" the value two when a single bit was being inverted. It should have returned a zero. I was able to work around this problem by zero extending the value manually. I found the bug by tracking the location of it down and dumping values using debug outputs.

Bugs in toolsets are often obvious. The most recent one caused the toolset to crash and quit running depending on how simulation was started. There was a work-around by restarting the simulation fresh every time which takes longer than the usual restart.

If you suspect a bug in the toolset try searching the web for information on it. If it's a common problem it's bound to be posted on the web somewhere. There are also usually forums on the web where one can post about problems, and even sometimes get replies.

# Glossary

## AMO

AMO stands for atomic memory operation. An atomic memory operation typically reads then writes to memory in a fashion that may not be interrupted by another processor. Some examples of AMO operations are swap, add, and, and or.

## Assembler

A program that translates mnemonics and operands into machine code OR a low-level language used by programmers to conveniently translate programs into machine code. Compilers are often capable of generating assembler code as an output.

## ATC

ATC stands for address translation cache. This buffer is used to cache address translations for fast memory access in a system with an mmu capable of performing address translations. The address translation cache is more commonly known as the TLB.

## Base Pointer

An alternate term for frame pointer.

## Burst Access

A burst access is several bus accesses that occur rapidly in a row in a known sequence. If hardware supports burst access the cycle time for access to the device is drastically reduced. For instance, dynamic RAM memory access is fast for sequential burst access, and somewhat slower for random access.

## BTB

An acronym for Branch Target Buffer. The branch target buffer is used to improve the performance of a processing core. The BTB is a table that stores the branch target from previously executed branch instructions. A typical table may contain 1024 entries. The table is typically indexed by part of the branch address. Since the target address of a branch type instruction may not be known at fetch time, the address is speculated to be the address in the branch target buffer. This allows the machine to fetch instructions in a continuous fashion without pipeline bubbles. In many cases the calculated branch address from a previously executed instruction remains the same the next time the same instruction is executed. If the address from the BTB turns out to be incorrect, then the machine will have to flush the instruction queue or pipeline and begin fetching instructions from the correct address.



## Card Memory

A card memory is a memory reserved to record the location of pointer stores in a garbage collection system. The card memory is much smaller than main memory; there may be card memory entry for a block of main memory addresses. Card memory covers memory in 128 to 512-byte sized blocks. Usually, a byte is dedicated to record the pointer store status even though a bit would be adequate, for performance reasons. The location of card memory to update is found by shifting the pointer value to the right some number of bits (7 to 9 bits) and then adding the base address of the table. The update to the card memory needs to be done with interrupts disabled.

## Commit

As in commit stage of processor. This is the stage where the processor is dedicated or committed to performing the operation. There are no prior outstanding exceptions or flow control changes to prevent the instruction from executing. The instruction may execute in the commit stage, but registers and memory are not updated until the retire stage of the processor.

## Decimal Floating Point

Floating point numbers encoded specially to allow processing as decimal numbers. Decimal floating point allows processing every-day decimal numbers rounding in the same manner as would be done by hand.

## Decode

The stage in a processor where instructions are decoded or broken up into simpler control signals. For instance, there is often a register file write signal that must be decoded from instructions that update the register file.

## Diadic

As in diadic instruction. An instruction with two operands.

## Endian

Computing machines are often referred to as big endian or little endian. The endian of the machine has to do with the order bits and bytes are labeled. Little endian machines label bits from right to left with the lowest bit at the right. Big endian machines label bits from left to right with the lowest numbered bit at the left.

## FIFO

An acronym standing for 'first-in first-out'. Fifo memories are used to aid data transfer when the rate of data exchange may have momentary differences. Usually when fifos transfer data the average data rate for input and output is the same. Data is stored in a buffer in order then retrieved from the buffer in order. Uarts often contain fifos.

## FPGA

An acronym for Field Programmable Gate Array. FPGA's consist of a large number of small RAM tables, flip-flops, and other logic. These are all connected with a programmable connection network. FPGA's are 'in the field' programmable, and usually re-programmable. An FPGA's re-programmability is typically RAM based. They are often used with configuration PROM's so they may be loaded to perform specific functions.

## Floating Point

A means of encoding numbers into binary code to allow processing. Floating point numbers have a range within which numbers may be processed, outside of this range the number will be marked as infinity or zero. The range is usually large enough that it is not a concern for most programs.

## Frame Pointer

A pointer to the current working area on the stack for a function. Local variables and parameters may be accessed relative to the frame pointer. As a program progresses a series of "frames" may build up on the stack. In many cases the frame pointer may be omitted, and the stack pointer used for references instead. Often a register from the general register file is used as a frame pointer.

## HDL

An acronym that stands for 'Hardware Description Language'. A hardware description language is used to describe hardware constructs at a high level.

## HLL

An acronym that stands for "High Level Language"

## Instruction Bundle

A group of instructions. It is sometimes required to group instructions together into bundle. For instance, all instructions in a bundle may be executed simultaneously on a processor as a unit. Instructions may also need to be grouped if they are oddball in size for example 41 bits, so that they can be fit evenly into memory. Typically, a bundle has some bits that are global to the bundle, such as template bits, in addition to the encoded instructions.

## Instruction Pointers

A processor register dedicated to addressing instructions in memory. It is also often called a program counter. The program counter got its name because it usually increments (or counts) automatically after an instruction is fetched. In early machines in some rare cases the program counter did not count in a sequential binary fashion, but instead used other forms of a counter such as a grey counter or linear feedback shift

register. In some machines the program counter addresses bundles of instructions rather than individual instructions. This is common with some stack machines where multiple instructions are packed into a memory word.

## Instruction Prefix

An instruction prefix applies to the following instruction to modify its operation. An instruction prefix may be used to add more bits to a following immediate constant, or to add additional register fields for the instruction. The prefix essentially extends the number of bits available to encode instructions. An instruction prefix usually locks out interrupts between the prefix and following instruction.

## Instruction Modifier

An instruction modifier is similar to an instruction prefix except that the modifier may apply to multiple following instructions.

## ISA

An acronym for Instruction Set Architecture. The group of instructions that an architecture supports. ISA's are sometimes categorized at extreme edges as RISC or CISC. RTF64 falls somewhere in between with features of both RISC and CISC architectures.

## Keyed Memory

A memory system that has a key associated with each page to protect access to the page. A process must have a matching key in its key list in order to access the memory page. The key is often 20 bits or larger. Keys for pages are usually cached in the processor for performance reasons. The key may be part of the paging tables.

## Linear Address

A linear address is the resulting address from a virtual address after segmentation has been applied.

## Machine Code

A code that the processing machine is able execute. Machine code is lowest form of code used for processing and is not usually dealt with by programmers except in debugging cases. While it is possible to assemble machine code by hand usually a tool called an assembler is used for this purpose.

## Milli-code

A short sequence of code that may be used to emulate a higher-level instruction. For instance, a garbage collection write barrier might be written as milli-code. Milli-code may use an alternate link register to return to obtain better performance.

# Monadic

An instruction with just a single operand.

# Opcode

A short form for operation code, a code that determines what operation the processor is going to perform. Instructions are typically made up of opcodes and operands.

# Operand

The data that an opcode operates on, or the result produced by the operation. Operands are often located in registers. Inputs to an operation are referred to as source operands, the result of an operation is a destination operand.

# Physical Address

A physical address is the final address seen by the memory system after both segmentation and paging have been applied to a virtual address. One can think of a physical address as one that is “physically” wired to the memory.

# Physical Memory Attributes (PMA)

Memory usually has several characteristics associated with it. In the memory system there may be several different types of memory, rom, static ram, dynamic ram, eeprom, memory mapped I/O devices, and others. Each type of memory device is likely to have different characteristics. These characteristics are called the physical memory attributes. Physical memory attributes are associated with address ranges that the memory is located in. There may be a hardware unit dedicated to verifying software is adhering to the attributes associated with the memory range. The hardware unit is called a physical memory attributes checker (PMA checker).

# Posits

An alternate representation of numbers.

# Program Counter

A processor register dedicated to addressing instructions in memory. It is also often and perhaps more aptly called an instruction pointer. The program counter got its name because it usually increments (or counts) automatically after an instruction is fetched. In early machines in some rare cases the program counter did not count in a sequential binary fashion, but instead used other forms of a counter such as a grey counter or linear feedback shift register. In some machines the program counter addresses bundles of instructions rather than individual instructions. This is common with some stack machines where multiple instructions are packed into a memory word.

## Retire

As in retire an instruction. This is the stage in processor in which the machine state is updated. Updates include the register file and memory. Buffers used for instruction storage are freed.

## ROB

An acronym for ReOrder Buffer. The re-order buffer allows instructions to execute out of order yet update the machine's state in order by tracking instruction state and variables. In FT64 the re-order buffer is a circular queue with a head and tail pointers. Instructions at the head are committed if done to the machine's state then the head advanced. New instructions are queued at the buffer's tail as long as there is room in the queue. Instructions in the queue may be processed out of the order that they entered the queue in depending on the availability of resources (register values and functional units).

## RSB

An acronym that stands for return stack buffer. A buffer of addresses used to predict the return address which increases processor performance. The RSB is usually small, typically 16 entries. When a return instruction is detected at time of fetch the RSB is accessed to determine the address of the next instruction to fetch. Predicting the return address allows the processing core to continuously fetch instructions in a speculative fashion without bubbles in the pipeline. The return address in the RSB may turn out to be detected as incorrect during execution of the return instruction, in which case the pipeline or instruction queue will need to be flushed and instructions fetched from the proper address.

## SIMD

An acronym that stands for 'Single Instruction Multiple Data'. SIMD instructions are usually implemented with extra wide registers. The registers contain multiple data items, such as a 128-bit register containing four 32-bit numbers. The same instruction is applied to all the data items in the register at the same time. For some applications SIMD instructions can enhance performance considerably.

## Stack Pointer

A processor register dedicated to addressing stack memory. Sometimes this register is assigned by convention from the general register pool. This register may also sometimes index into a small dedicated stack memory that is not part of the main memory system. Sometimes machines have multiple stack pointers for different purposes, but they all work on the idea of a stack. For instance, in Forth machines there are typically two stacks, one for data and one for return addresses.

## Telescopic Memory

A memory system composed of layers where each layer contains simplified data from the topmost layer downwards. At the topmost layer data is represented verbatim. At the

bottom layer there may be only a single bit to represent the presence of data. Each layer of the telescopic memory uses far less memory than the layer above. A telescopic memory could be used in garbage collection systems. Normally however the extra overhead of updating multiple layers of memory is not warranted.

## TLB

TLB stands for translation look-aside buffer. This buffer is used to store address translations for fast memory access in a system with an mmu capable of performing address translations.

## Trace Memory

A memory that traces instructions or data. As instructions are executed the address of the executing instruction is stored in a trace memory. The trace memory may then be dumped to allow debugging of software. The trace memory may compress the storage of addresses by storing branch status (taken or not taken) for consecutive branches rather than storing all addresses. It typically requires only a single bit to store the branch status. However, even when branches are traced, periodically the entire address of the program executing is stored. Often trace buffers support tracing thousands of instructions.

## Triadic

An instruction with three operands.

## Vector Length (VL register)

The vector length register controls the maximum number of elements of a vector that are processed. The vector length register may not be set to a value greater than the number of elements supported by hardware. Vector registers often contain more elements than are required by program code. It would be wasteful to process all elements when only a few are needed. To improve the processing performance only the elements up to the vector length are examined.

## Vector Mask (VM)

A vector mask is used to restrict which elements of a vector are processed during a vector operation. A one bit in a mask register enables the processing for that element, a zero bit disables it. The mask register is commonly set using a vector set operation.

## Virtual Address

The address before segmentation and paging has been applied. This is the primary type of address a program will work with. Different programs may use the same virtual address range without being concerned about data being overwritten by another program. Although the virtual address may be the same the final physical addresses used will be different.

# Writeback

A stage in a pipelined processing core where the machine state is updated. Values are 'written back' to the register file.

# Miscellaneous

## Reference Material

Below is a short list of some of the reading material the author has studied. The author has downloaded a fair number of documents on computer architecture from the web. Too many to list.

Modern Processor Design Fundamentals of Superscalar Processors by John Paul Shen, Mikko H. Lipasti. Waveland Press, Inc.

Computer Architecture A Quantitative Approach, Second Edition, by John L Hennessy & David Patterson, published by Morgan Kaufman Publishers, Inc. San Francisco, California is a good book on computer architecture. There is a newer edition of the book available.

Memory Systems Cache, DRAM, Disk by Bruce Jacob, Spencer W. Ng., David T. Wang, Samuel Rodriguez, Morgan Kaufman Publishers

PowerPC Microprocessor Developer's Guide, SAMS publishing. 201 West 103<sup>rd</sup> Street, Indianapolis, Indiana, 46290

80386/80486 Programming Guide by Ross P. Nelson, Microsoft Press

Programming the 286, C. Vieillefond, SYBEX, 2021 Challenger Drive #100, Alameda, CA 94501

Tech. Report UMD-SCA-2000-02 ENEE 446: Digital Computer Design — An Out-of-Order RiSC-16

Programming the 65C816, David Eyes and Ron Lichty, Western Design Centre Inc.

Microprocessor Manuals from Motorola, and Intel,

The SPARC Architecture Manual Version 8, SPARC International Inc, 535 Middlefield Road. Suite 210 Menlo Park California, CA 94025

The SPARC Architecture Manual Version 9, SPARC International Inc, San Jose California, PTR Prentice Hall, Englewood Cliffs, New Jersey, 07632

The MMIX processor: <http://mmix.cs.hm.edu/doc/instructions-en.html>

RISCV 2.0 Spec, Andrew Waterman, Yunsup Lee, David Patterson, Krste Asanović CS Division, EECS Department, University of California, Berkeley {[waterman](mailto:waterman@eecs.berkeley.edu)|[yunsup](mailto:yunsup@eecs.berkeley.edu)|[patterson](mailto:patterson@eecs.berkeley.edu)|[krste](mailto:krste@eecs.berkeley.edu)}@eecs.berkeley.edu

The Garbage Collection Handbook, Richard Jones, Antony Hosking, Eliot Moss published by CRC Press 2012

RISC-V Cryptography Extensions Volume I Scalar & Entropy Source Instructions See [github.com/riscv/riscv-crypto](https://github.com/riscv/riscv-crypto) for more information.

## Trademarks

IBM® is a registered trademark of International Business Machines Corporation. Intel® is a registered trademark of Intel Corporation. HP® is a registered trademark of Hewlett-Packard Development Company. "SPARC® is a registered trademark of SPARC International, Inc.



